

# High-level Kafka Architecture

---



**Paweł Kordek**

DATA ENGINEER

@pawel\_kordek <https://kordek.github.io>



# Overview



**Minimal deployment**

**Message flow**

**Topics and partitions**

**Basic settings**



# Demo




Minimal local Kafka deployment

Admin client

Producers and consumers



# Minimal Kafka Deployment



ZooKeeper



Kafka broker



# ZooKeeper

**Coordination service**

**Used in distributed services**

**Consensus**



# Kafka Broker

**Manages topics and messages**

**Clients connect to the broker**

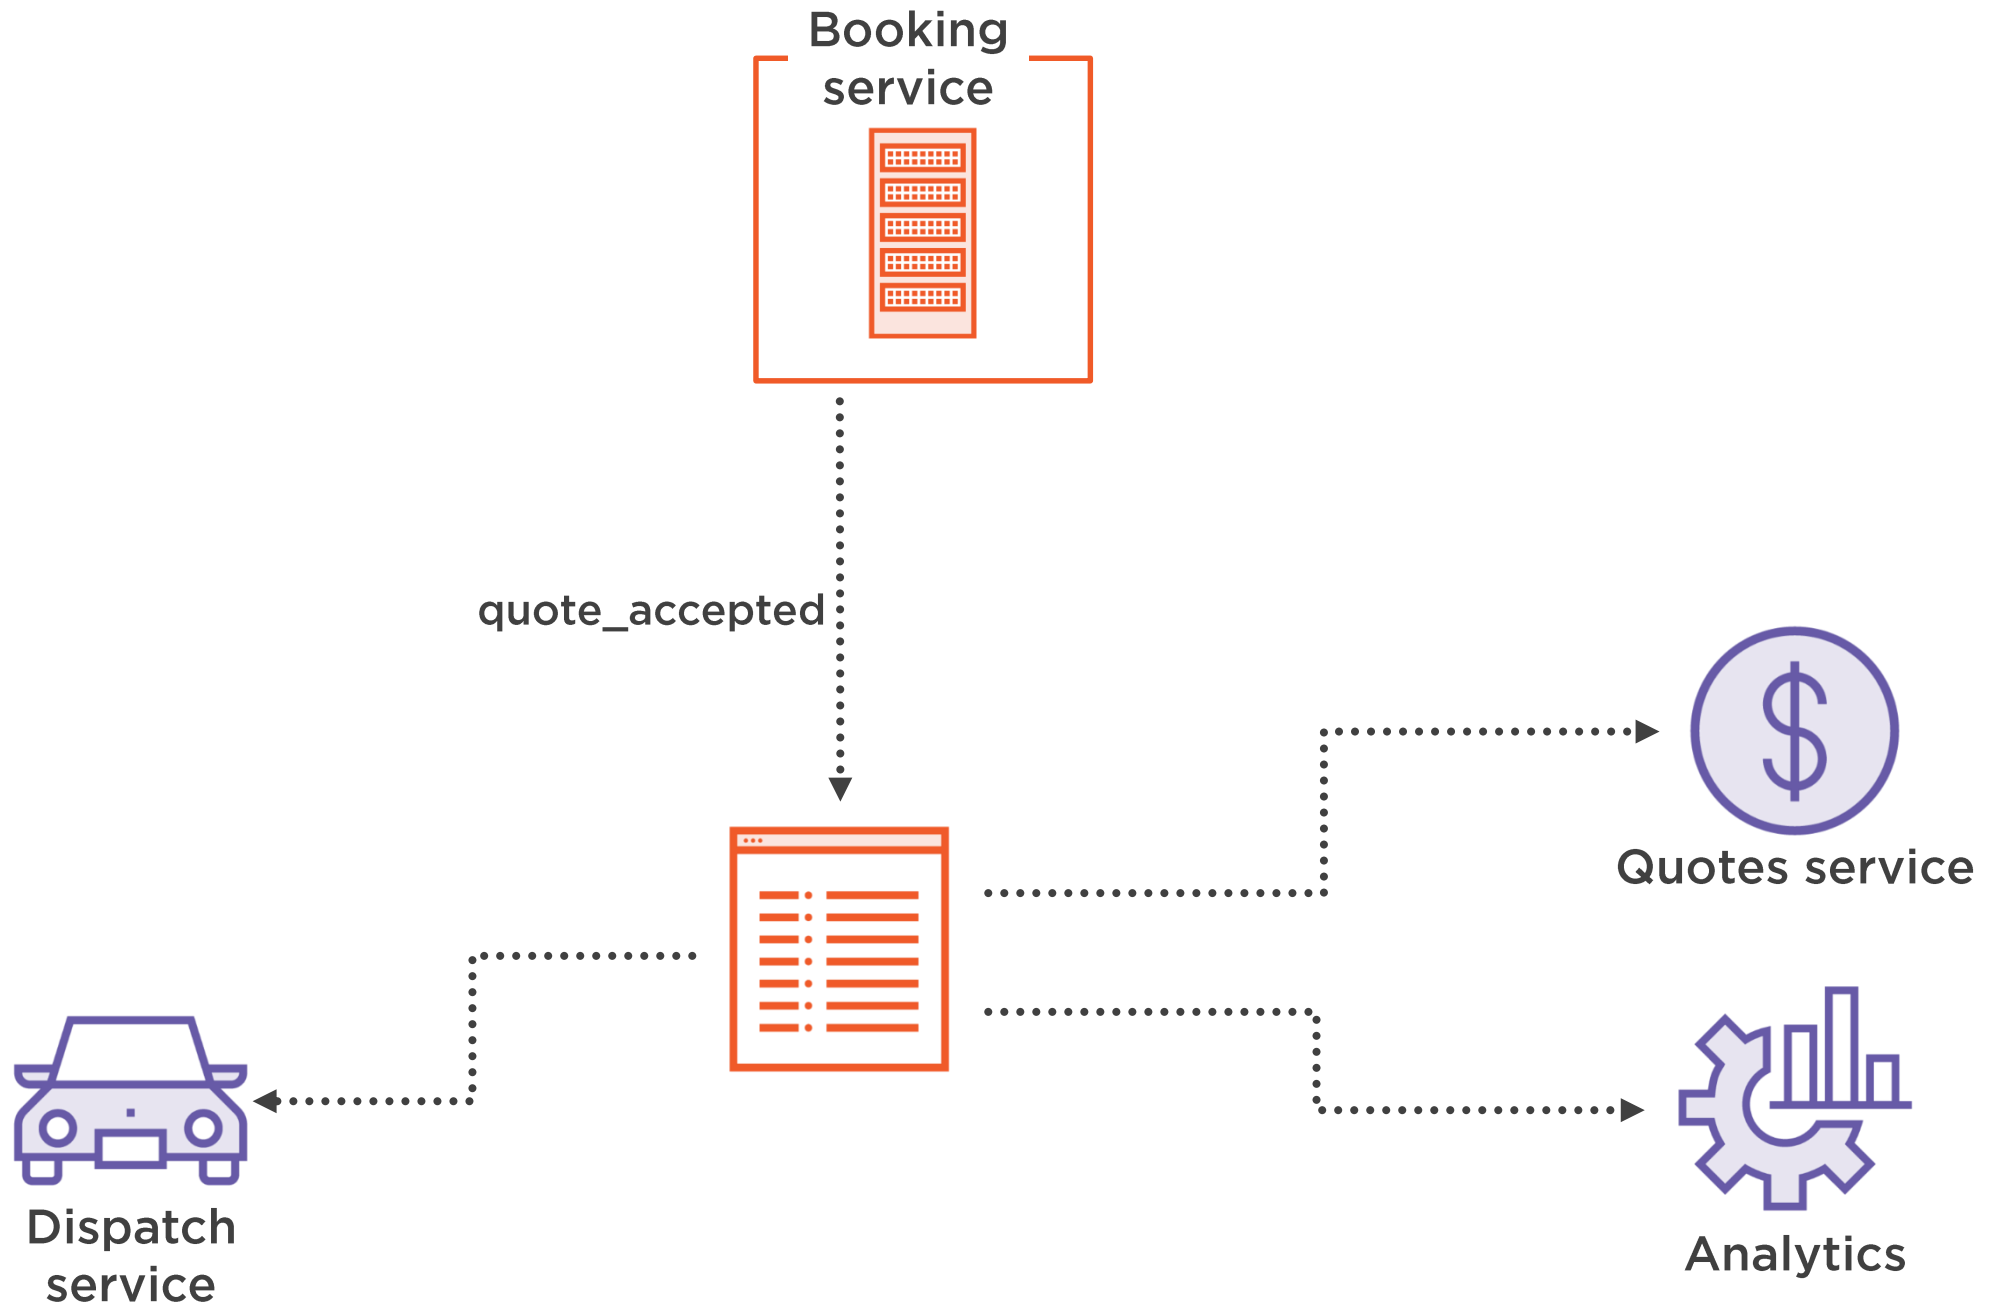
**Core component**



# Interacting with Kafka Broker

---




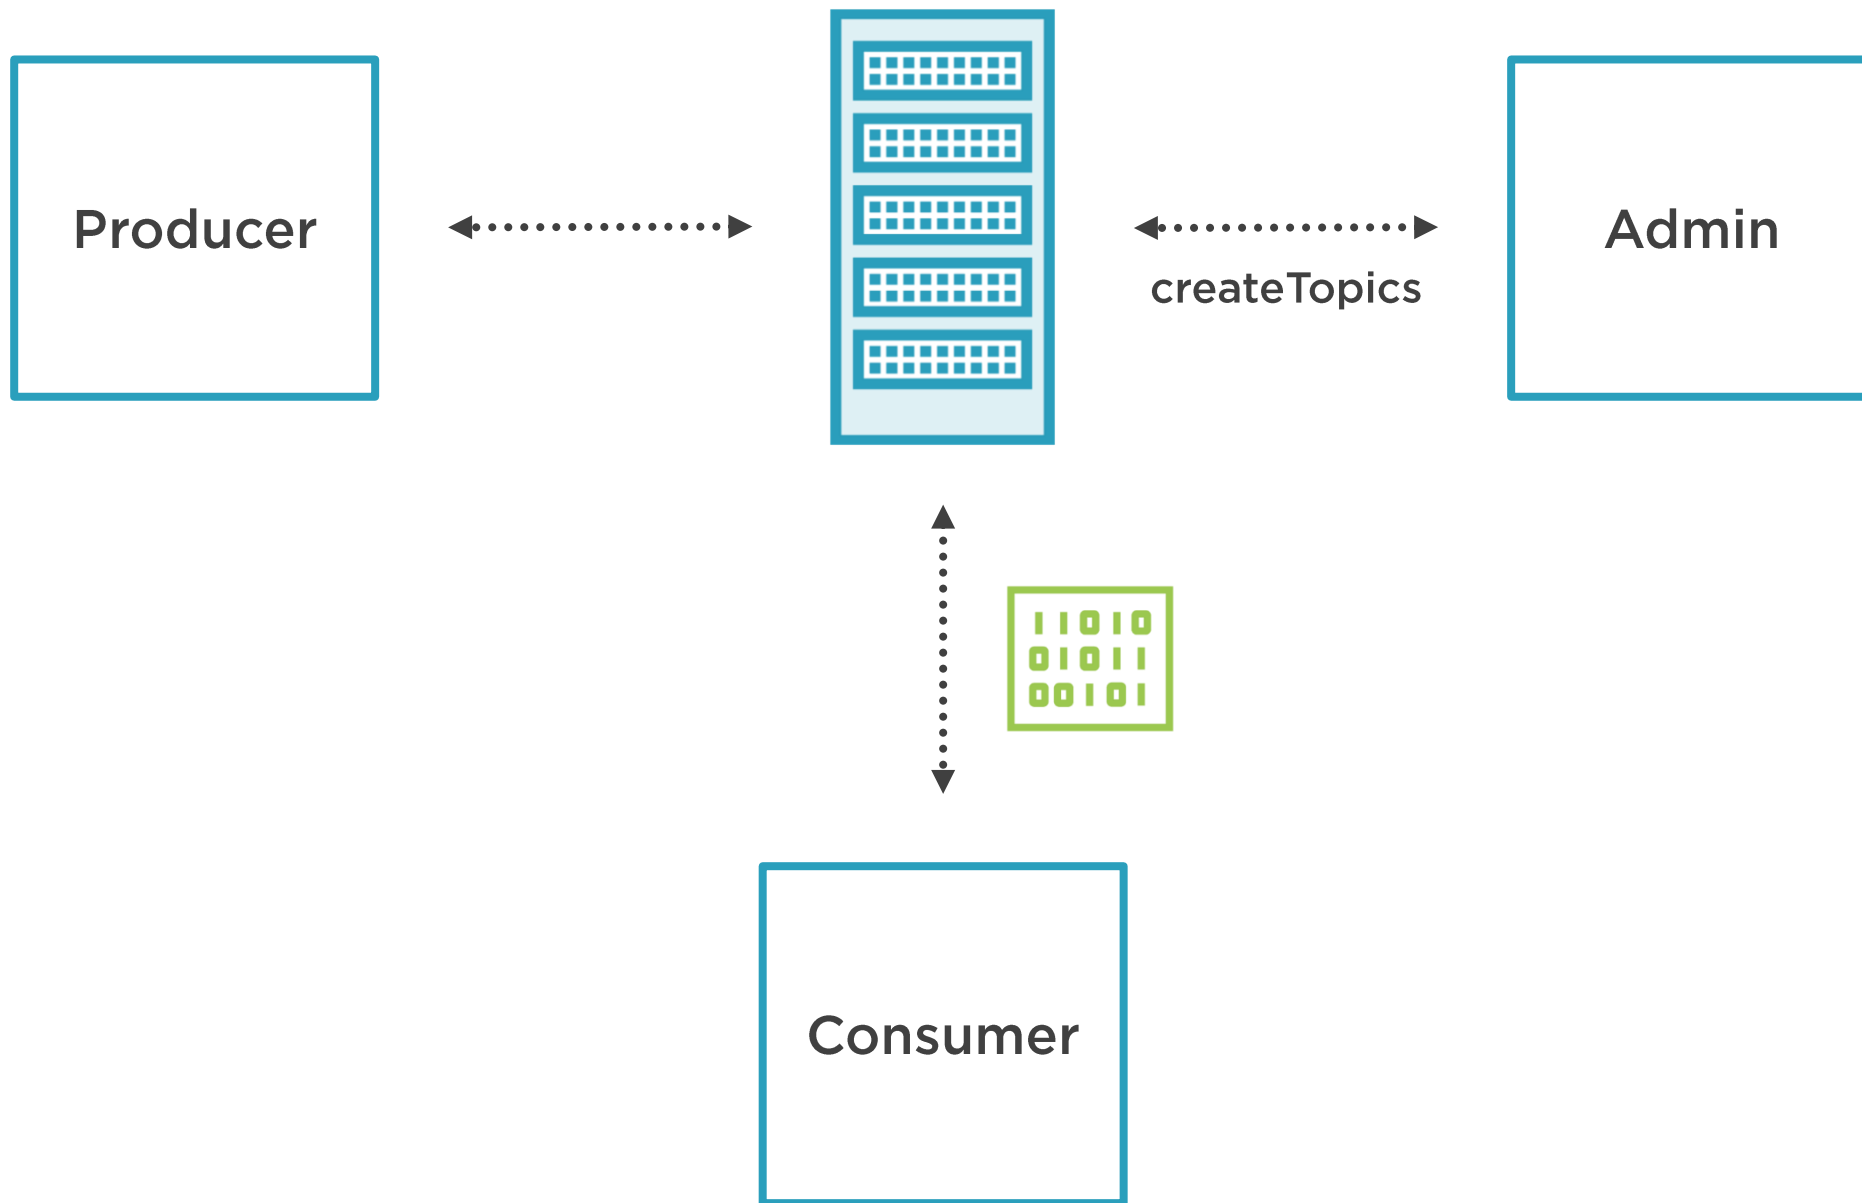




Kafka  
broker

quote-feedback






# Client Libraries

[docs.confluent.io/current/clients](https://docs.confluent.io/current/clients)



# Creating a New Topic

```
Admin admin = Admin.create(  
    Map.of("bootstrap.servers", "localhost:9092")  
);  
  
NewTopic topic = new NewTopic("quote-feedback", ...)  
admin.createTopics(List.of(topic))
```



Returns a **KafkaFuture**



# Creating a New Topic

```
Admin admin = Admin.create(  
    Map.of("bootstrap.servers", "localhost:9092")  
);  
  
NewTopic topic = new NewTopic("quote-feedback", ...)   
admin.createTopics(List.of(topic)).get()
```



# Creating a New Topic

```
Admin admin = Admin.create(  
    Map.of("bootstrap.servers", "localhost:9092")  
);  
  
NewTopic topic = new NewTopic("quote-feedback", ...)   
admin.createTopics(List.of(topic)).all().get()
```



# Running Kafka

---







Make sure you have read the instructions

I will code in my preferred editor

But you will have no problem using  
whatever you like







**Orchestrates  
individual containers  
Deployments defined  
using YAML files**

# Docker Compose



## Kafka cluster

PLAINTEXT\_HOST:  
**kafka1.abc.xyz**

Internal address:  
**192.168.0.14**



External address:  
**kafka2.abc.xyz**



*Hey **kafka2.abc.xyz**, I'd like  
to connect to the cluster!*

Only knows about  
**kafka2.abc.xyz**



## Kafka cluster

PLAINTEXT\_HOST:  
**kafka1.abc.xyz**

Internal address:  
**192.168.0.14**



External address:  
**kafka2.abc.xyz**



Only knows about  
**kafka2.abc.xyz**



*Hey, you need to connect to  
**192.168.0.14**, it says you will  
find it at **kafka1.abc.xyz***



# Partitions, Consumers, and Producers

---



# Partitions




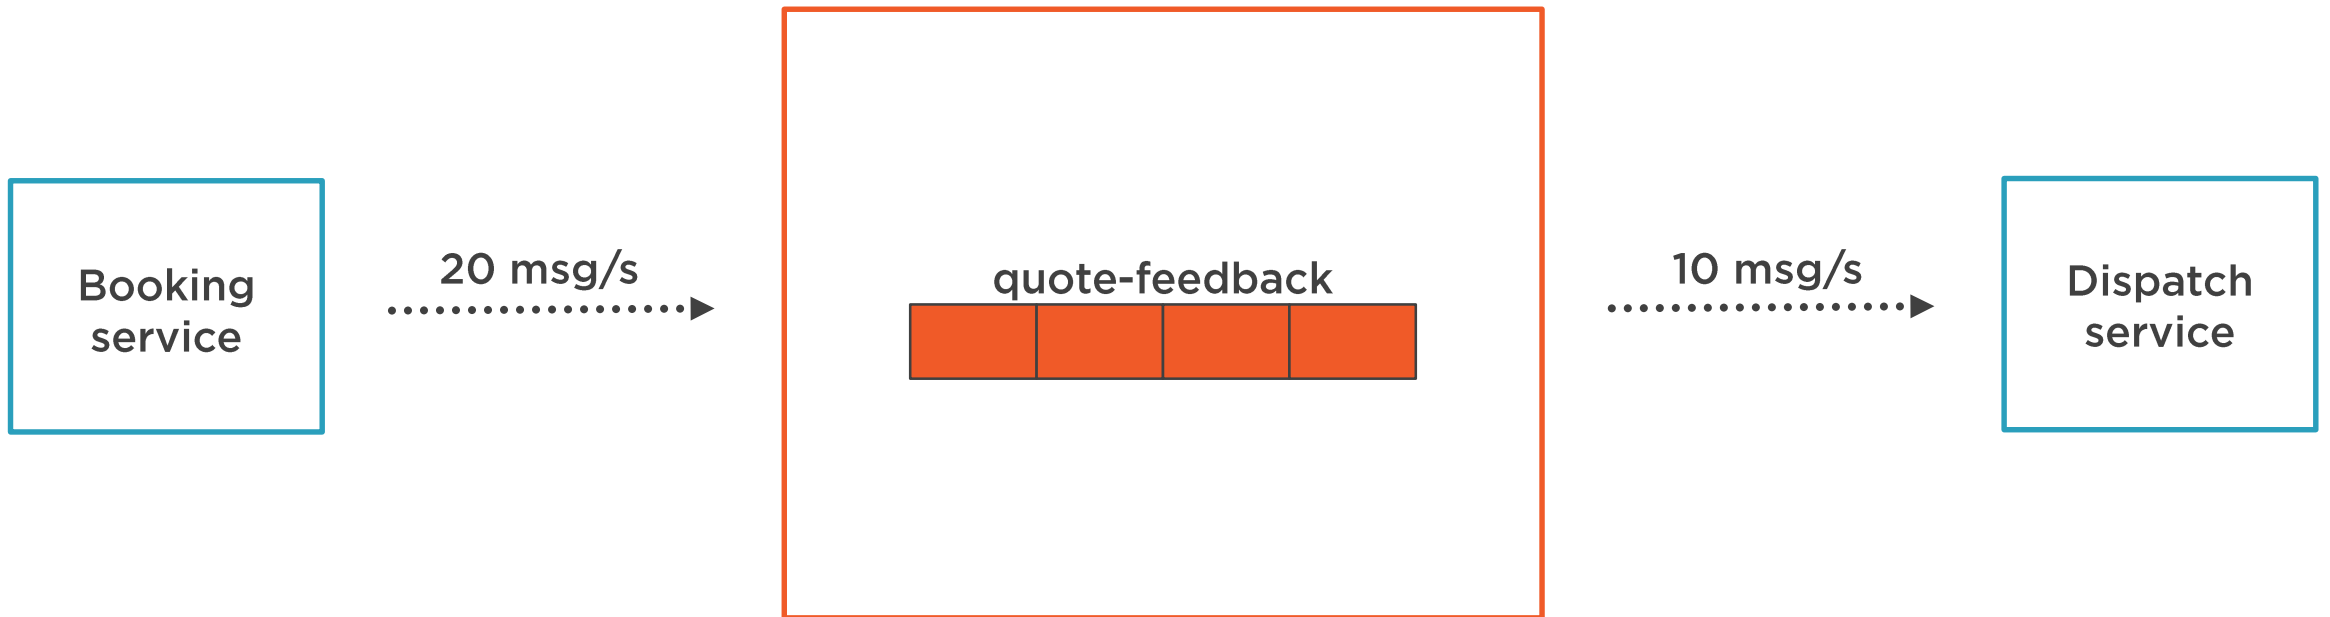
Logical and physical topics' component

Each message belongs to one partition

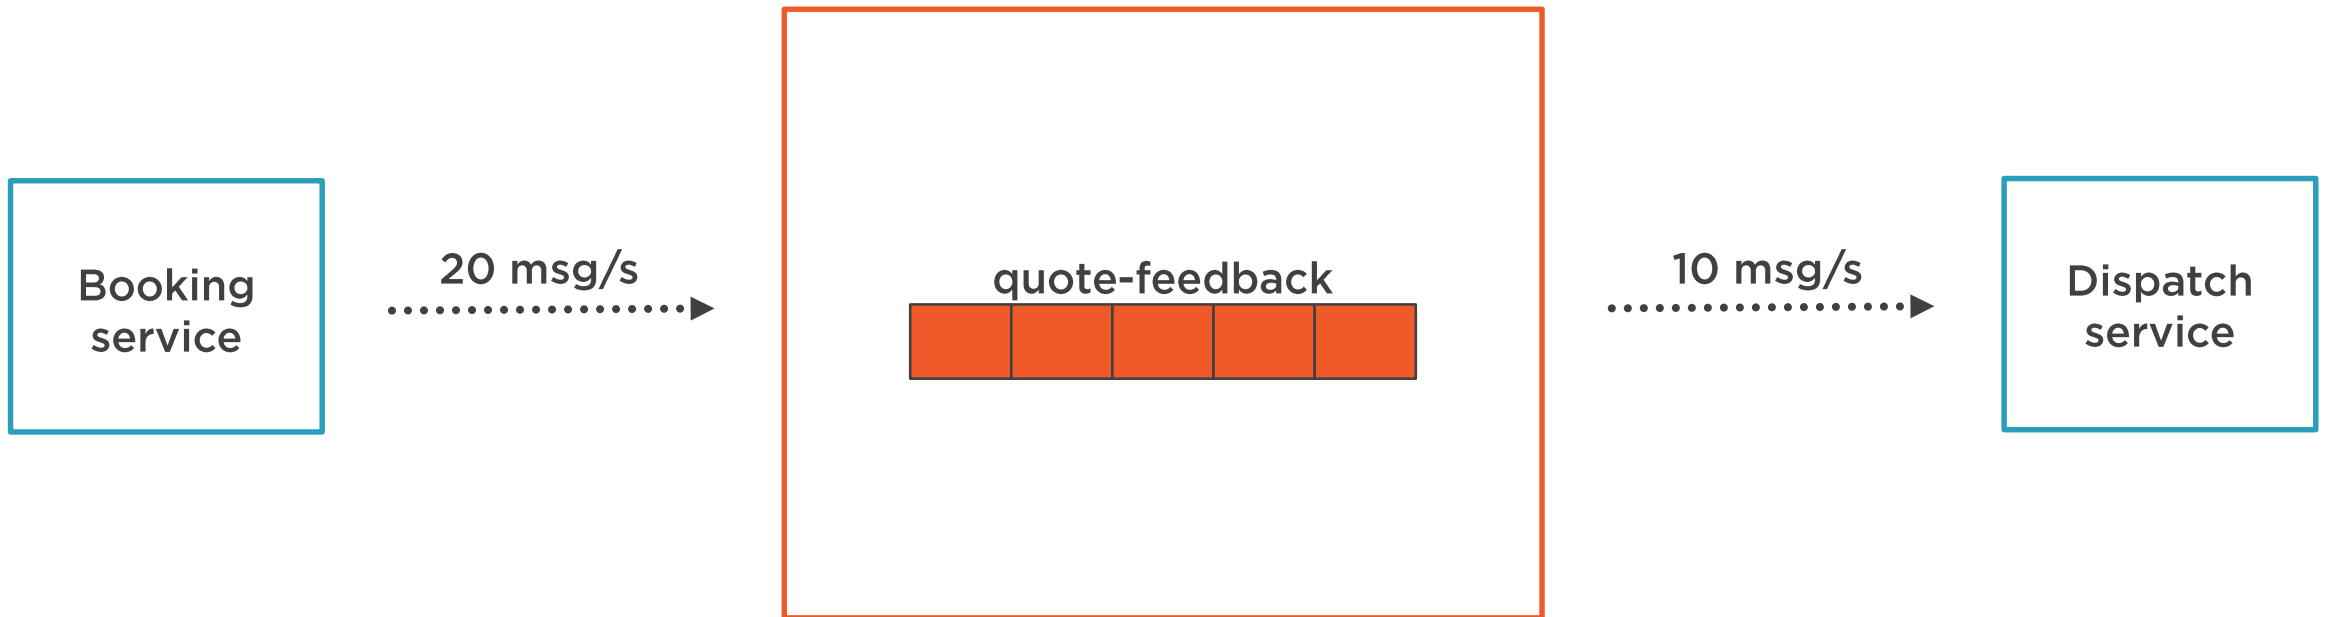
Key to scalability

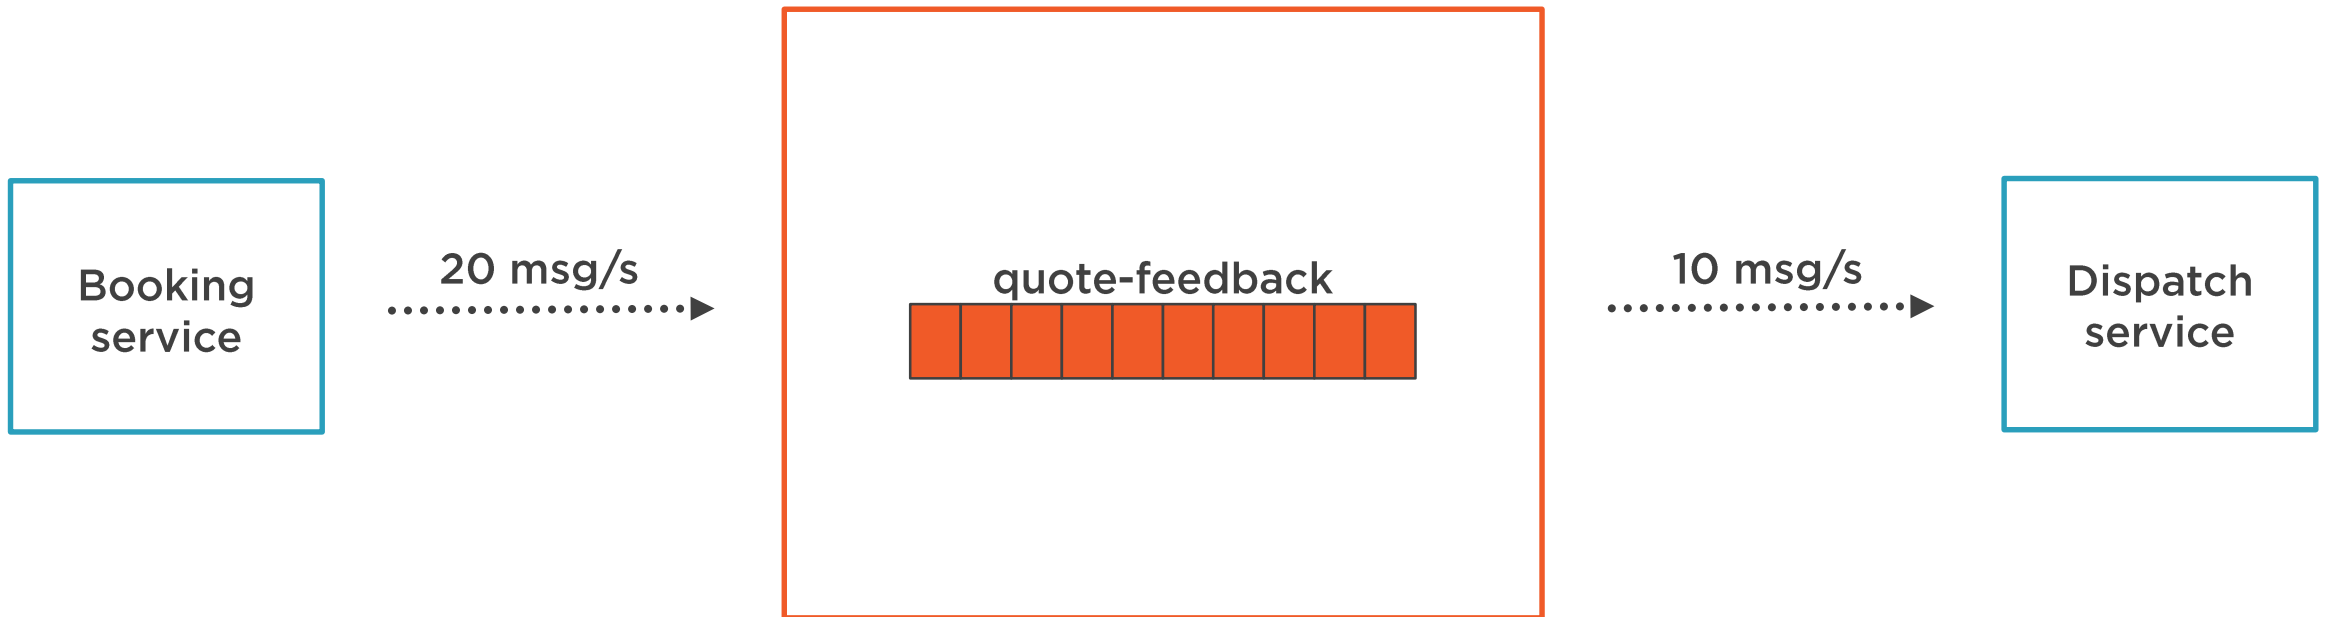


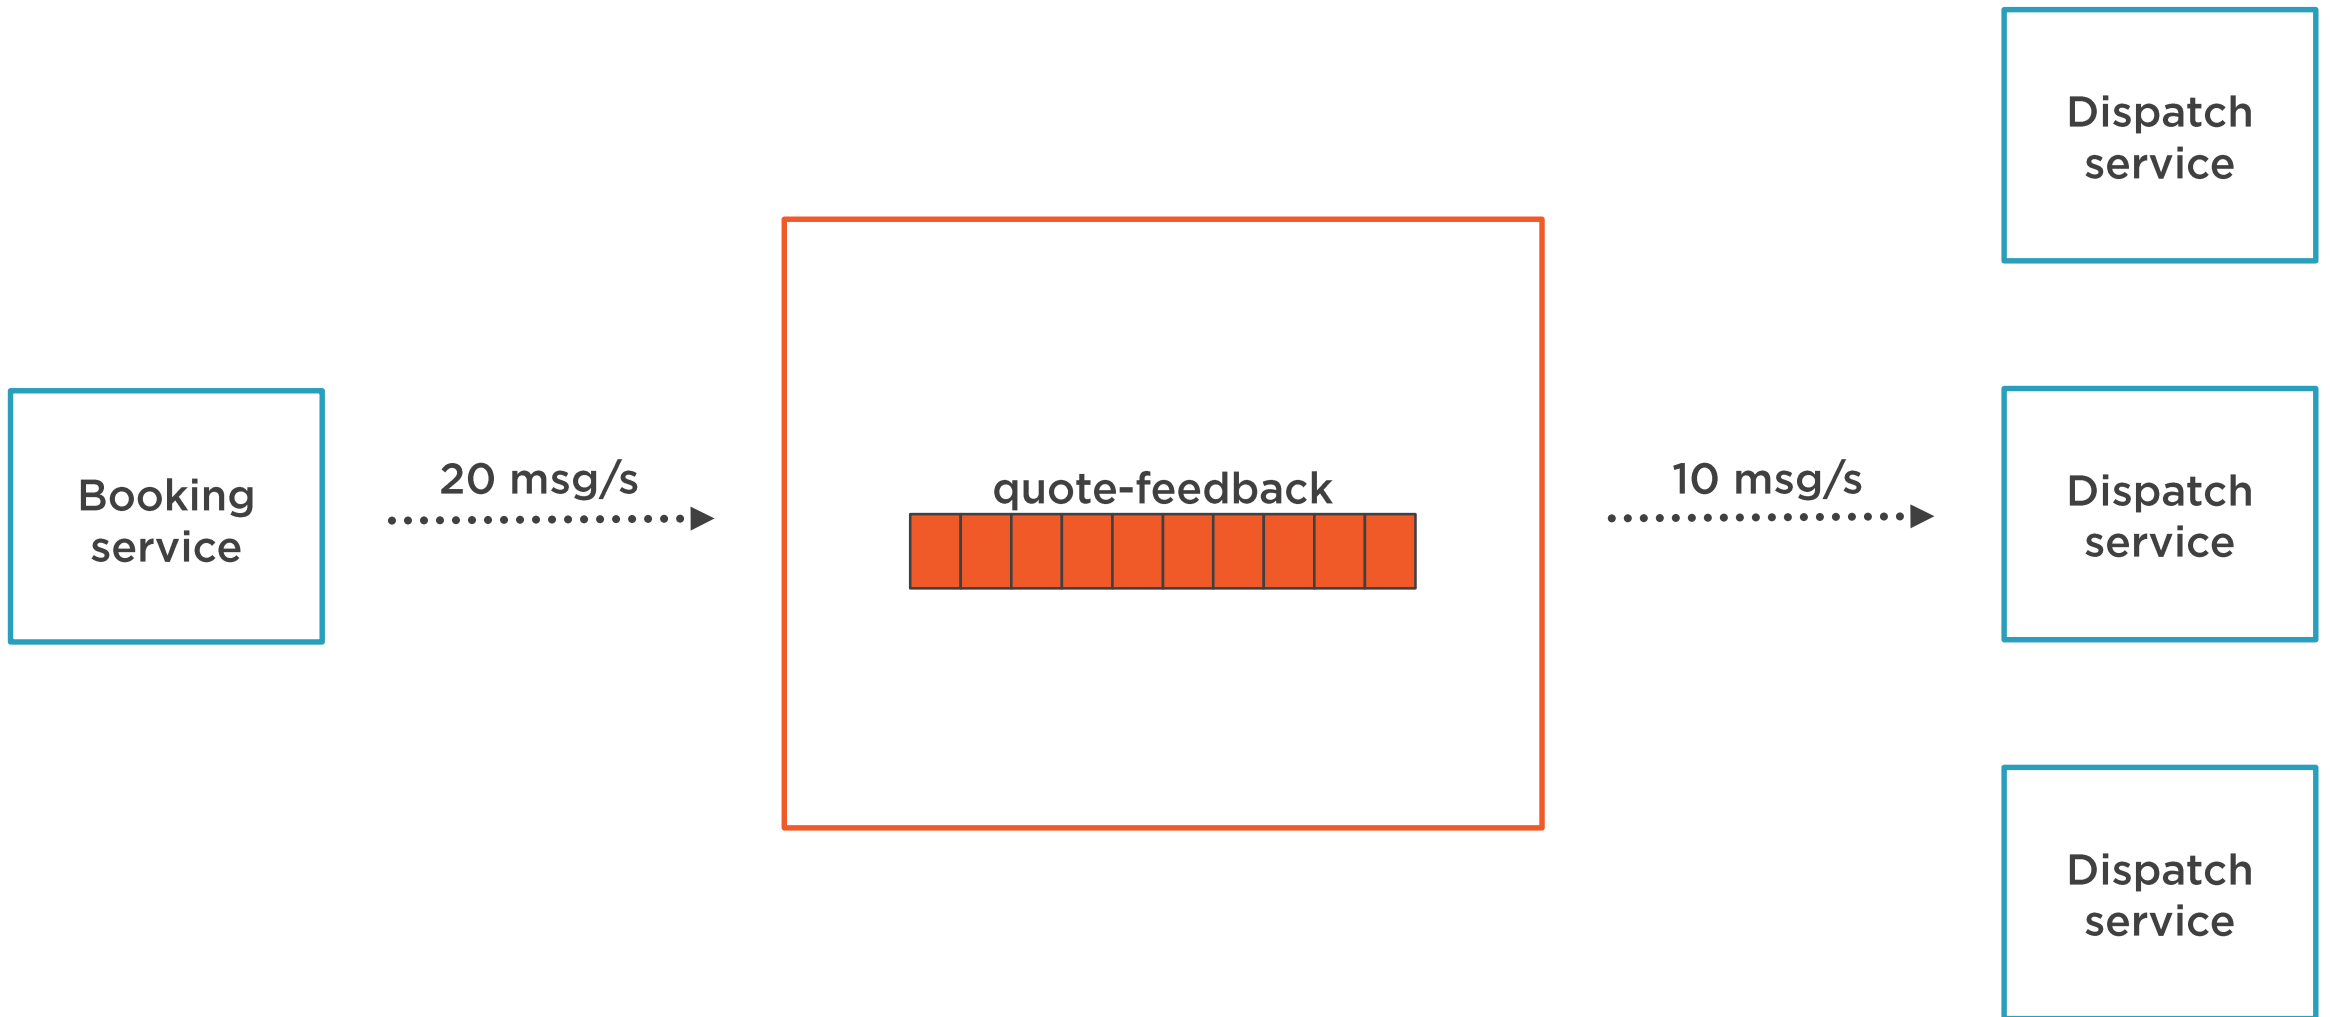


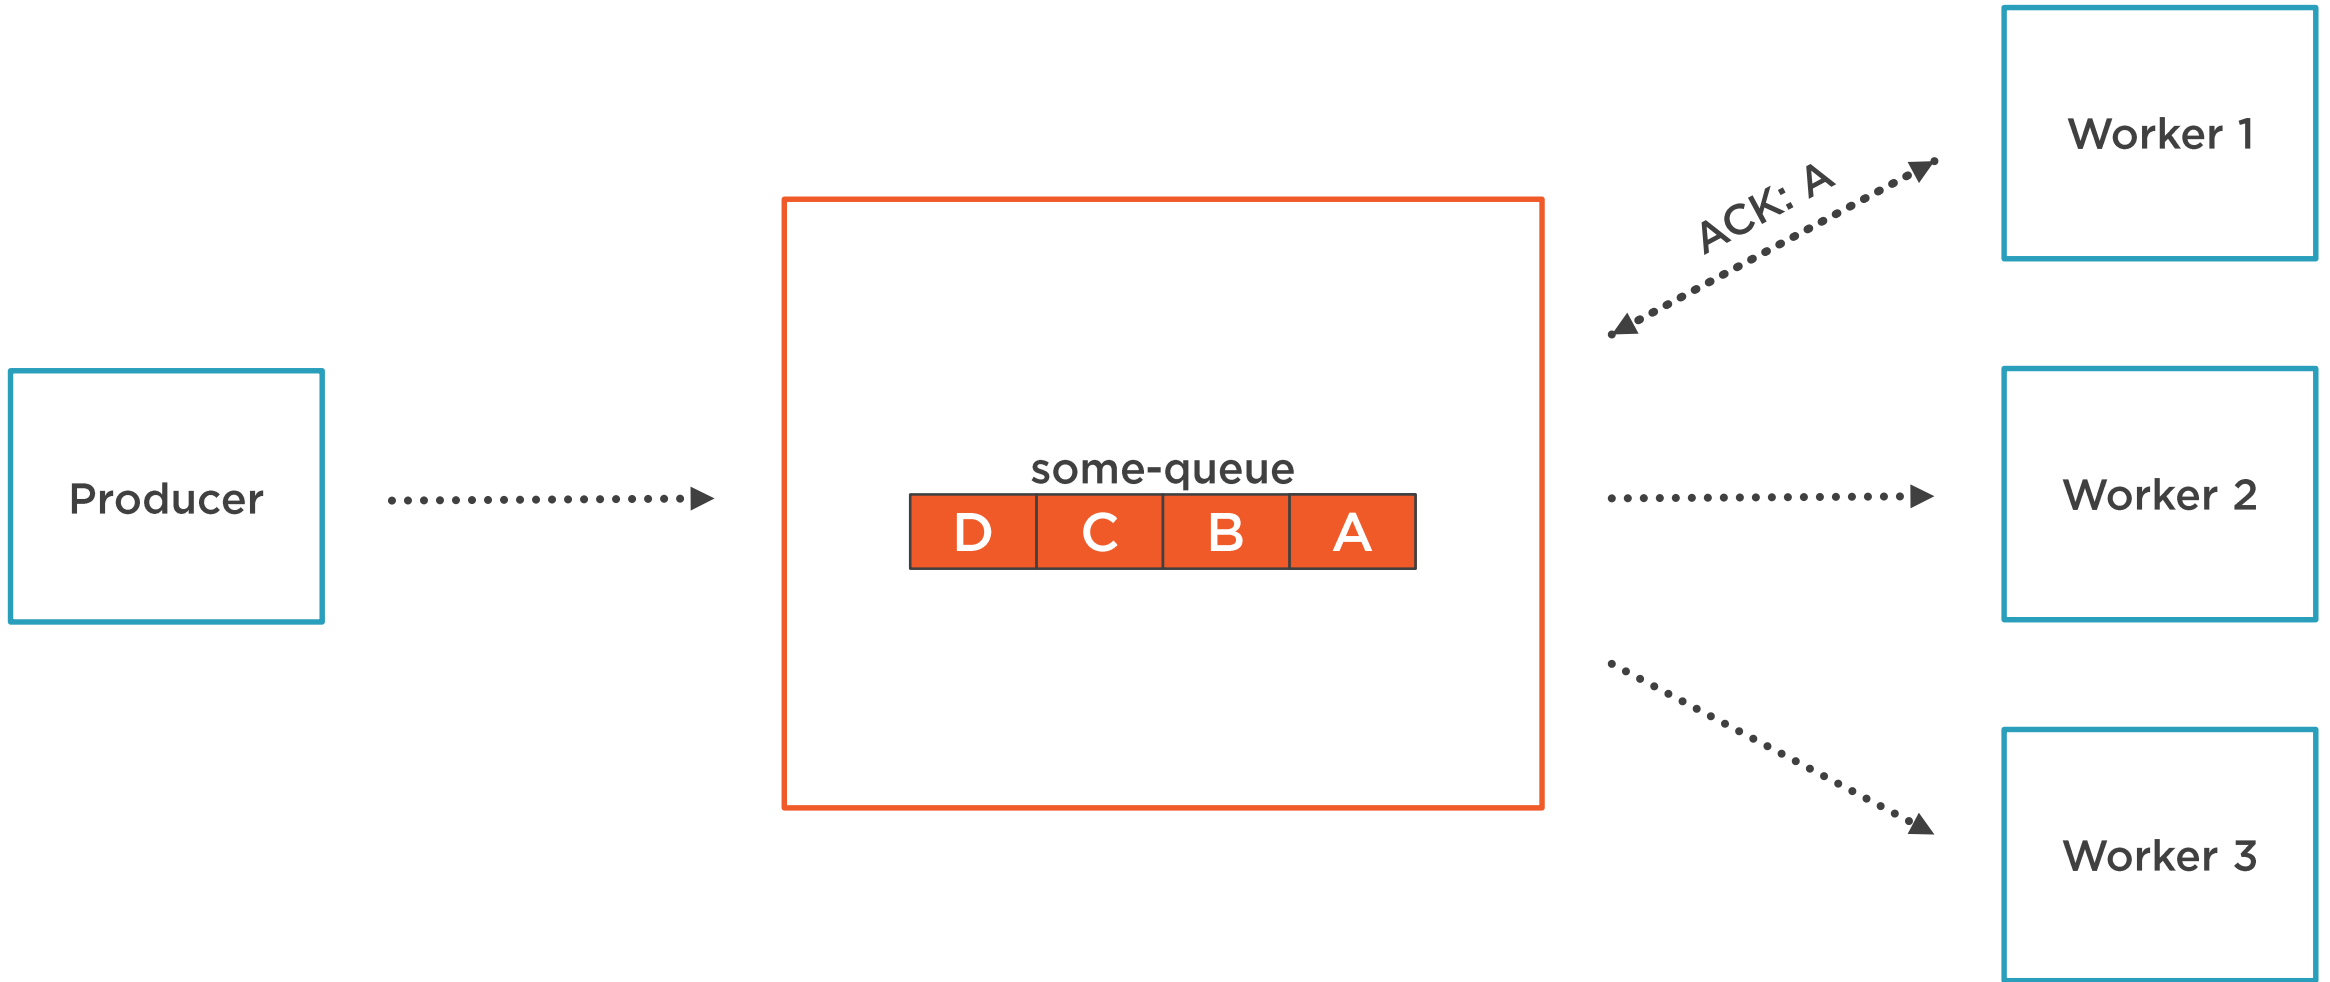


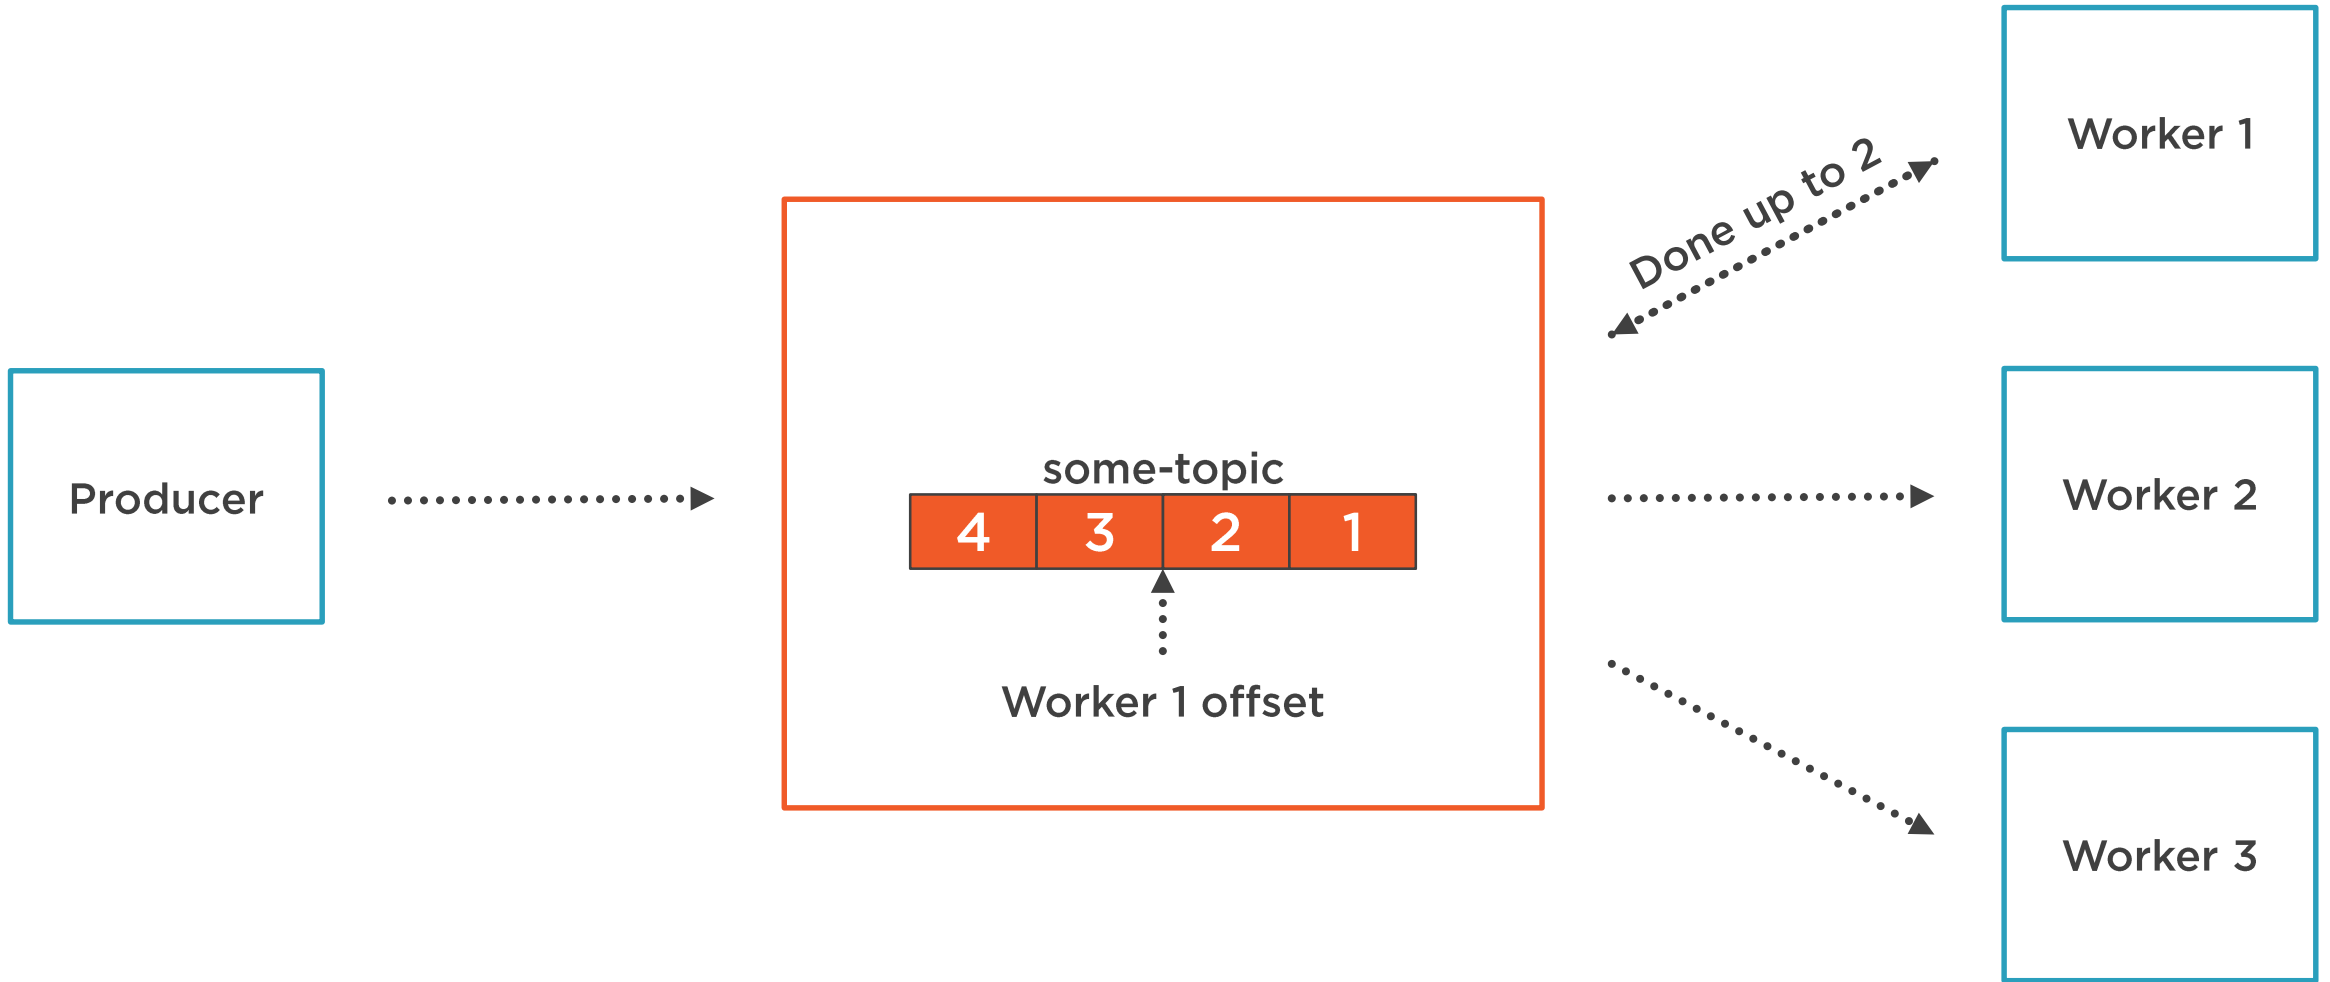




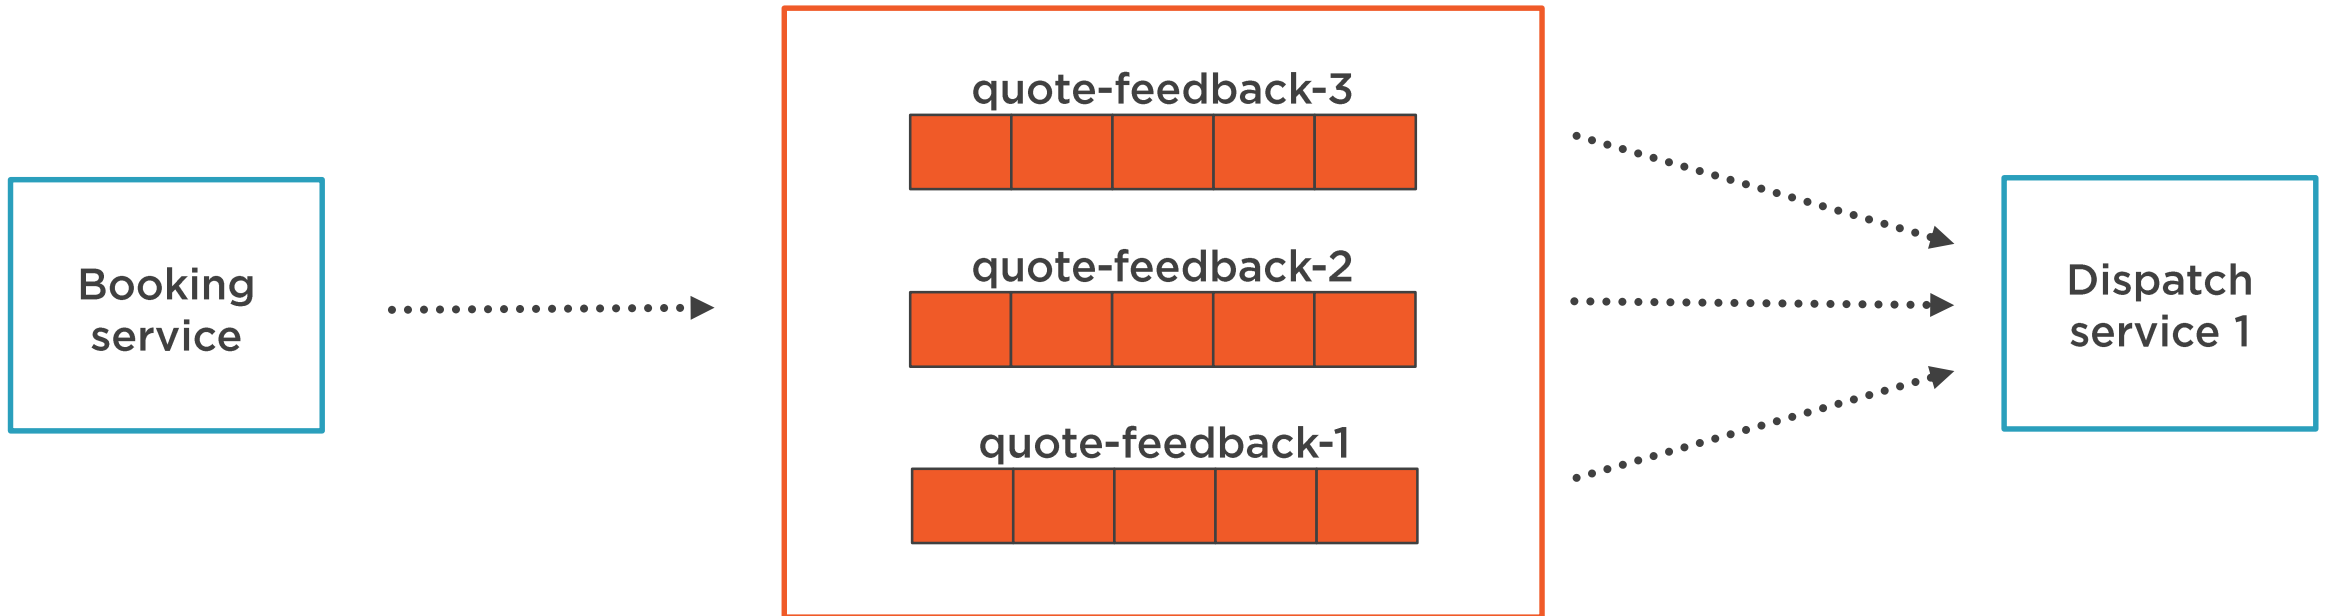








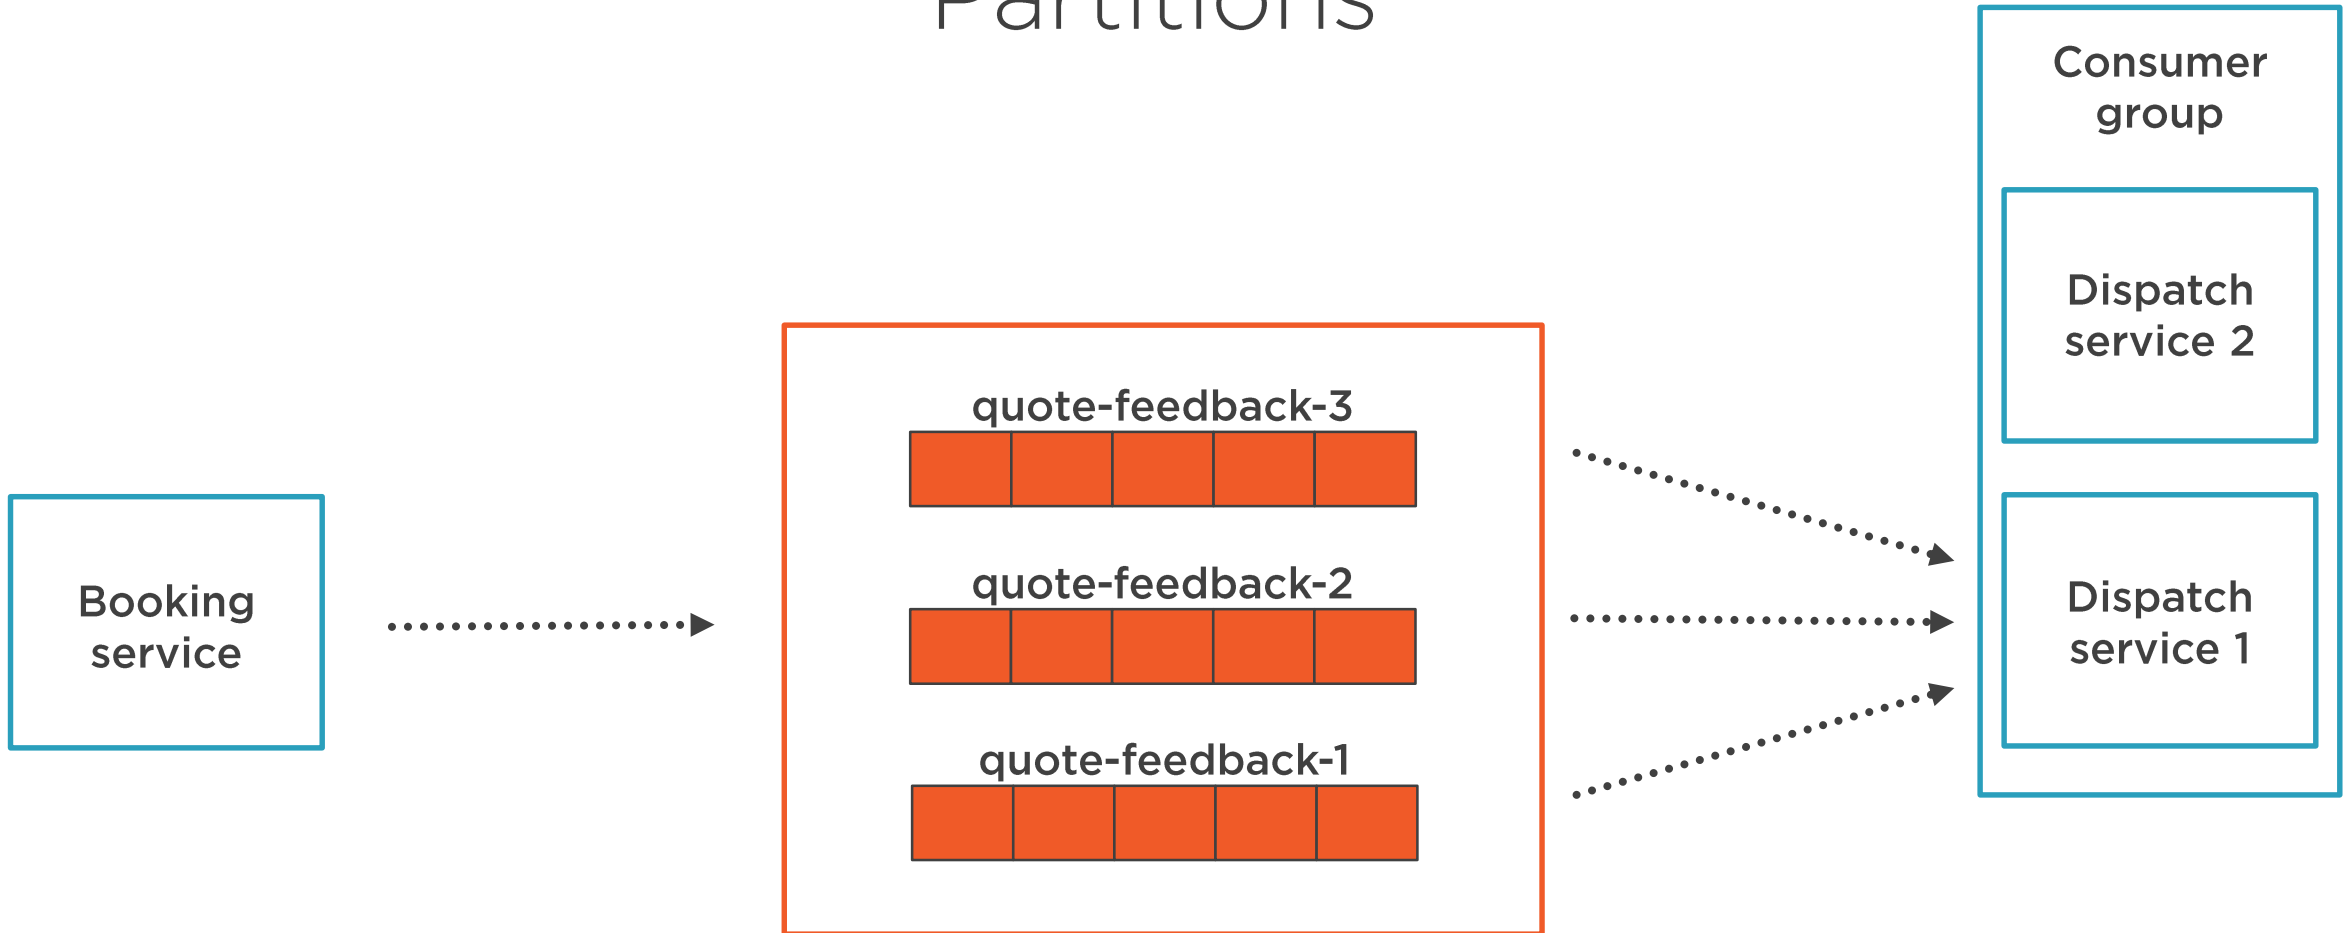
# Partitions



# Partitions

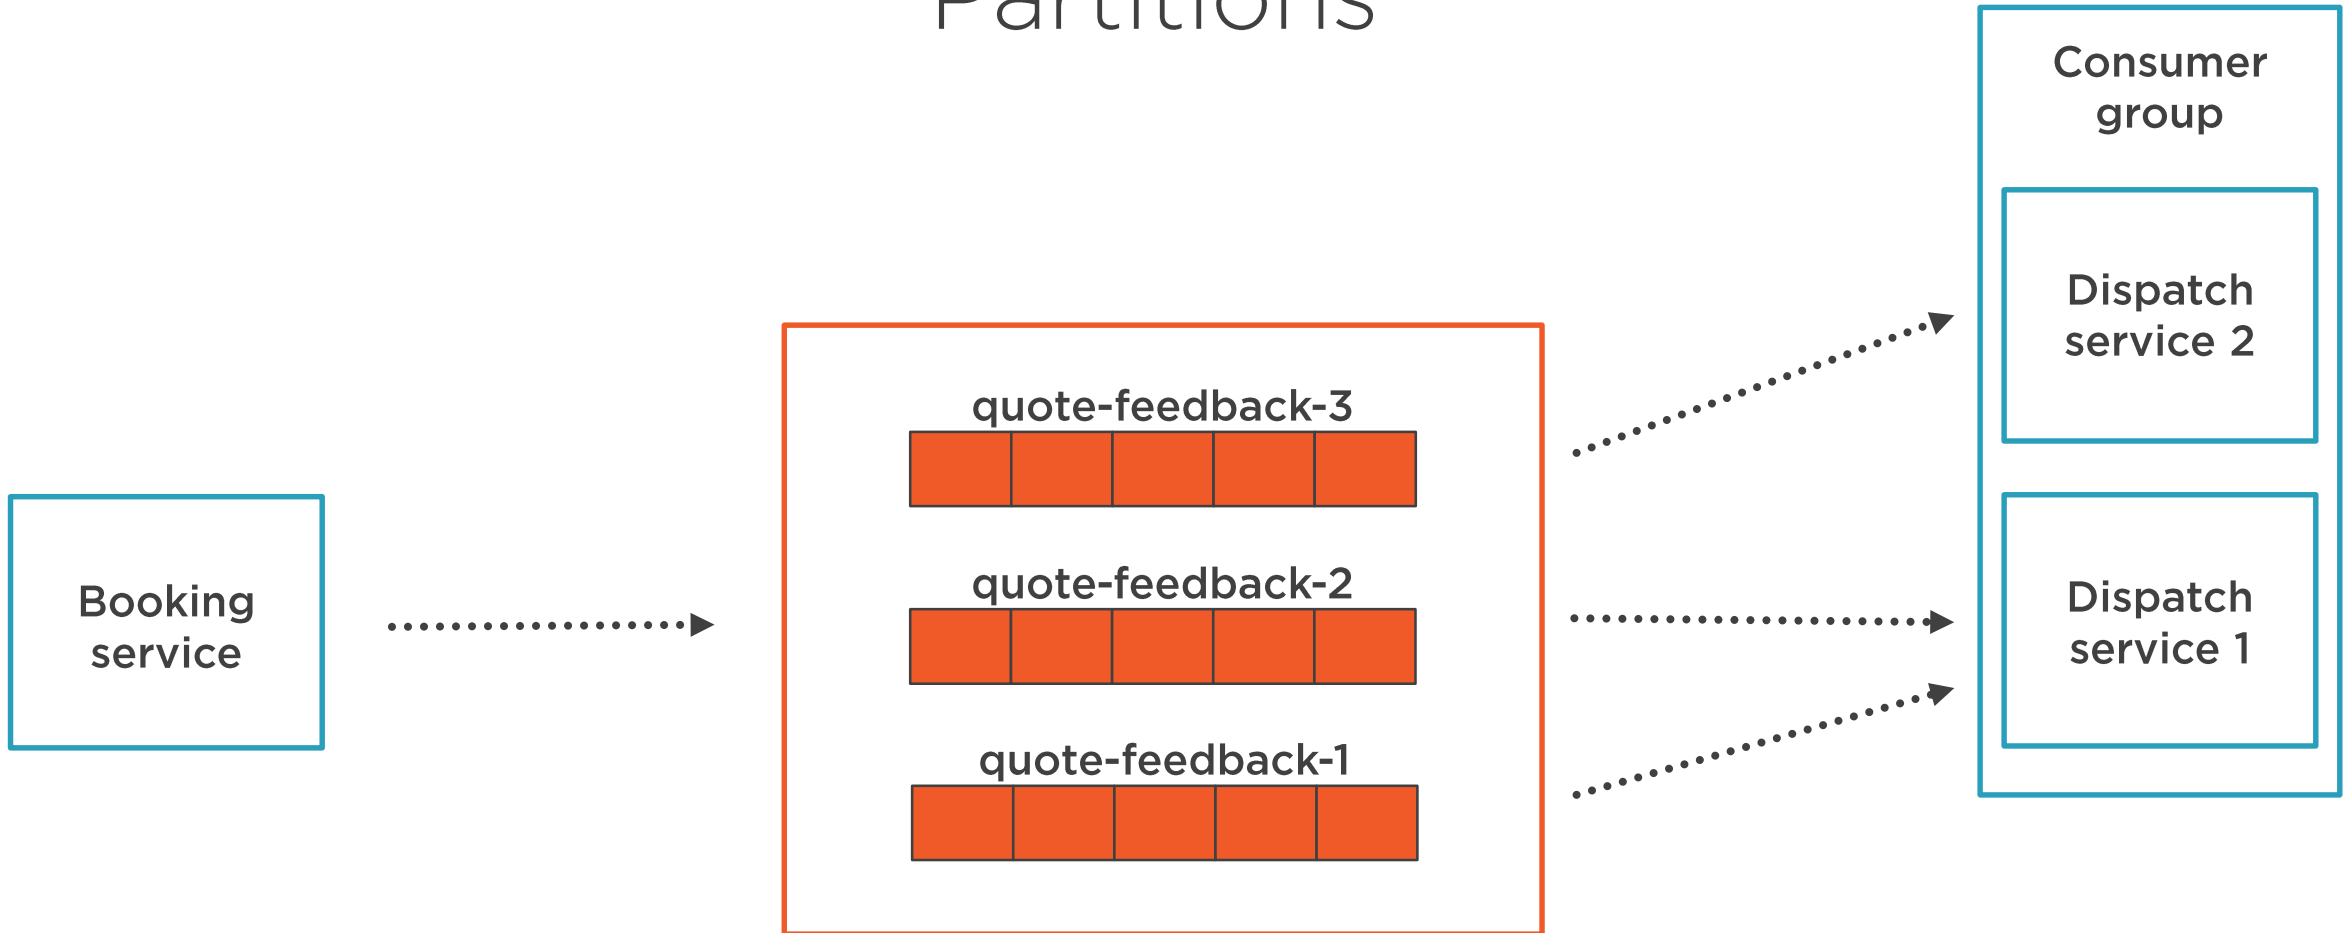


# Partitions

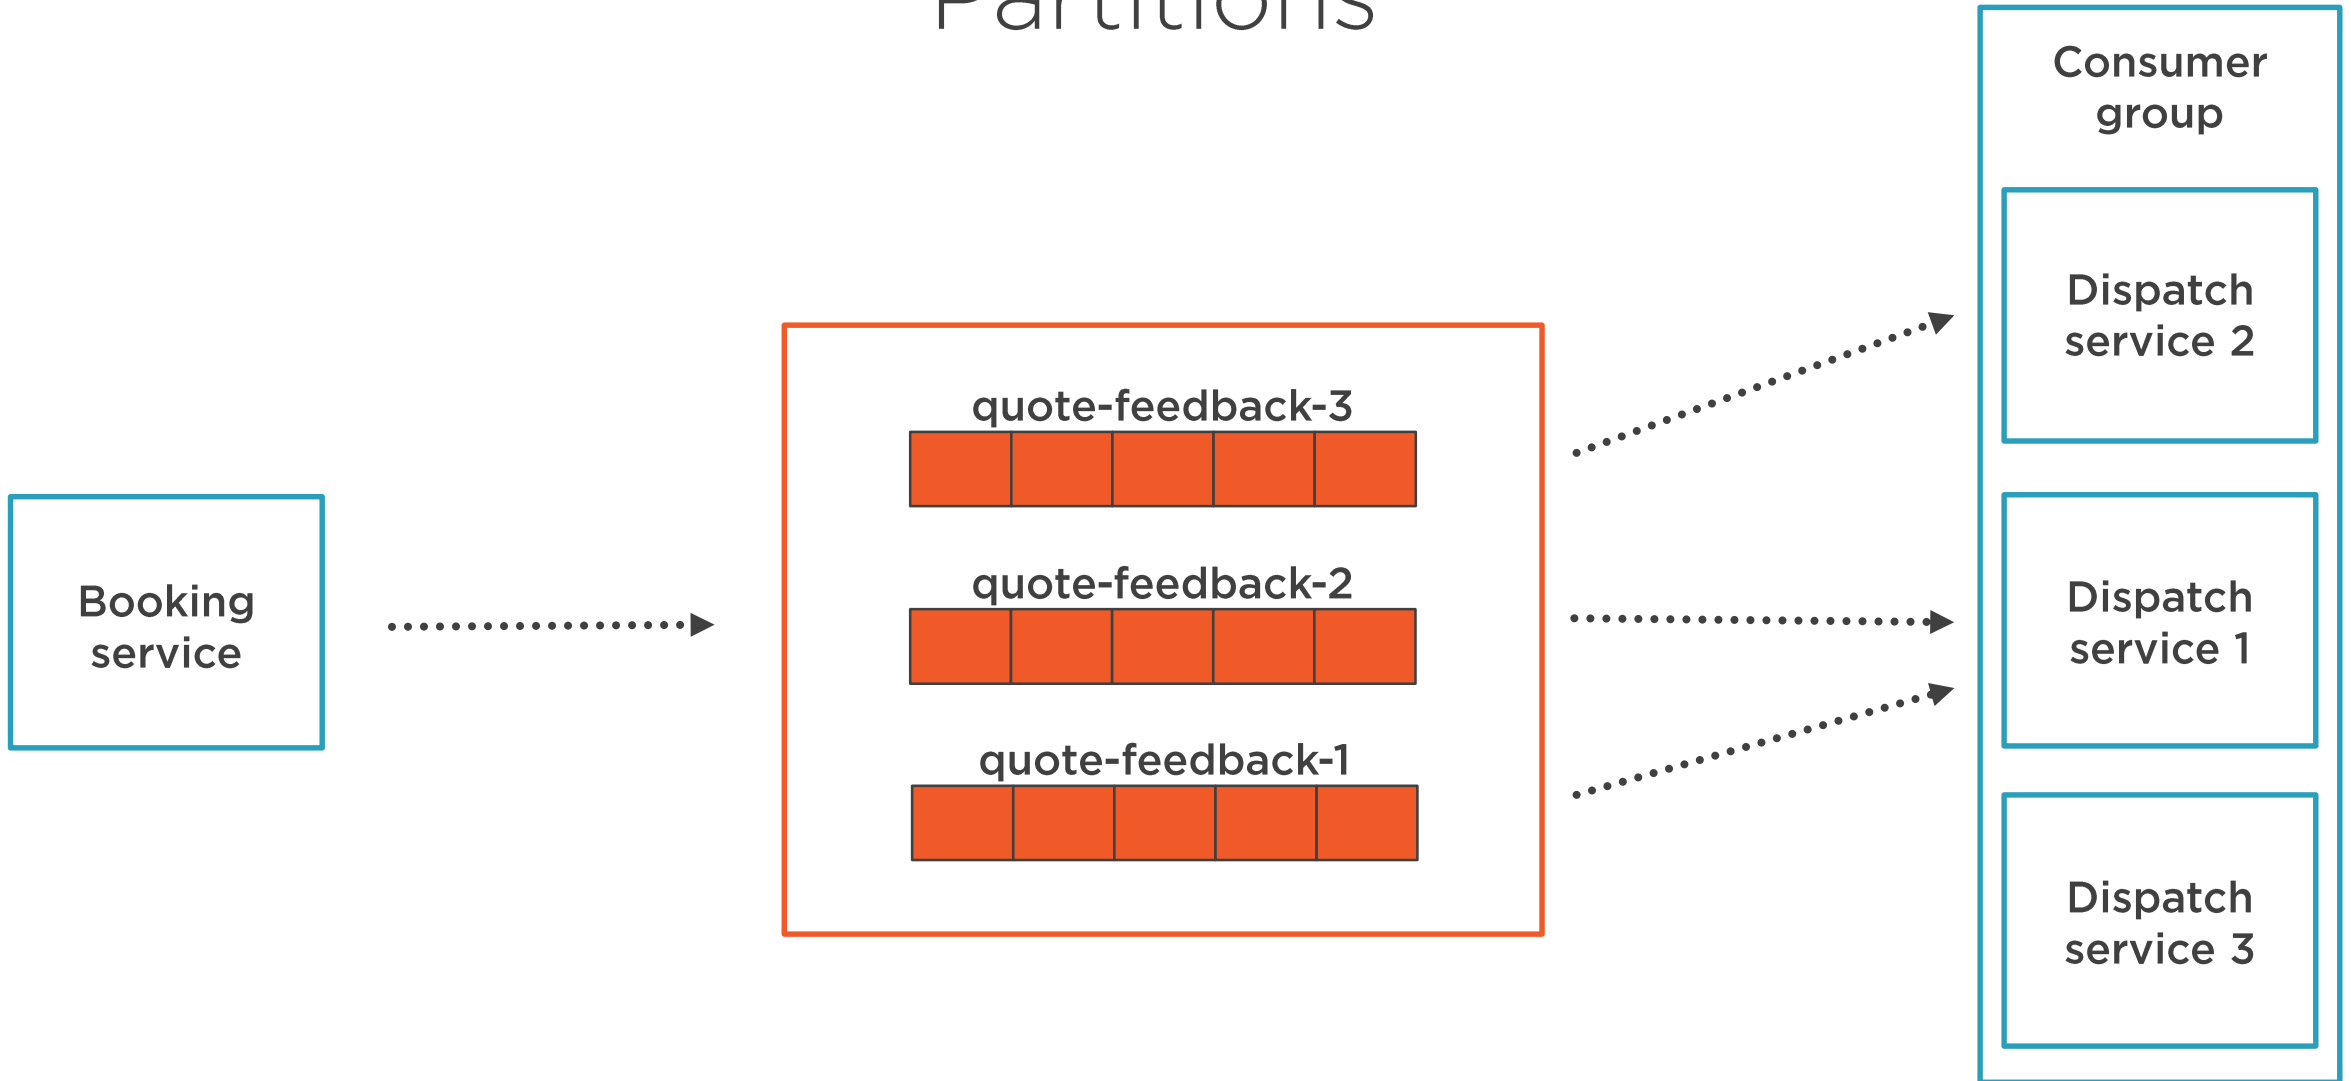




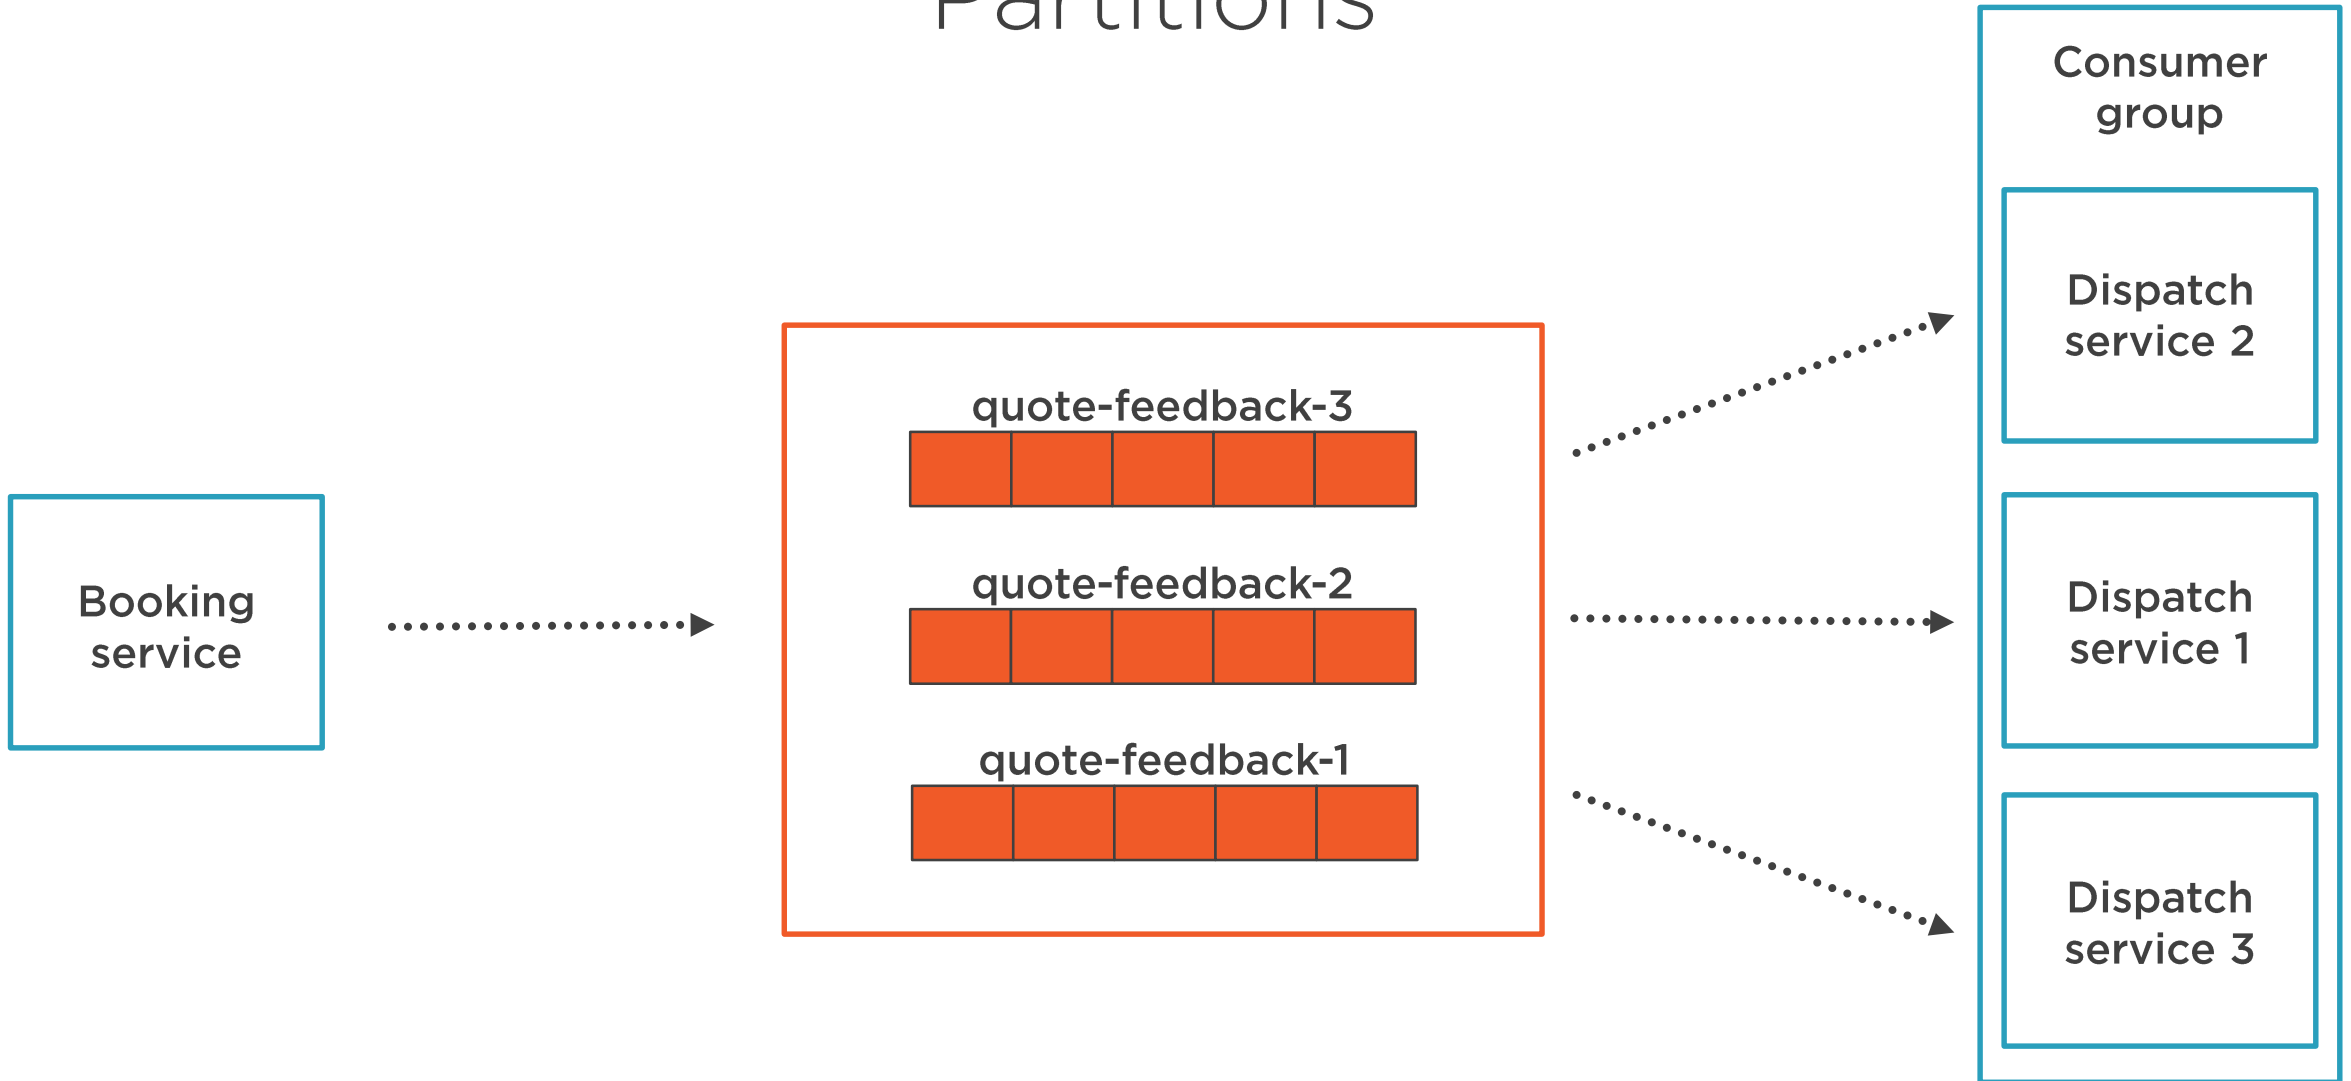
# Partitions



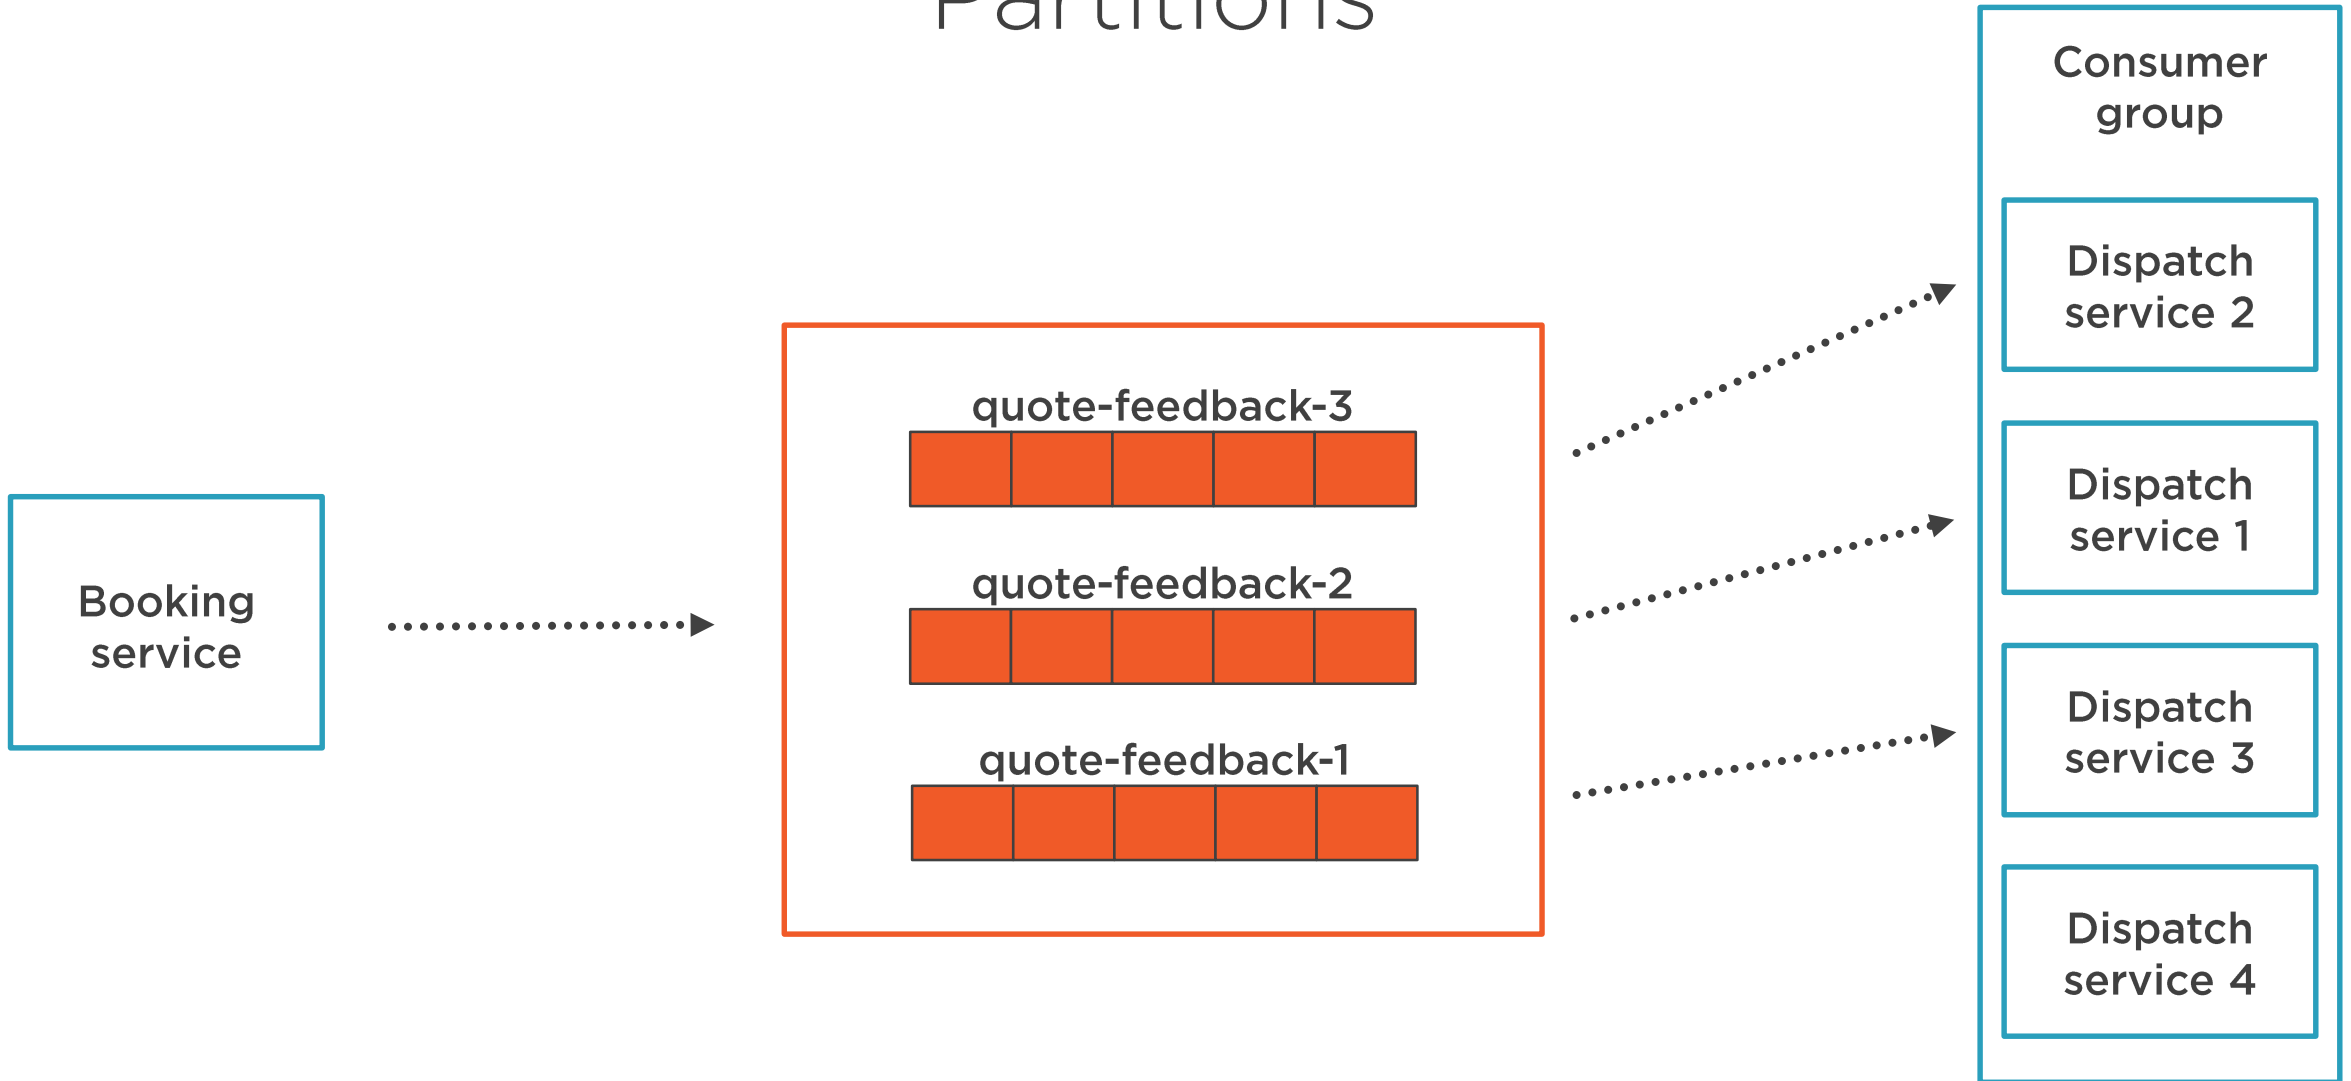
# Partitions



# Partitions



# Partitions



Number of partitions should  
be chosen up front



# End-to-end Demo

---



# Consuming Messages

```
Properties props = new Properties();
props.setProperty("bootstrap.servers", "localhost:9092");
props.setProperty("group.id", "consumer-group-name")
props.setProperty("value.deserializer", ...);

KafkaConsumer<String, String> = new KafkaConsumer<>(props);
consumer.subscribe(List.of("quote-feedback"));

while(true) {
    var records = consumer.poll(...);
    for(var record : records) {
        System.out.println(record.value());
    }
}
```



# Producing Messages

```
Properties props = new Properties();  
props.setProperty("bootstrap.servers", "localhost:9092");  
props.setProperty("value.serializer", ...);  
  
KafkaProducer<String, String> = new KafkaProducer<>(props);  
  
ProducerRecord<String, String> r = new ProducerRecord("topic", "value");  
  
producer.send();
```





# Summary



**Running Kafka**

**Creating topics**

**Understanding partitions**

**Running end-to-end scenario**

