

Producers and Consumers



Paweł Kordek

DATA ENGINEER

@pawel_kordek <https://kordek.github.io>



Some questions







How are partitions chosen?

How are they assigned to consumers?

How are messages not lost?

How do consumers commit offsets?





Booking Service

```
Properties props = new Properties();
props.setProperty("bootstrap.servers", "localhost:9092");
props.setProperty("key.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");
props.setProperty("value.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");

KafkaProducer<String, String> = new KafkaProducer<>(props);

ProducerRecord<String, String> r =
    new ProducerRecord("quote-feedback", "Booking 44321 accepted.");

producer.send(r);
```



Key in Kafka



Optional part of a `ProducerRecord`

Used for partitioning

Choosing a partition based on its value is producer's job



Booking Service

```
Properties props = new Properties();
props.setProperty("bootstrap.servers", "localhost:9092");
props.setProperty("key.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");
props.setProperty("value.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");

KafkaProducer<String, String> = new KafkaProducer<>(props);

ProducerRecord<String, String> r =
    new ProducerRecord("quote-feedback", 1094, "Booking 44321 accepted.");

producer.send(r);
```



Booking Service

```
Properties props = new Properties();
props.setProperty("bootstrap.servers", "localhost:9092");
props.setProperty("key.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");
props.setProperty("value.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");

KafkaProducer<Integer, String> = new KafkaProducer<>(props);

ProducerRecord<Integer, String> r =
    new ProducerRecord("quote-feedback", 1094, "Booking 44321 accepted.");

producer.send(r);
```



Booking Service

```
Properties props = new Properties();
props.setProperty("bootstrap.servers", "localhost:9092");
props.setProperty("key.serializer", IntegerSerializer.class.getName());
props.setProperty("value.serializer",
    "org.apache.kafka.common.serialization.StringSerializer");

KafkaProducer<Integer, String> = new KafkaProducer<>(props);

ProducerRecord<Integer, String> r =
    new ProducerRecord("quote-feedback", 1094, "Booking 44321 accepted.");

producer.send(r);
```



Booking Service

```
Properties props = new Properties();
props.setProperty("bootstrap.servers", "localhost:9092");
props.setProperty("key.serializer", IntegerSerializer.class.getName());
props.setProperty("value.serializer", StringSerializer.class.getName());

KafkaProducer<Integer, String> = new KafkaProducer<>(props);

ProducerRecord<Integer, String> r =
    new ProducerRecord("quote-feedback", 1094, "Booking 44321 accepted.");

producer.send(r);
```



Demo



Verify how producer assigns partitions

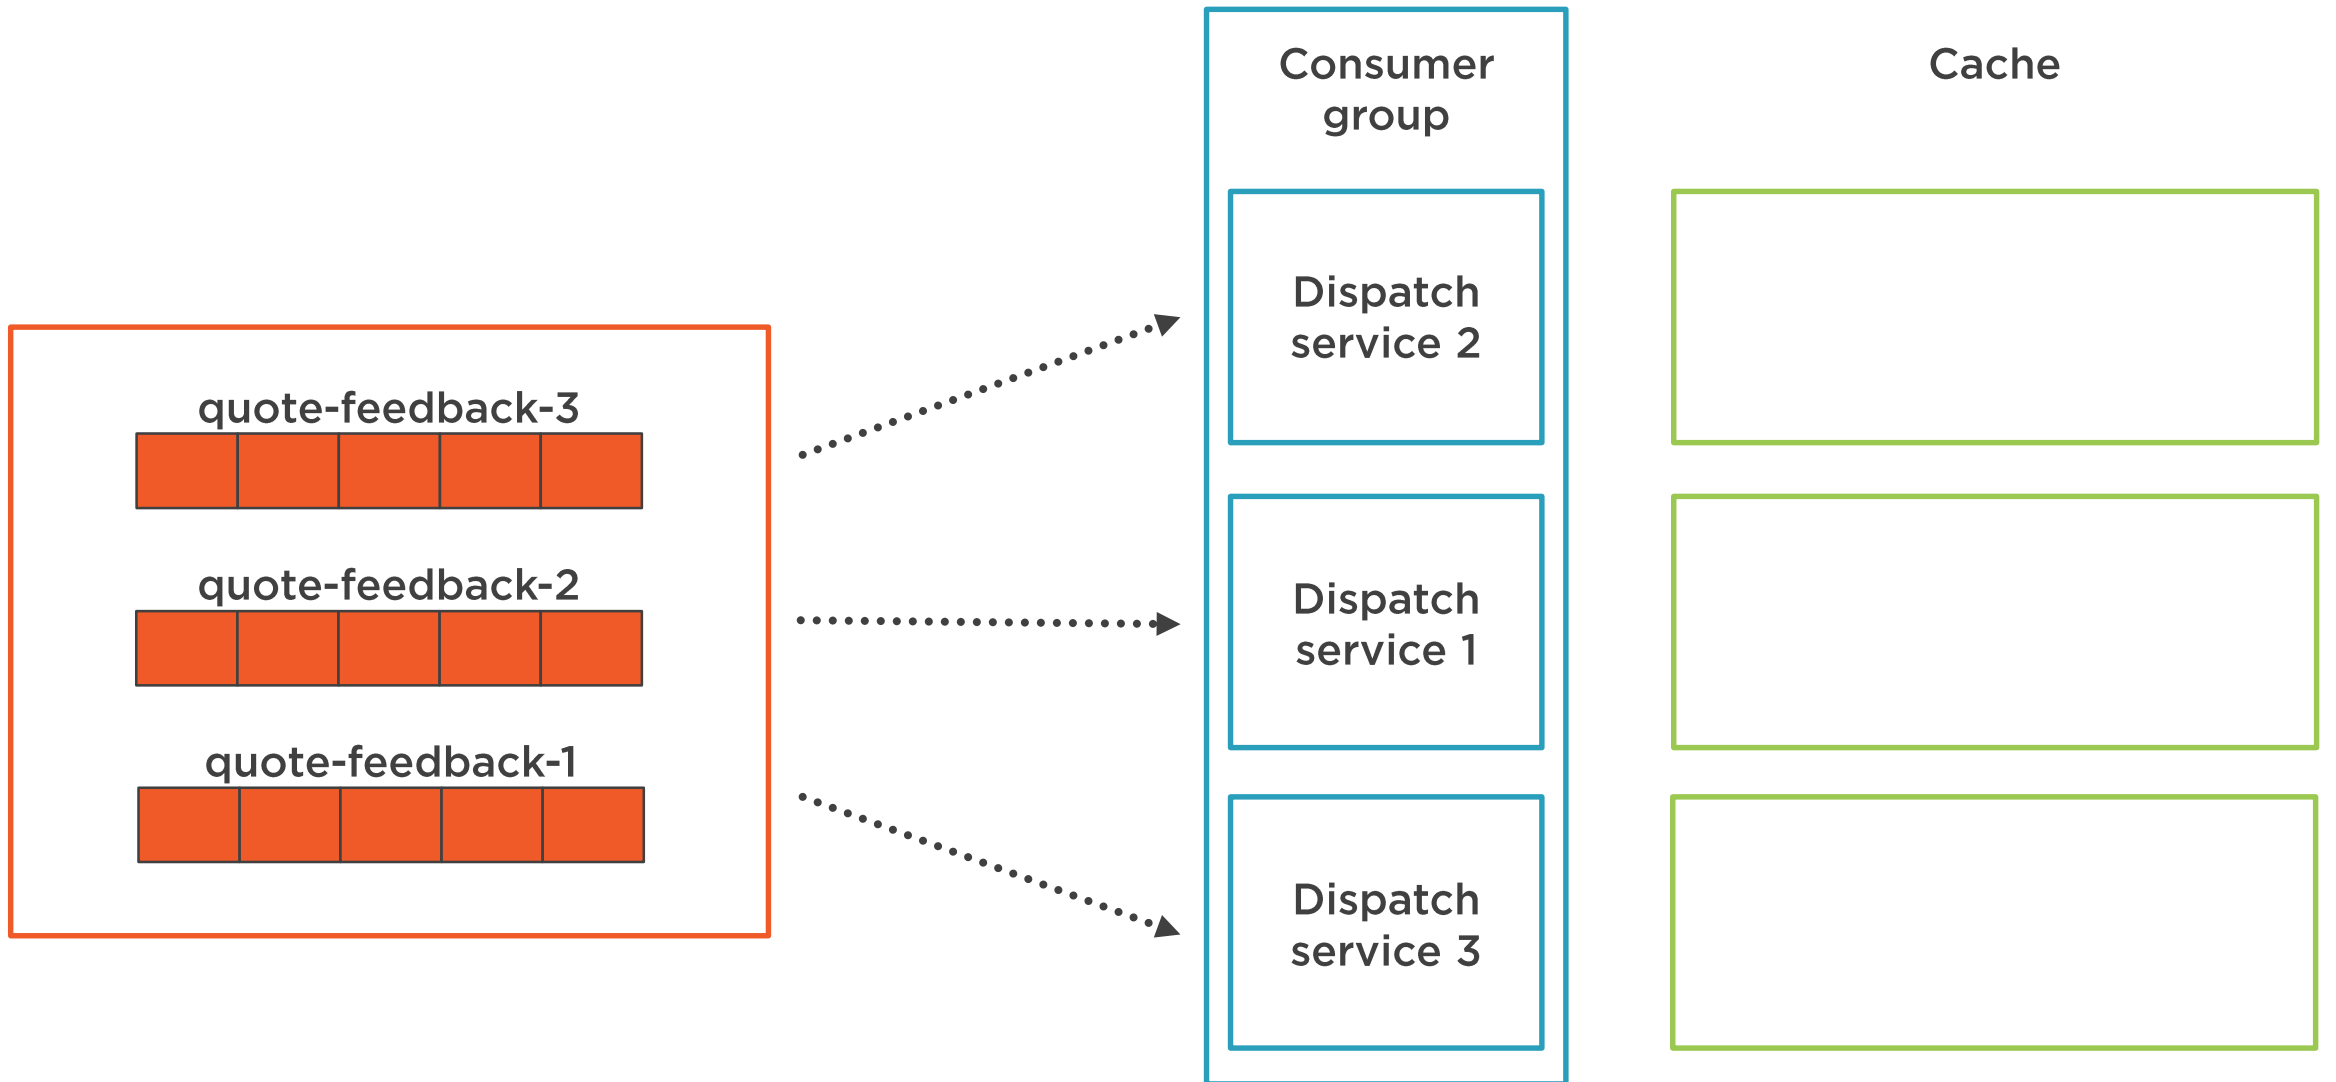


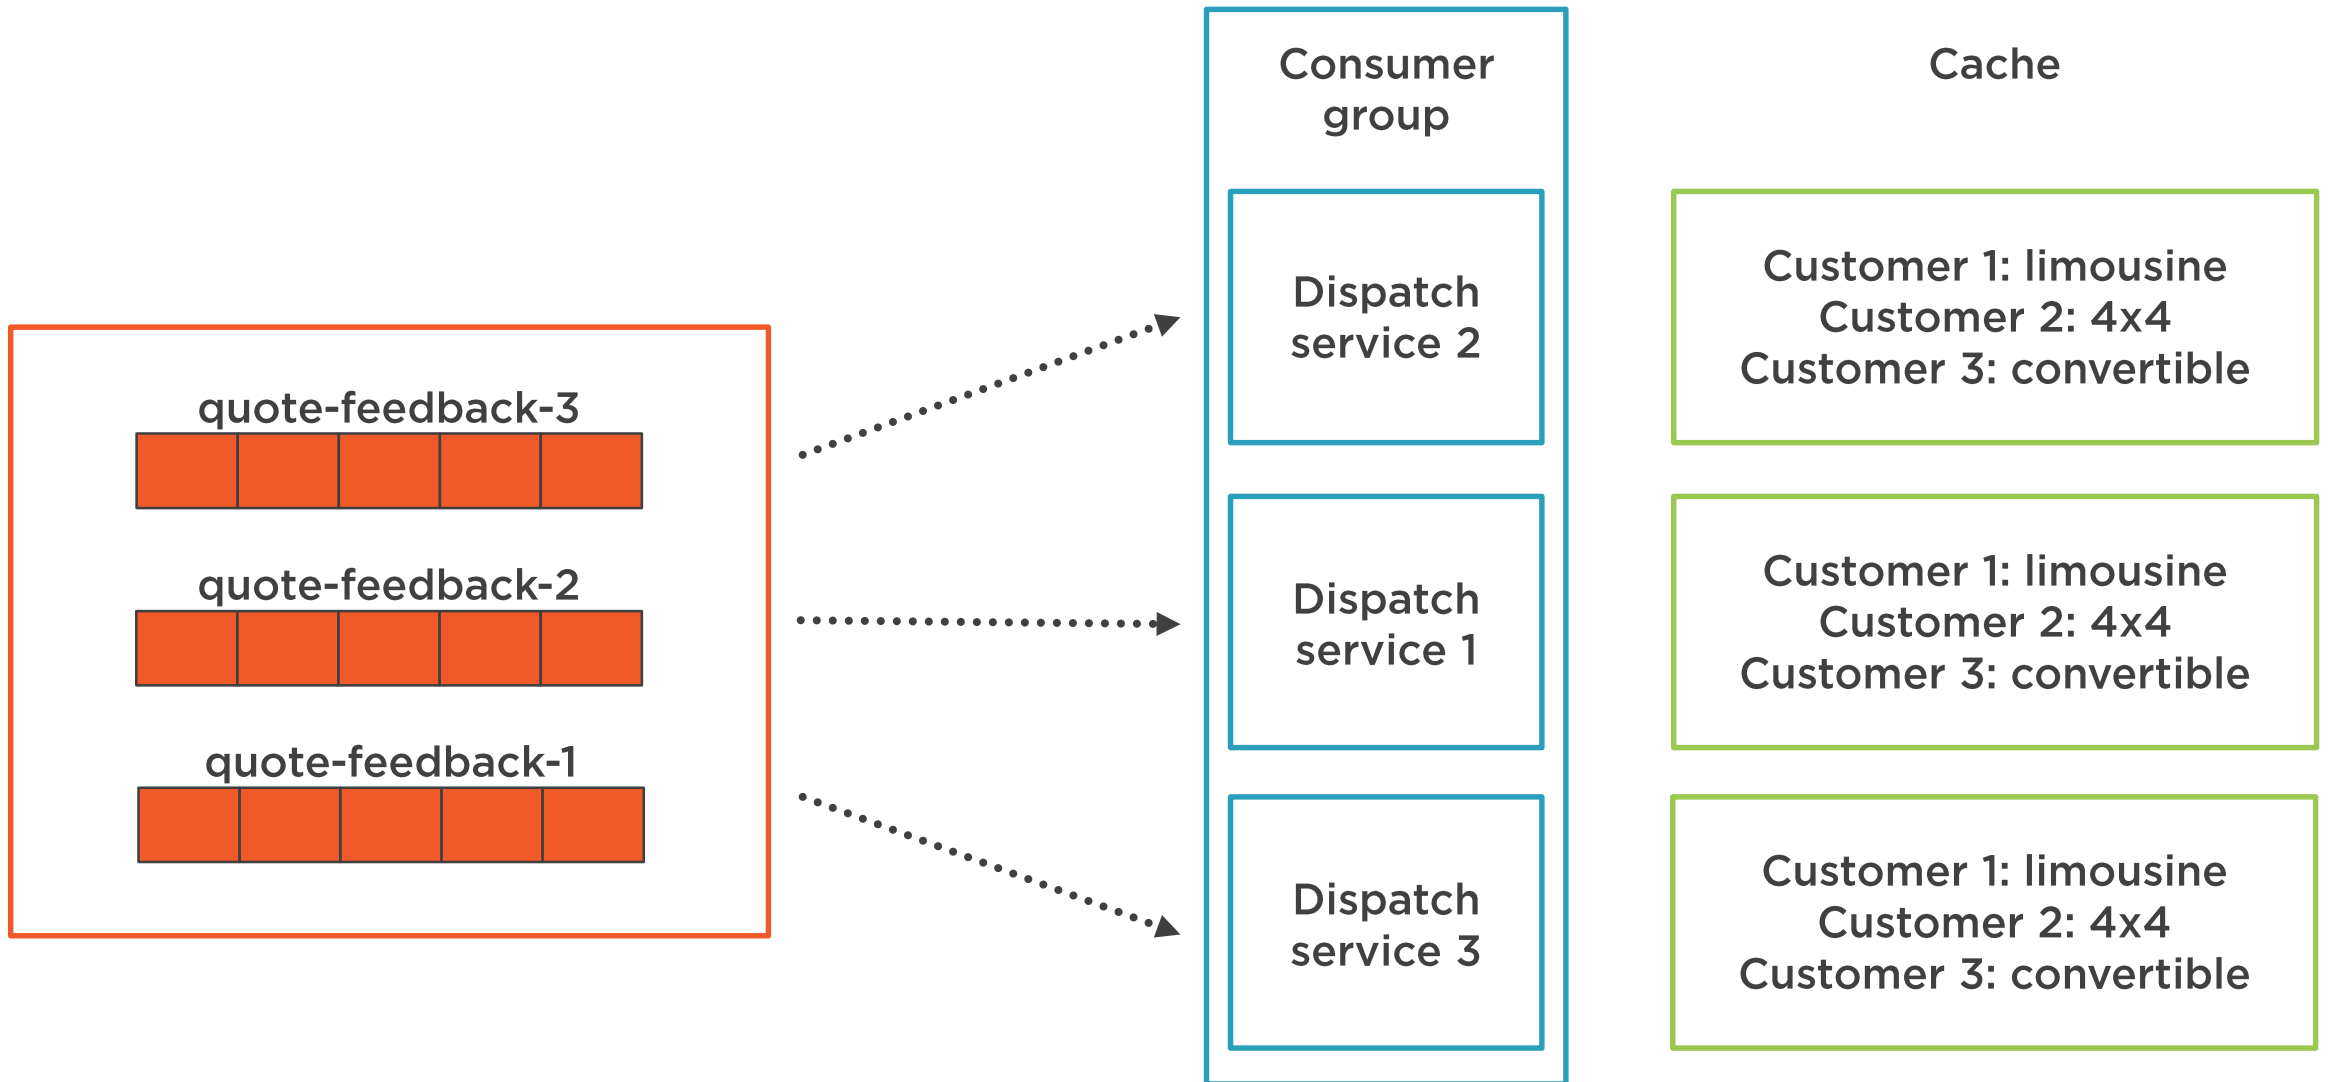
Choosing the Right Key

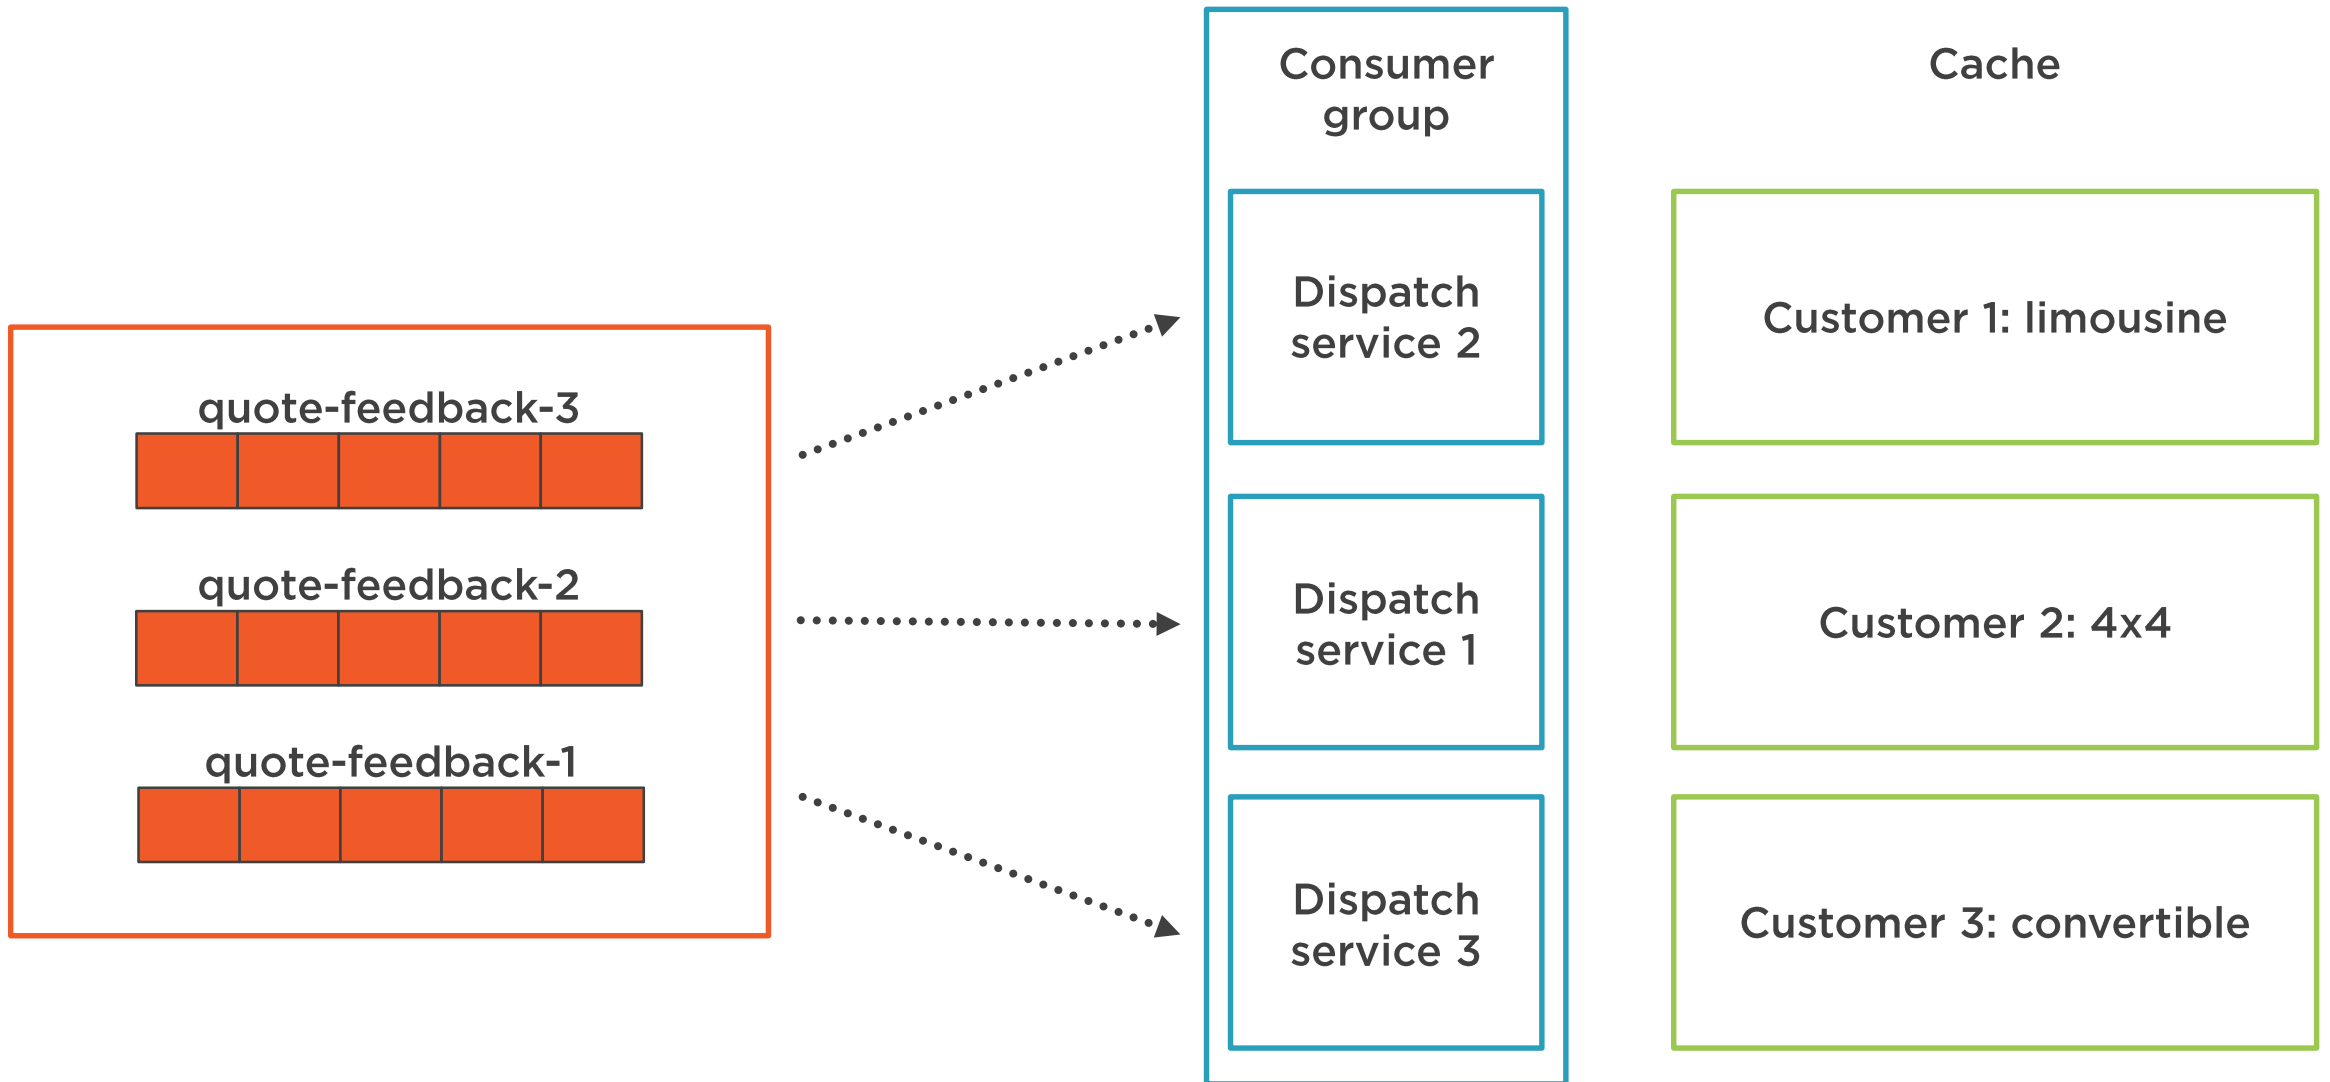


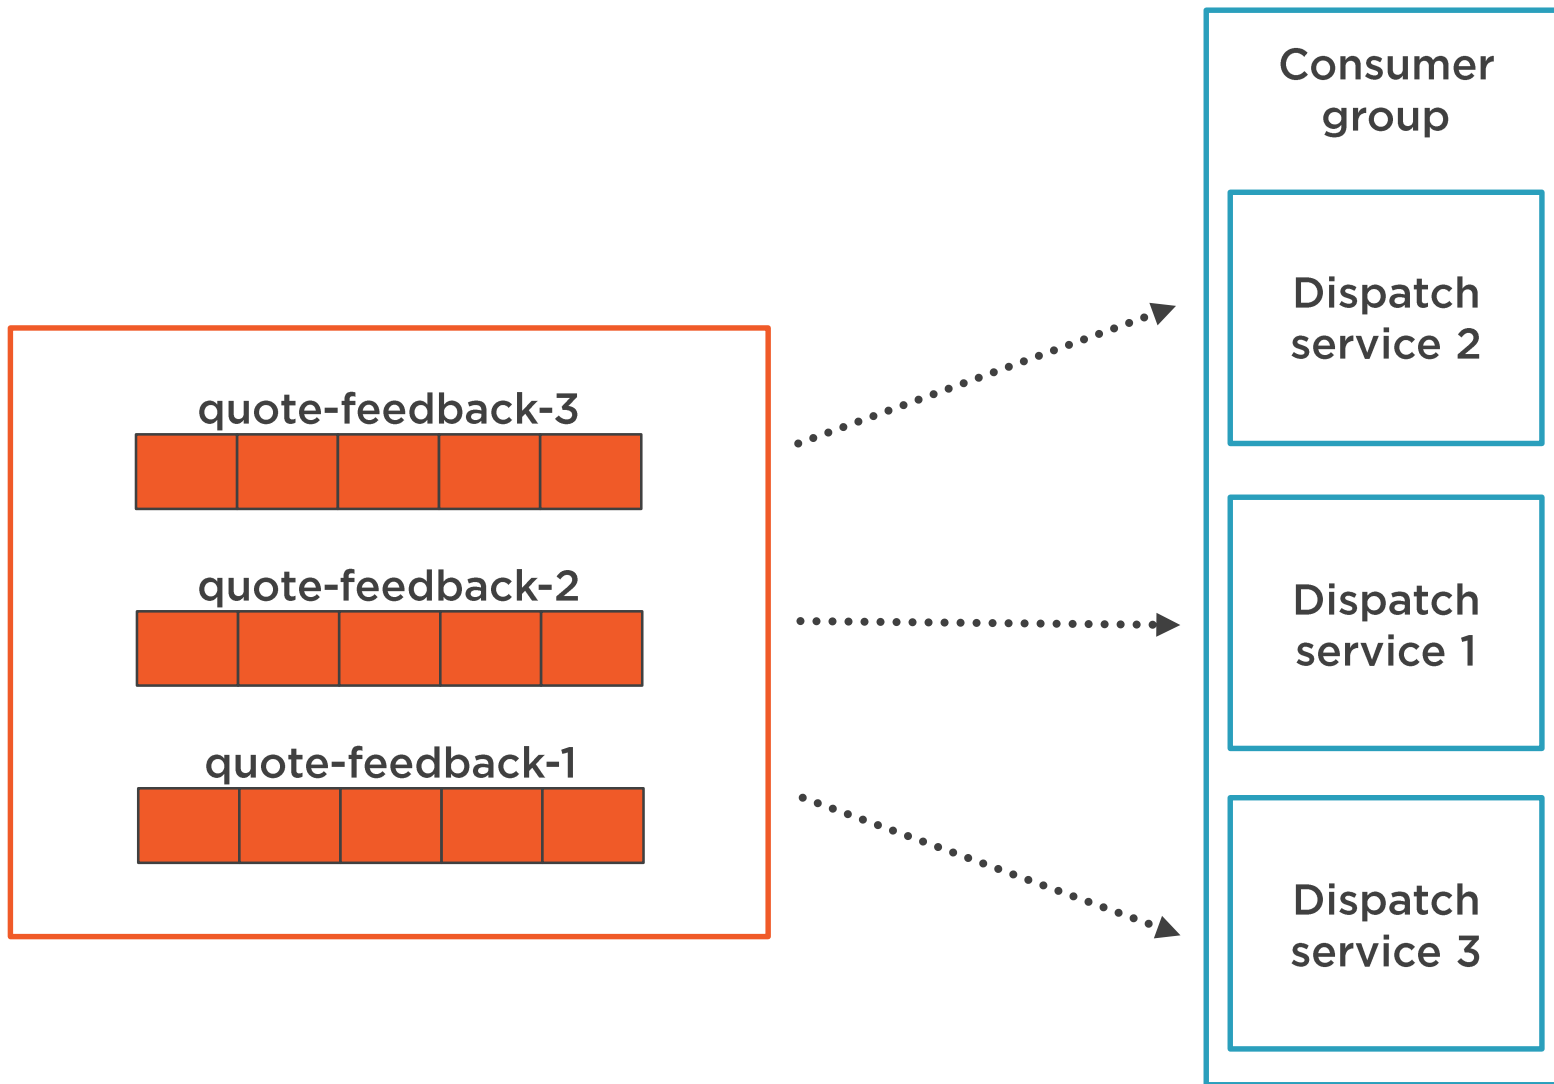
Think about your
consumption patterns

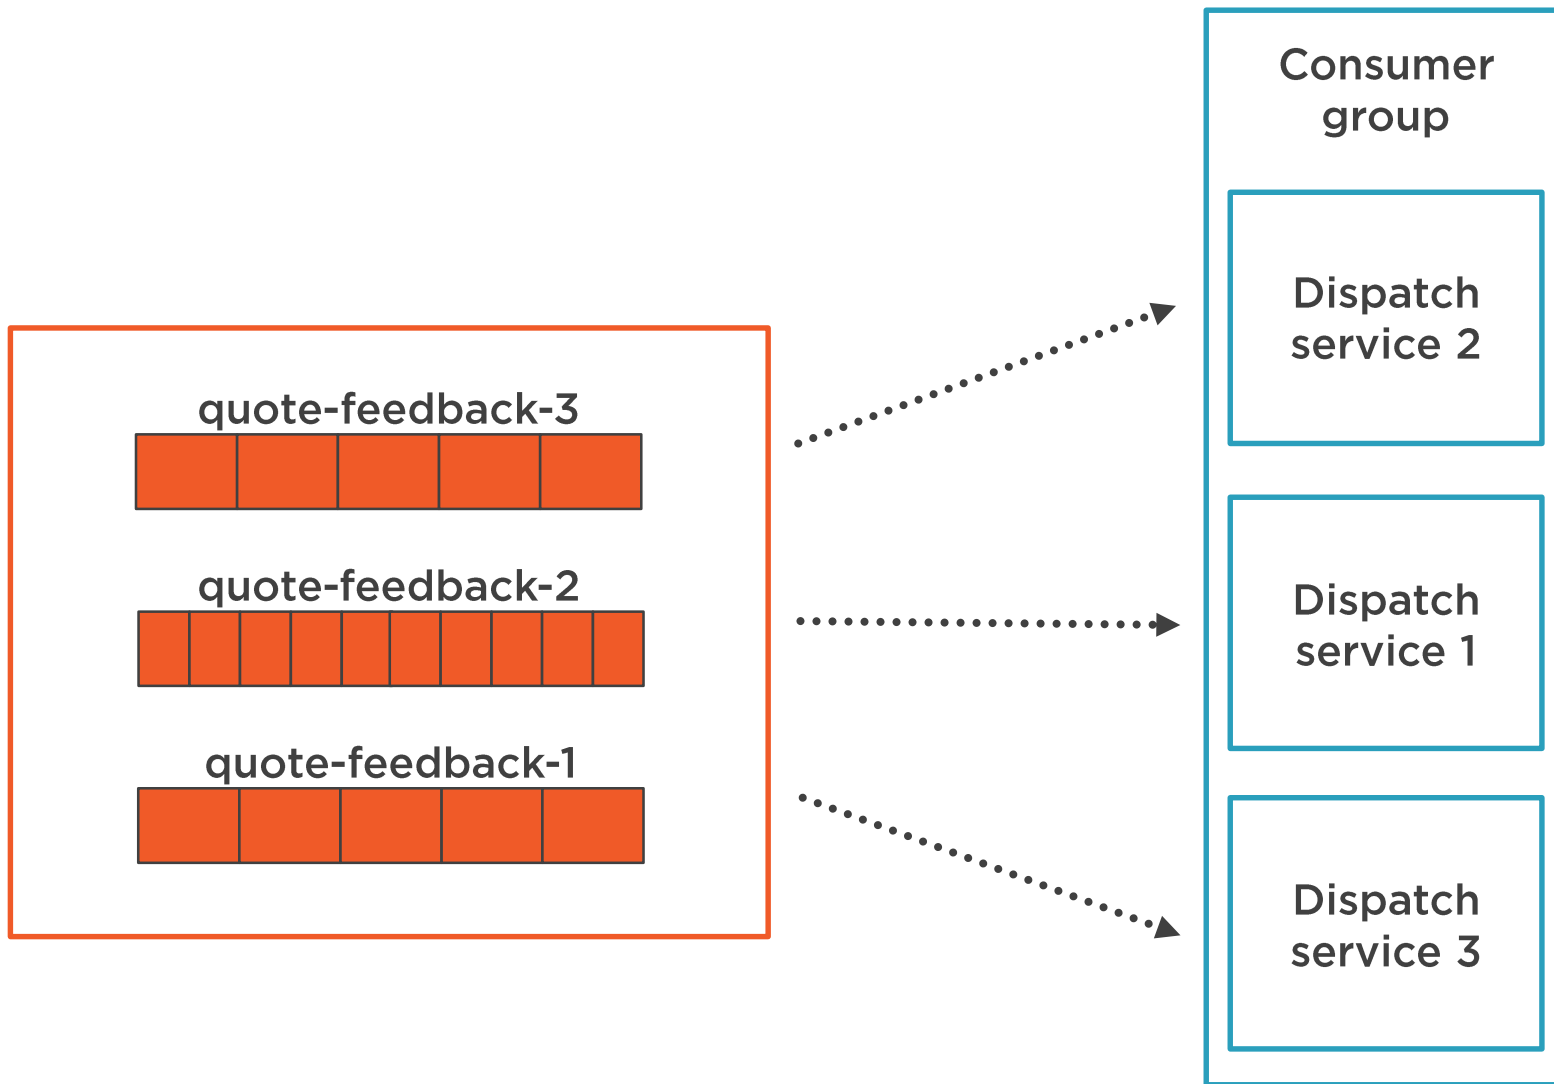


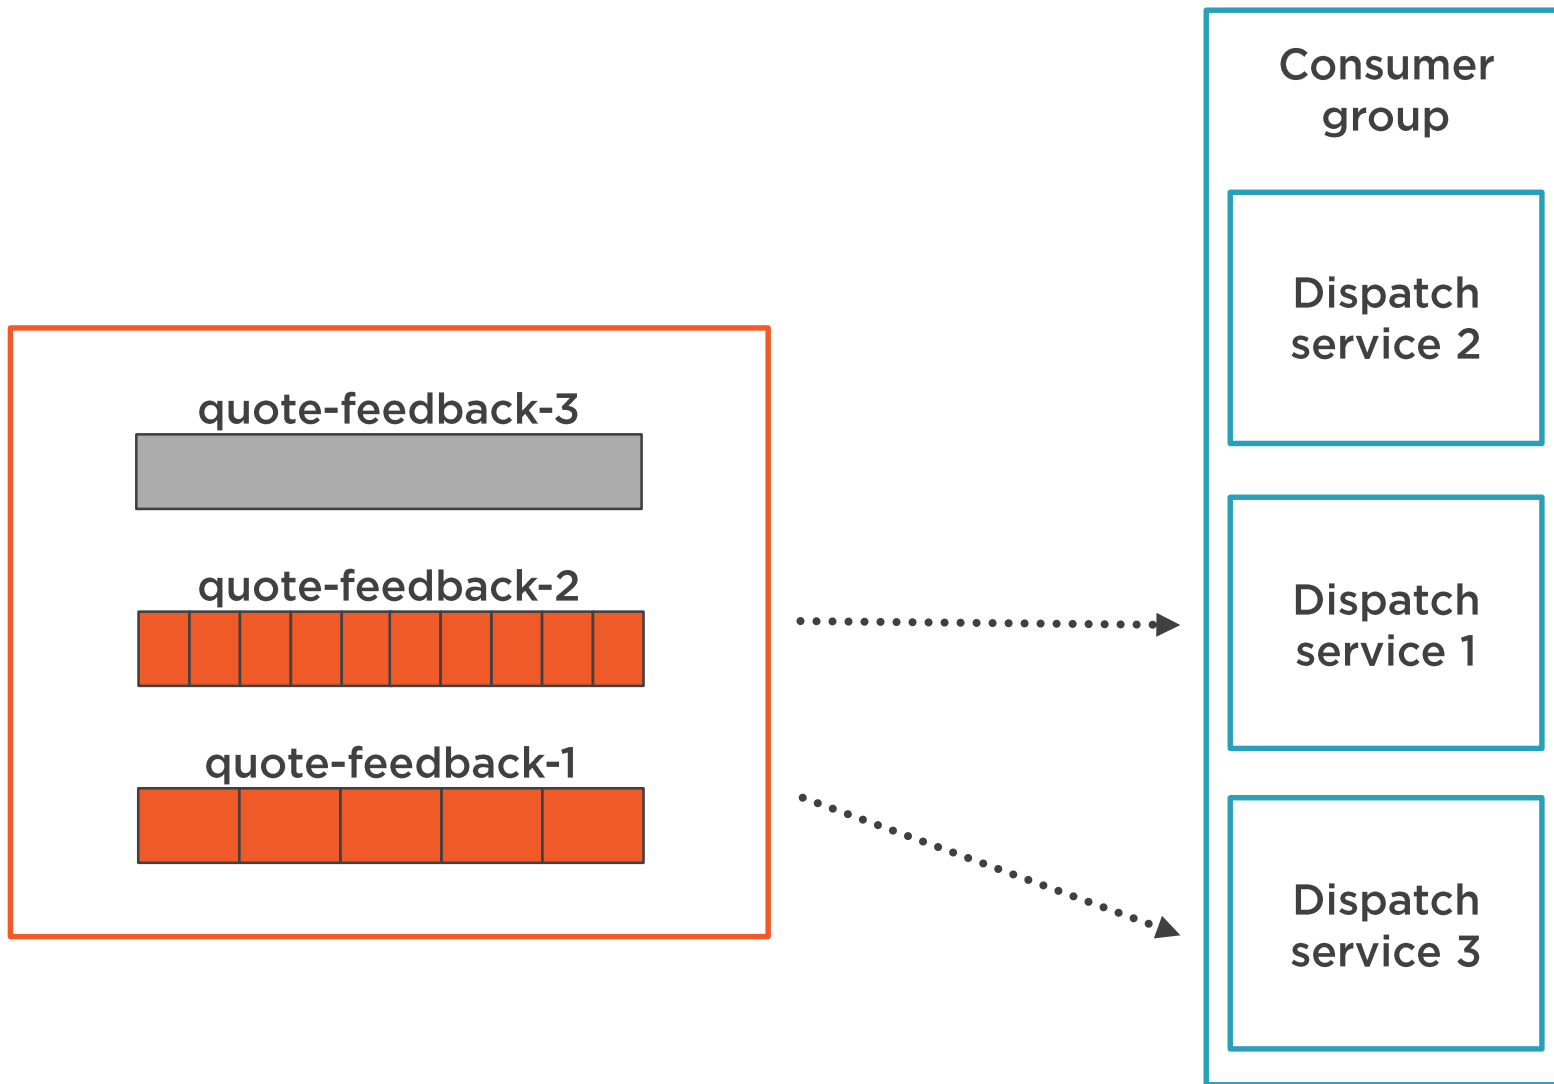














Order of messages can be guaranteed only within a single partition

Custom partitioning is possible, but should be reserved for special cases



Automatic Retries





Maximum number limited by the 'retries' property

Timeout specified by 'delivery.timeout.ms'



Batching and Compression



Booking Service

```
ProducerRecord<Integer, String> r =  
    new ProducerRecord("quote-feedback", 1094, "Booking 44321 accepted.");  
  
producer.send(r);  
  
producer.flush();
```




Producer Flow



Customize using:

- 'batch.size'
- 'linger.ms' (by default 0)

Kafka
Broker



Send buffer (per partition)

101...	111100...
100...	110100...
111...	101101...
101...	110101...



Latency vs. Throughput

Quote feedback

Latency is important

Volume is low

We can keep 'linger.ms' at 0

Location data

High-volume

Increased latency admissible

'linger.ms' larger than 0 - beneficial



Compression



Enabled by setting '`compression.type`'
Optimizes network and broker disk usage
Although disabled by default, usually recommended



Demo



Retries


Batching

Compression



Consumers





Only consumer – gets all partitions



Group
coordinator



Group
members

- C1



Only consumer – gets all partitions



Group
coordinator



Group
members

- C1
- C2





Group
coordinator



Revoke partitions



Group
members

- C1
- C2



Group leader



Member and
partition lists



Group
coordinator



Group
members

- C1
- C2



Group leader



Partition
assignments



Group
coordinator



Group
members

- C1
- C2



Group leader




Group coordinator



Group members

- C1
- C2



Assigned partitions



Rebalance



There is always a gap in processing



Some partitions will change ownership



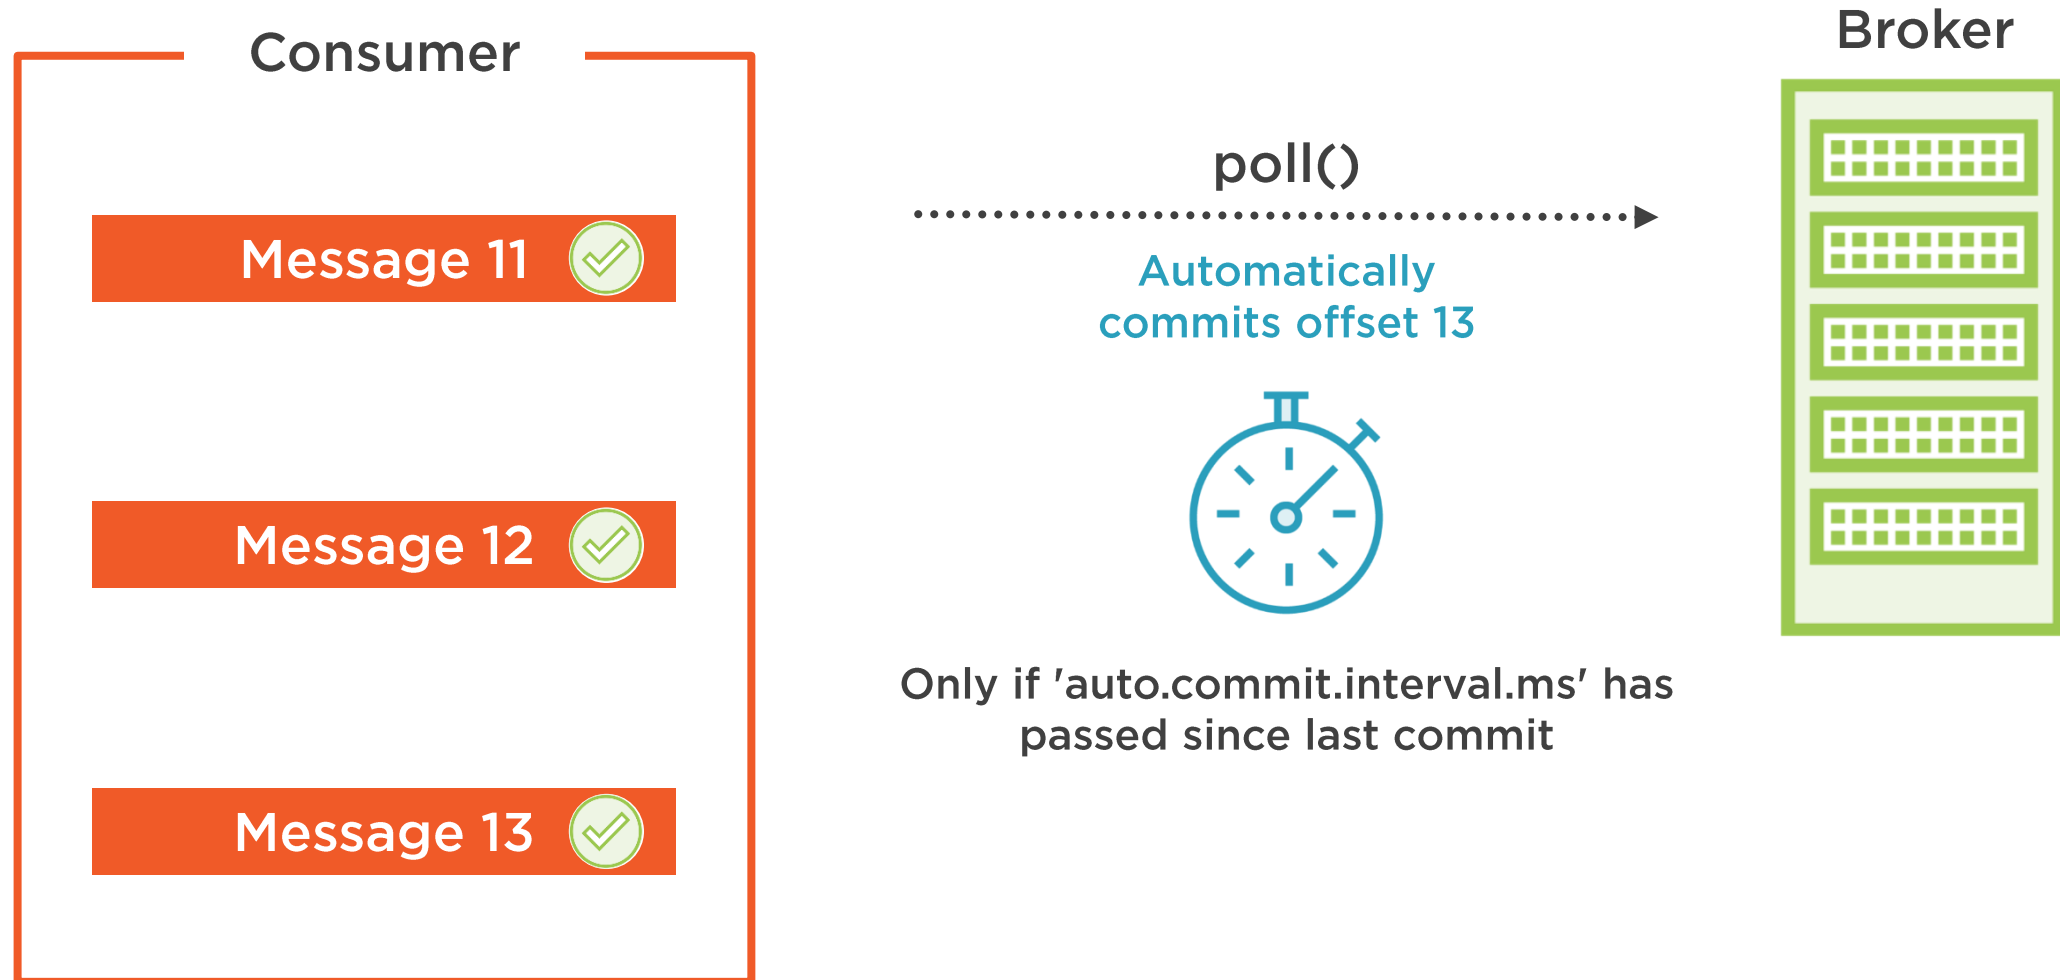
Static assignment is possible



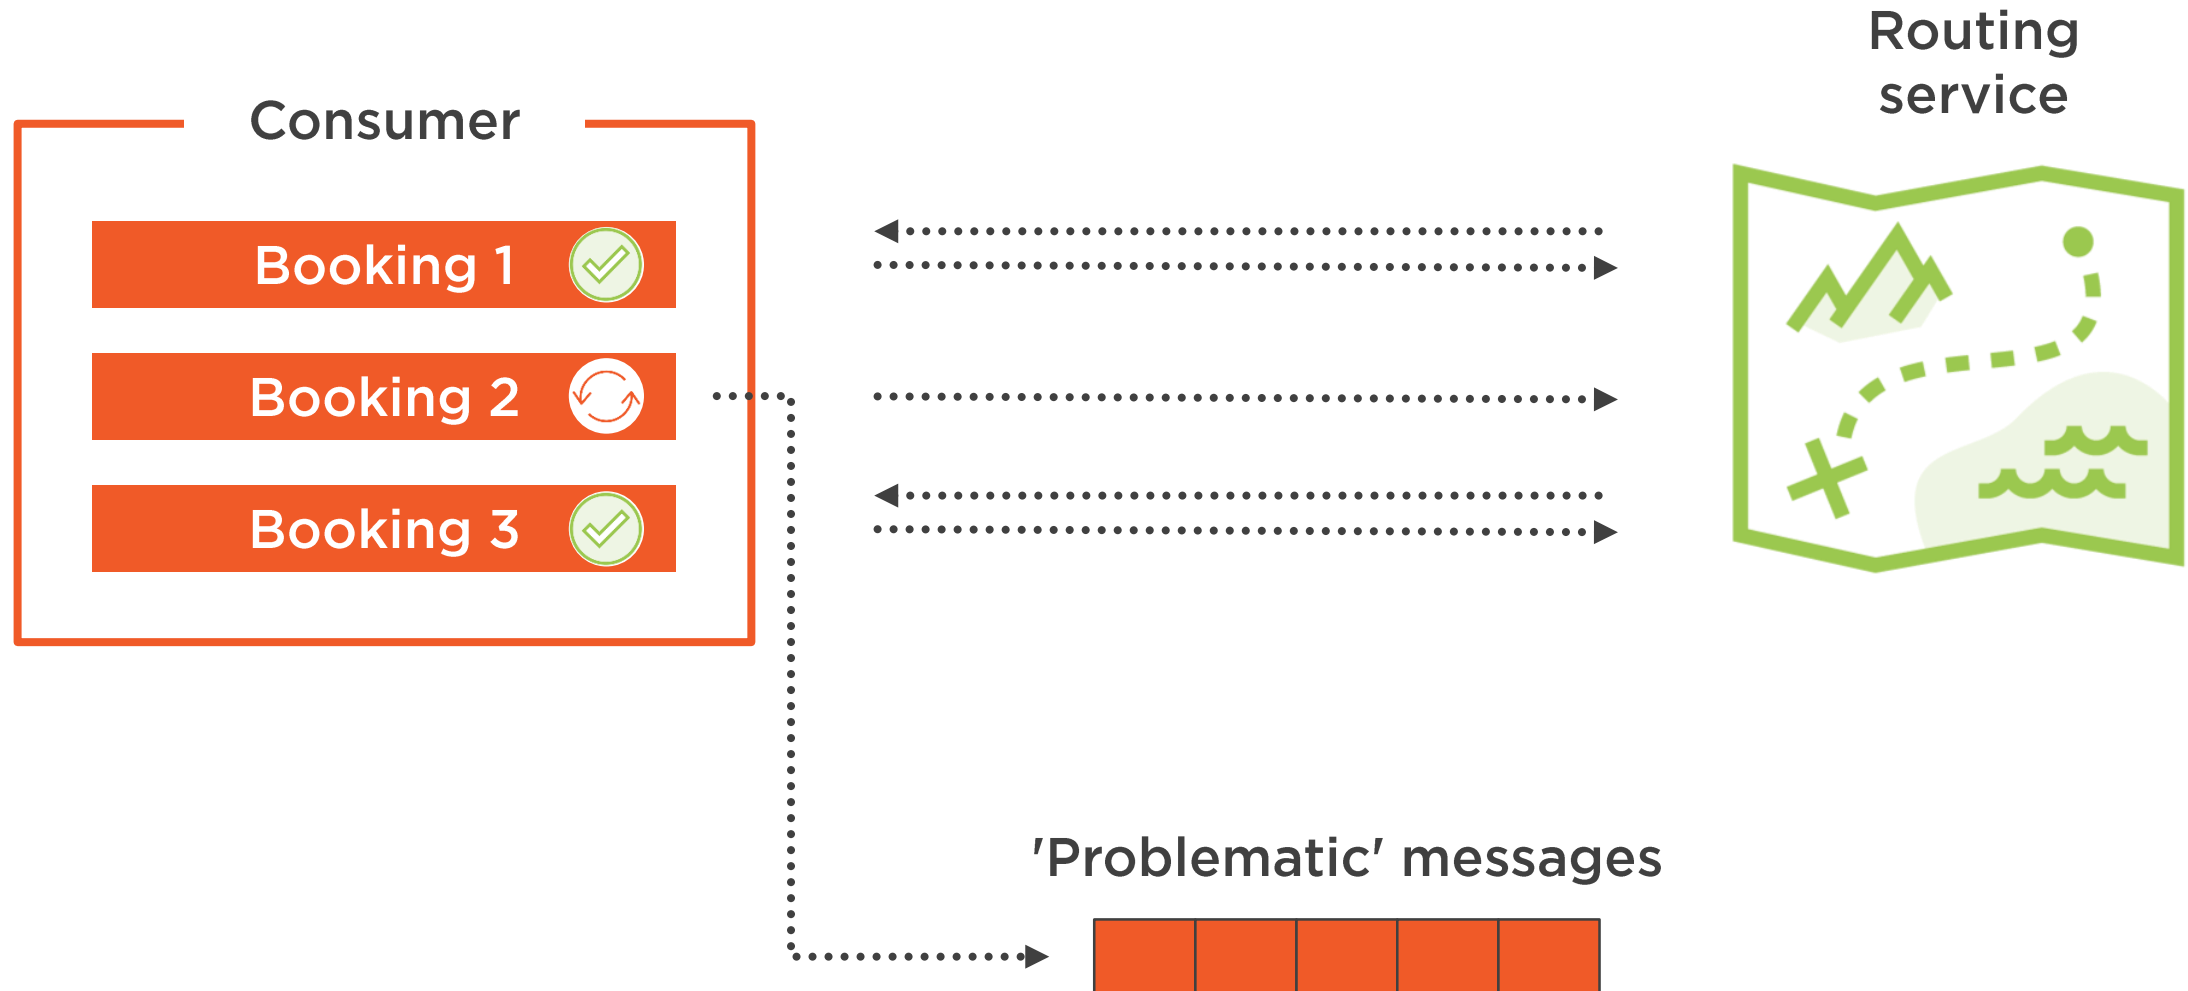
Offsets



Auto Commit



Auto Commit




Consumer Failures





heartbeat.interval.ms

Group
coordinator



consumer.close()



'I am leaving!'



session.timeout.ms





Heartbeats run in the background

They do not protect from consumer being stuck and not polling

Therefore the '`max.poll.interval.ms`' property

Demo



Offsets management

Consumer failure



Duplicate Messages

At-least-once guarantee

Exactly-once is tricky



Summary



Kafka is a complex technology

Core concepts necessary for reasoning about a system has been laid out

You always need to consult the docs

