Fault Tolerance and High Availability

Paweł Kordek
DATA ENGINEER

@pawel_kordek https://kordek.github.io

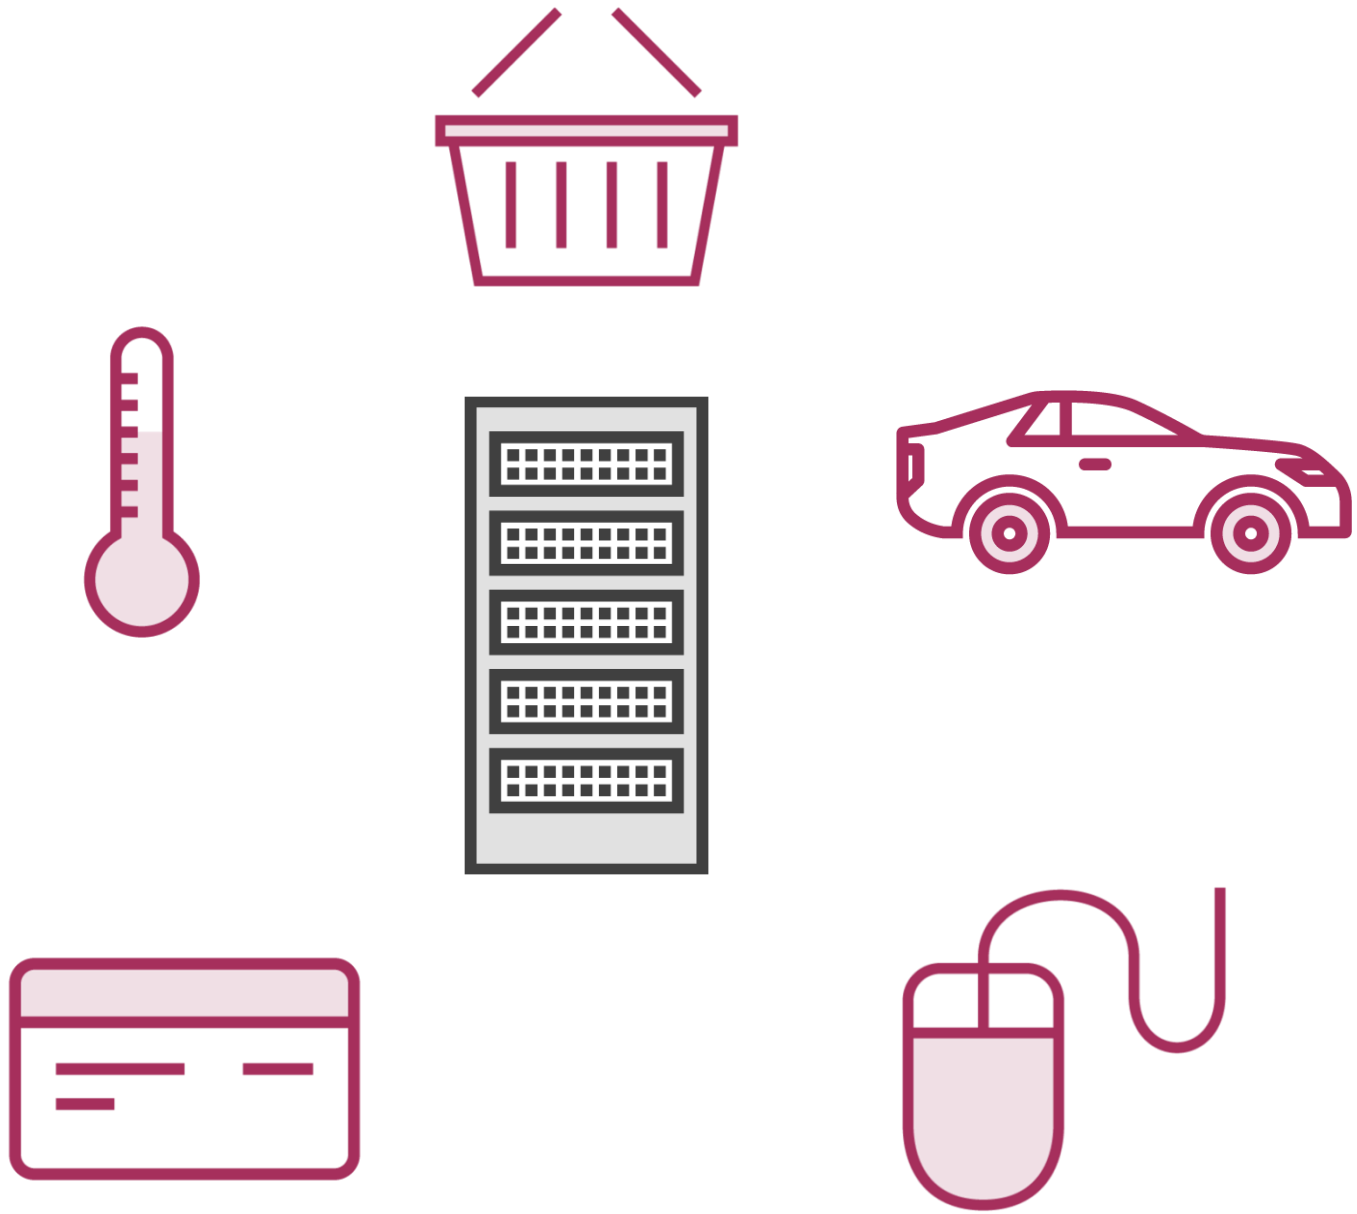

Using a single broker means accepting:

- Message loss
- Downtime

Desired Properties

Fault Tolerance

High Availability

Consistency

ZooKeeper

- Source of truth about cluster members
- Point of contact for new brokers
- Single ensemble can be shared with other applications

Message Persistence

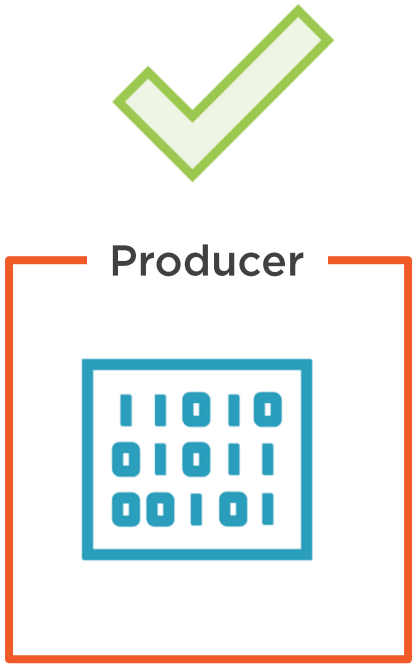

"acks" = "none"

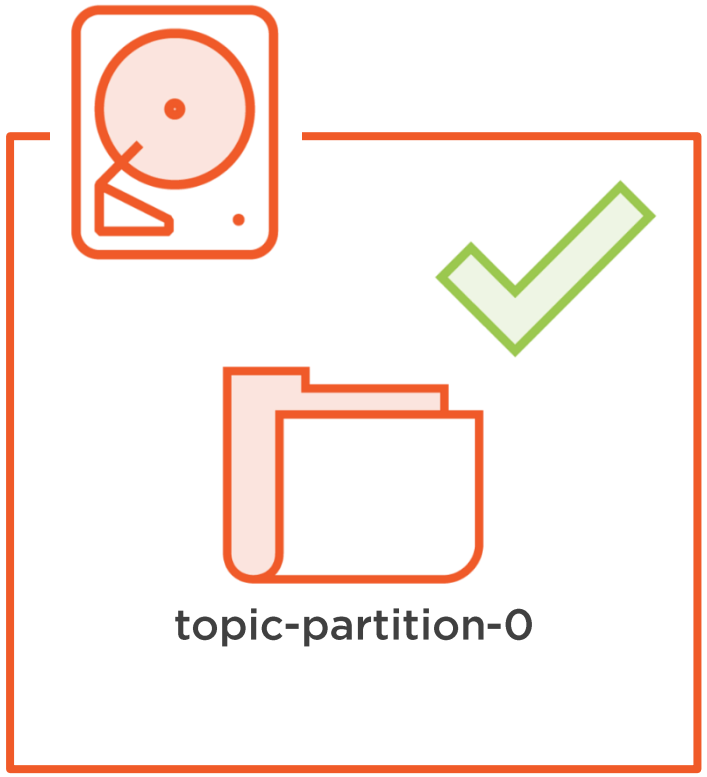

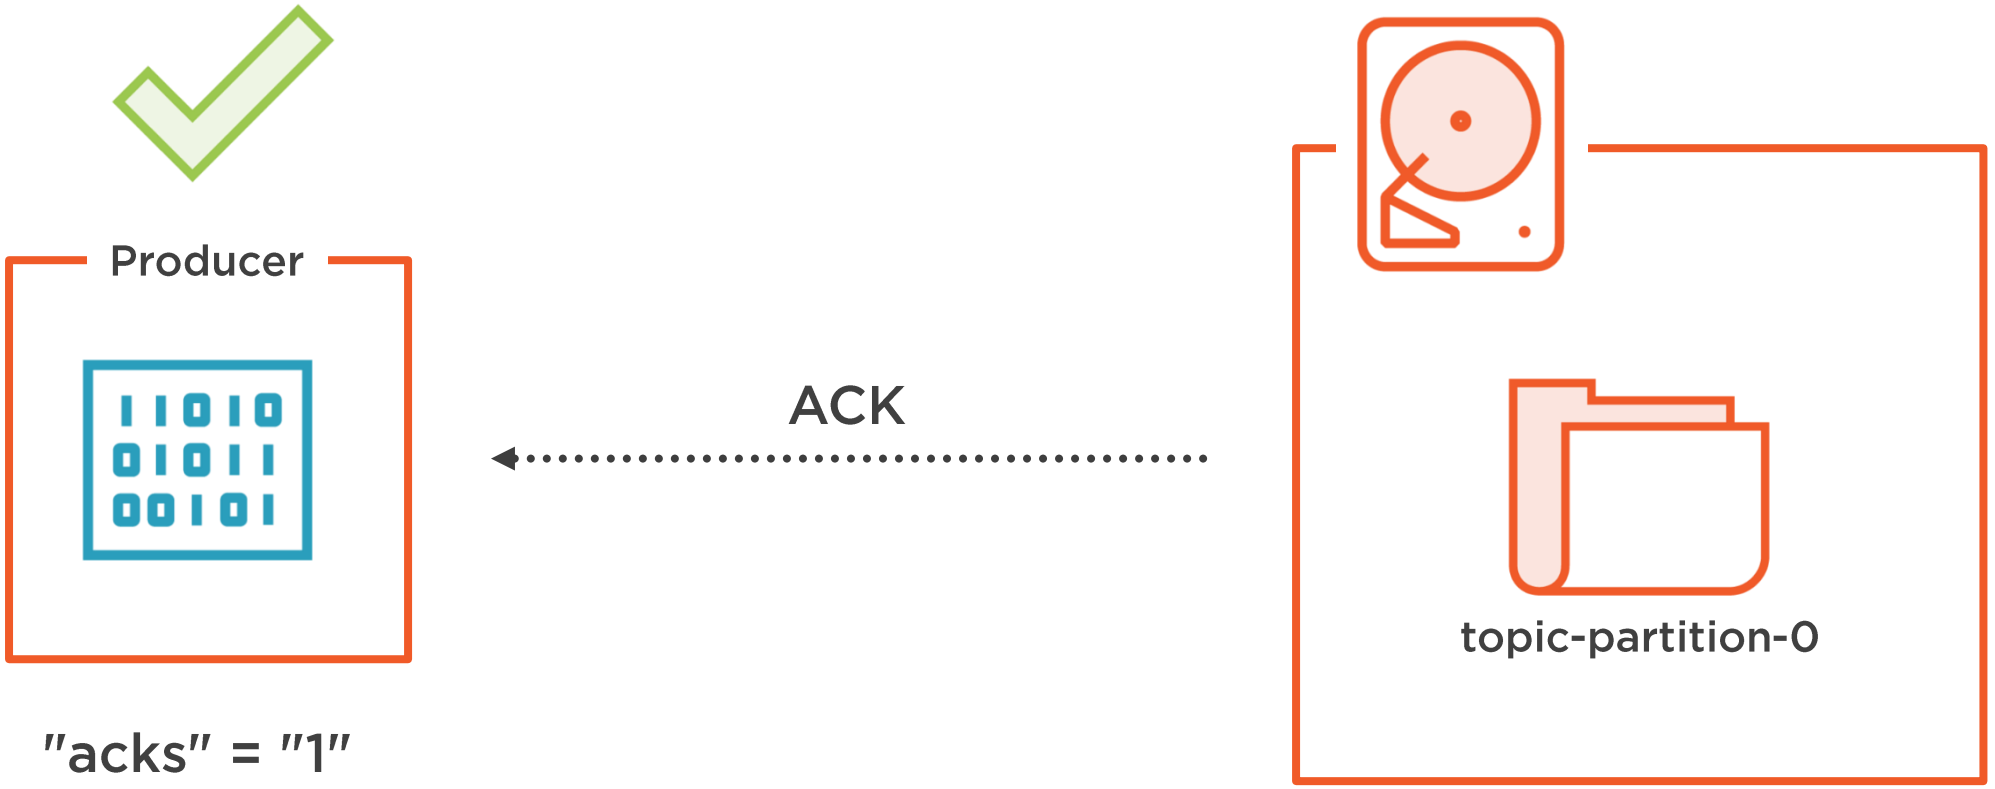
Producer
ACK
topic-partition-0
"acks" = "1"

Producer Delivery Guarantees

```
Properties props = new Properties();

. . .

KafkaProducer<String, String> = new KafkaProducer<>(props);

ProducerRecord<String, String> r = new ProducerRecord("topic", "key", "value");

RecordMetadata rm = producer.send(r).get();
```

When it returns depends on 'acks' value

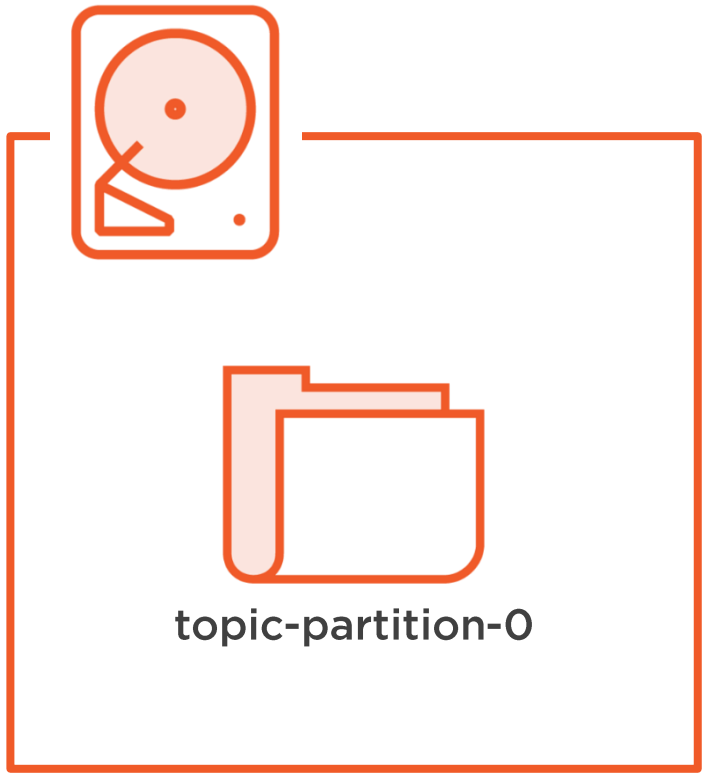

Only after a successful write, data is committed and made available to consumers.

Writing to Disk

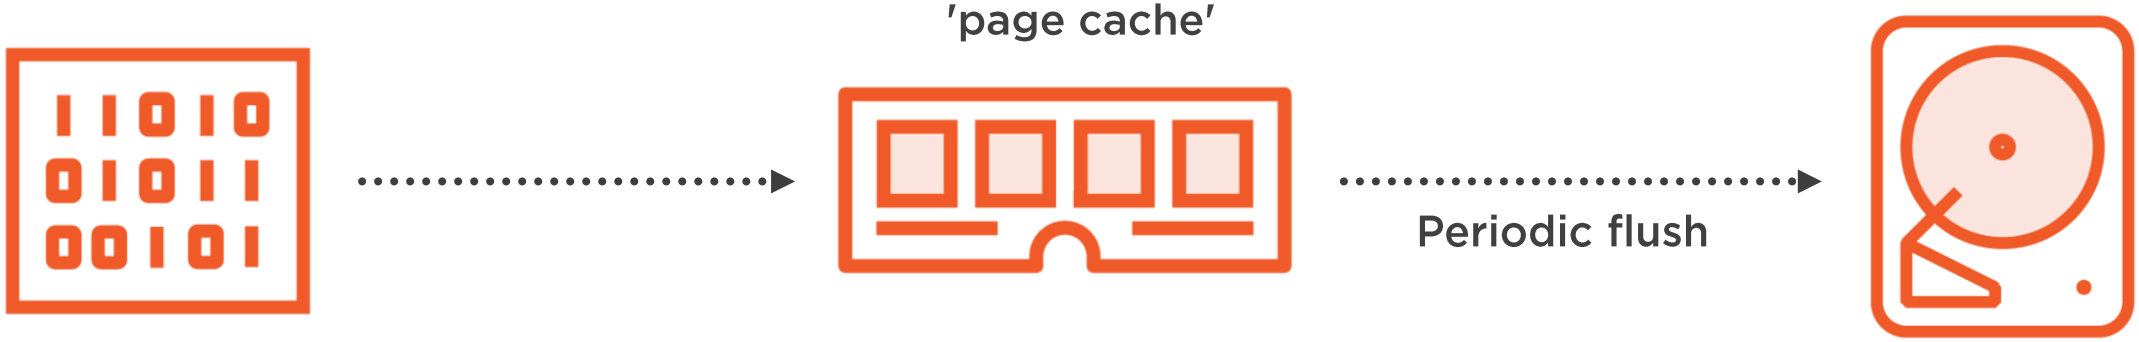

Replication

Creating a New Topic

```
Admin admin = Admin.create(
        Map.of("bootstrap.servers", "localhost:9092")
);

NewTopic topic = new NewTopic("quote-feedback", 2, (short) 1)

admin.createTopics(List.of(topic)).get()
```

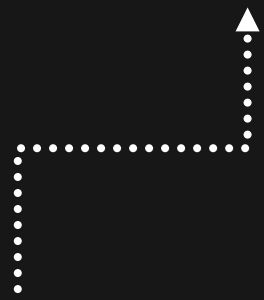

Creating a New Topic

```
Admin admin = Admin.create(
        Map.of("bootstrap.servers", "localhost:9092")
);

NewTopic topic = new NewTopic("quote-feedback", 2, (short) 3)

admin.createTopics(List.of(topic)).get()
```

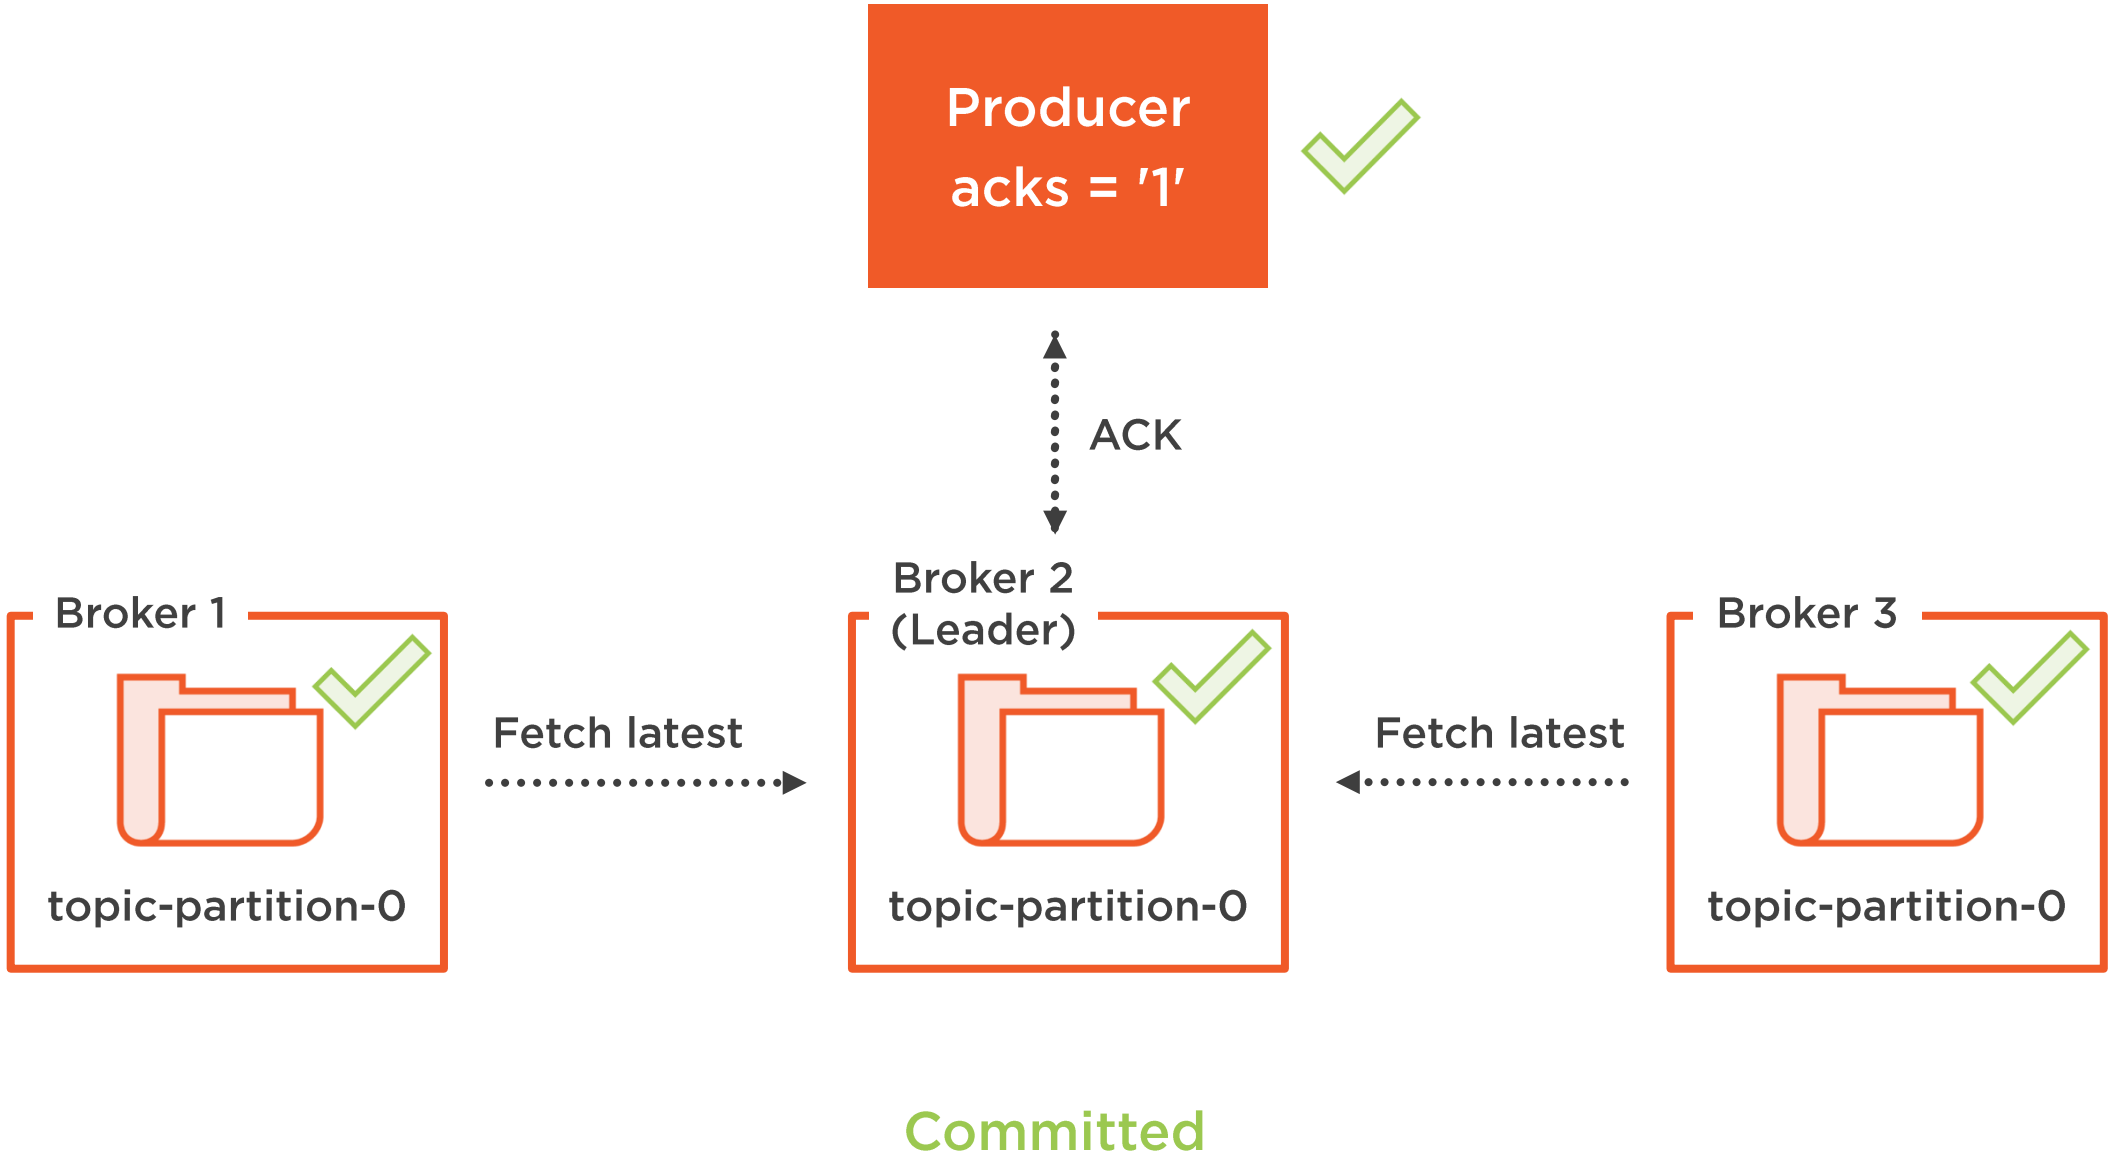
Producer
acks = '1'
ACK
Broker 1
topic-partition-0
Fetch latest
Broker 2
(Leader)
topic-partition-0
Fetch latest
Broker 3
topic-partition-0
Committed

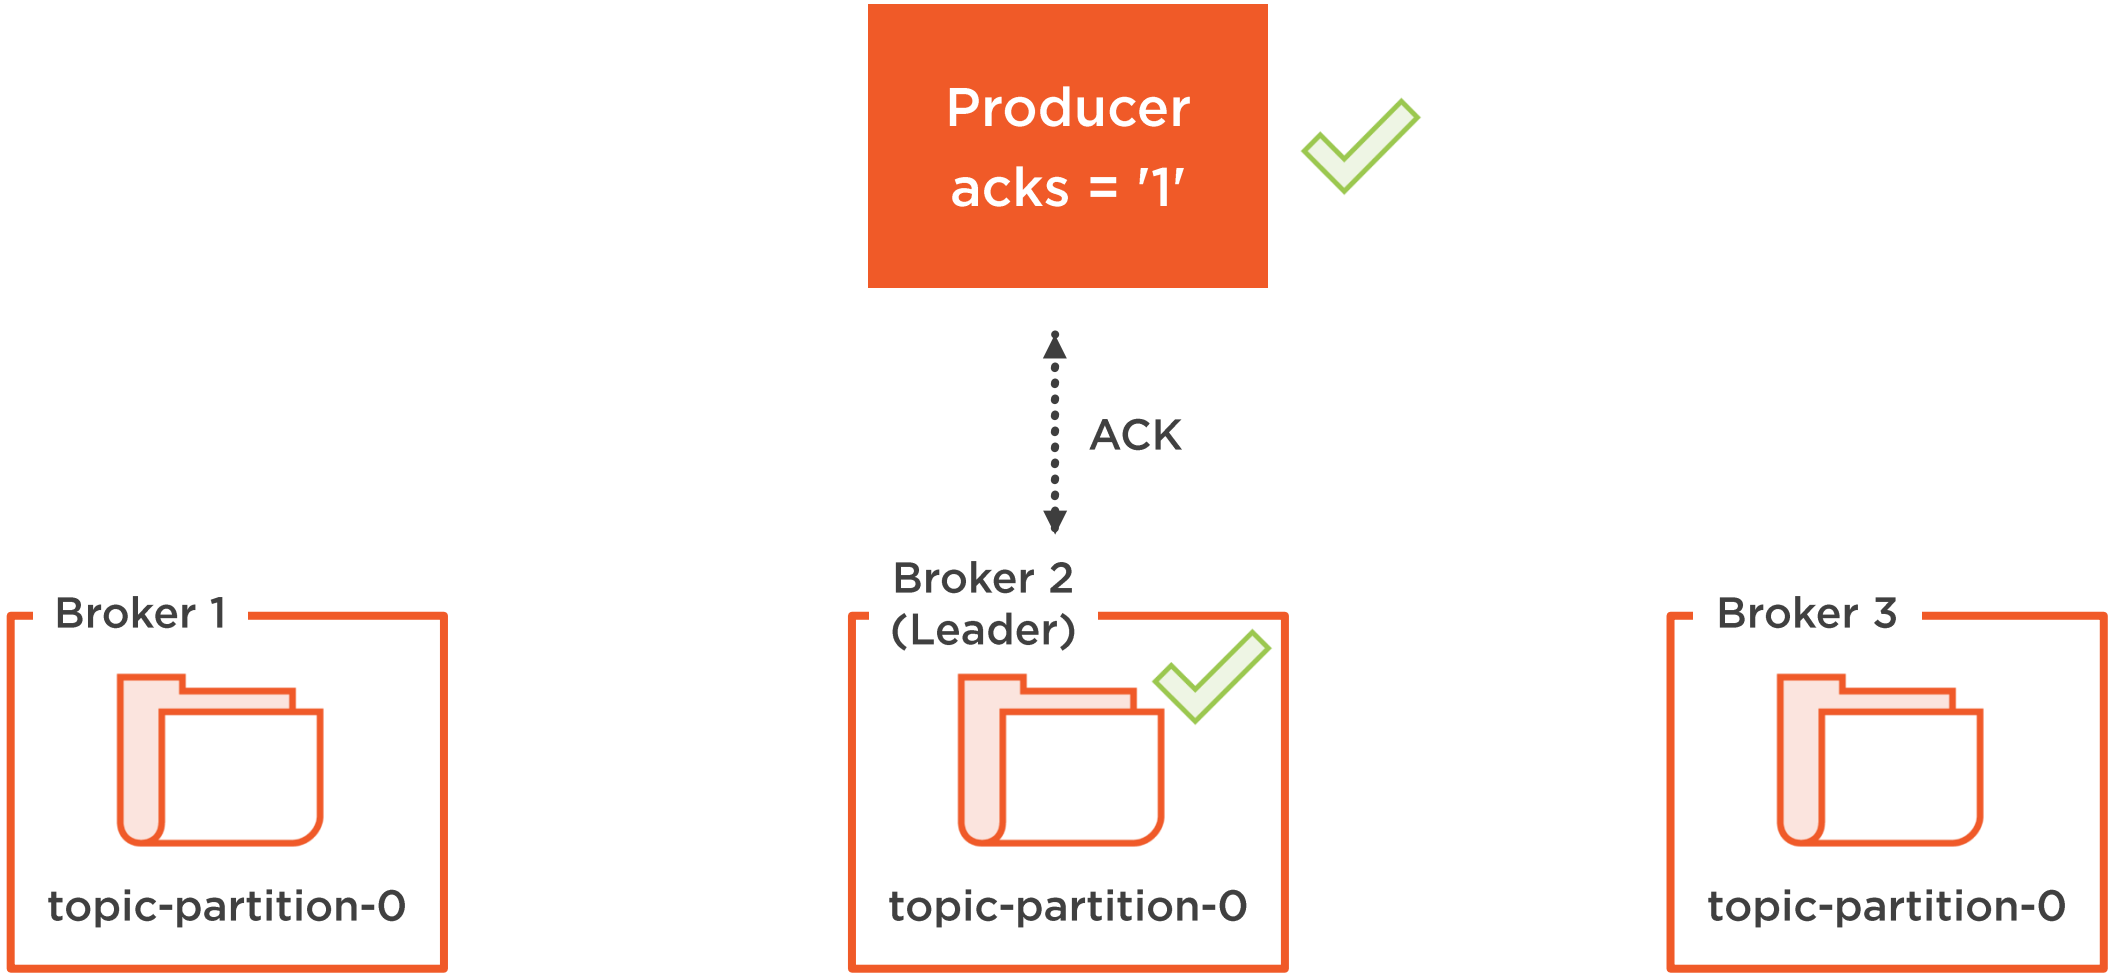
Producer
acks = '1'
ACK
Broker 1
topic-partition-0
Broker 2
(Leader)
topic-partition-0
Broker 3
topic-partition-0

Producer
acks = '1'

Producer
acks = '1'

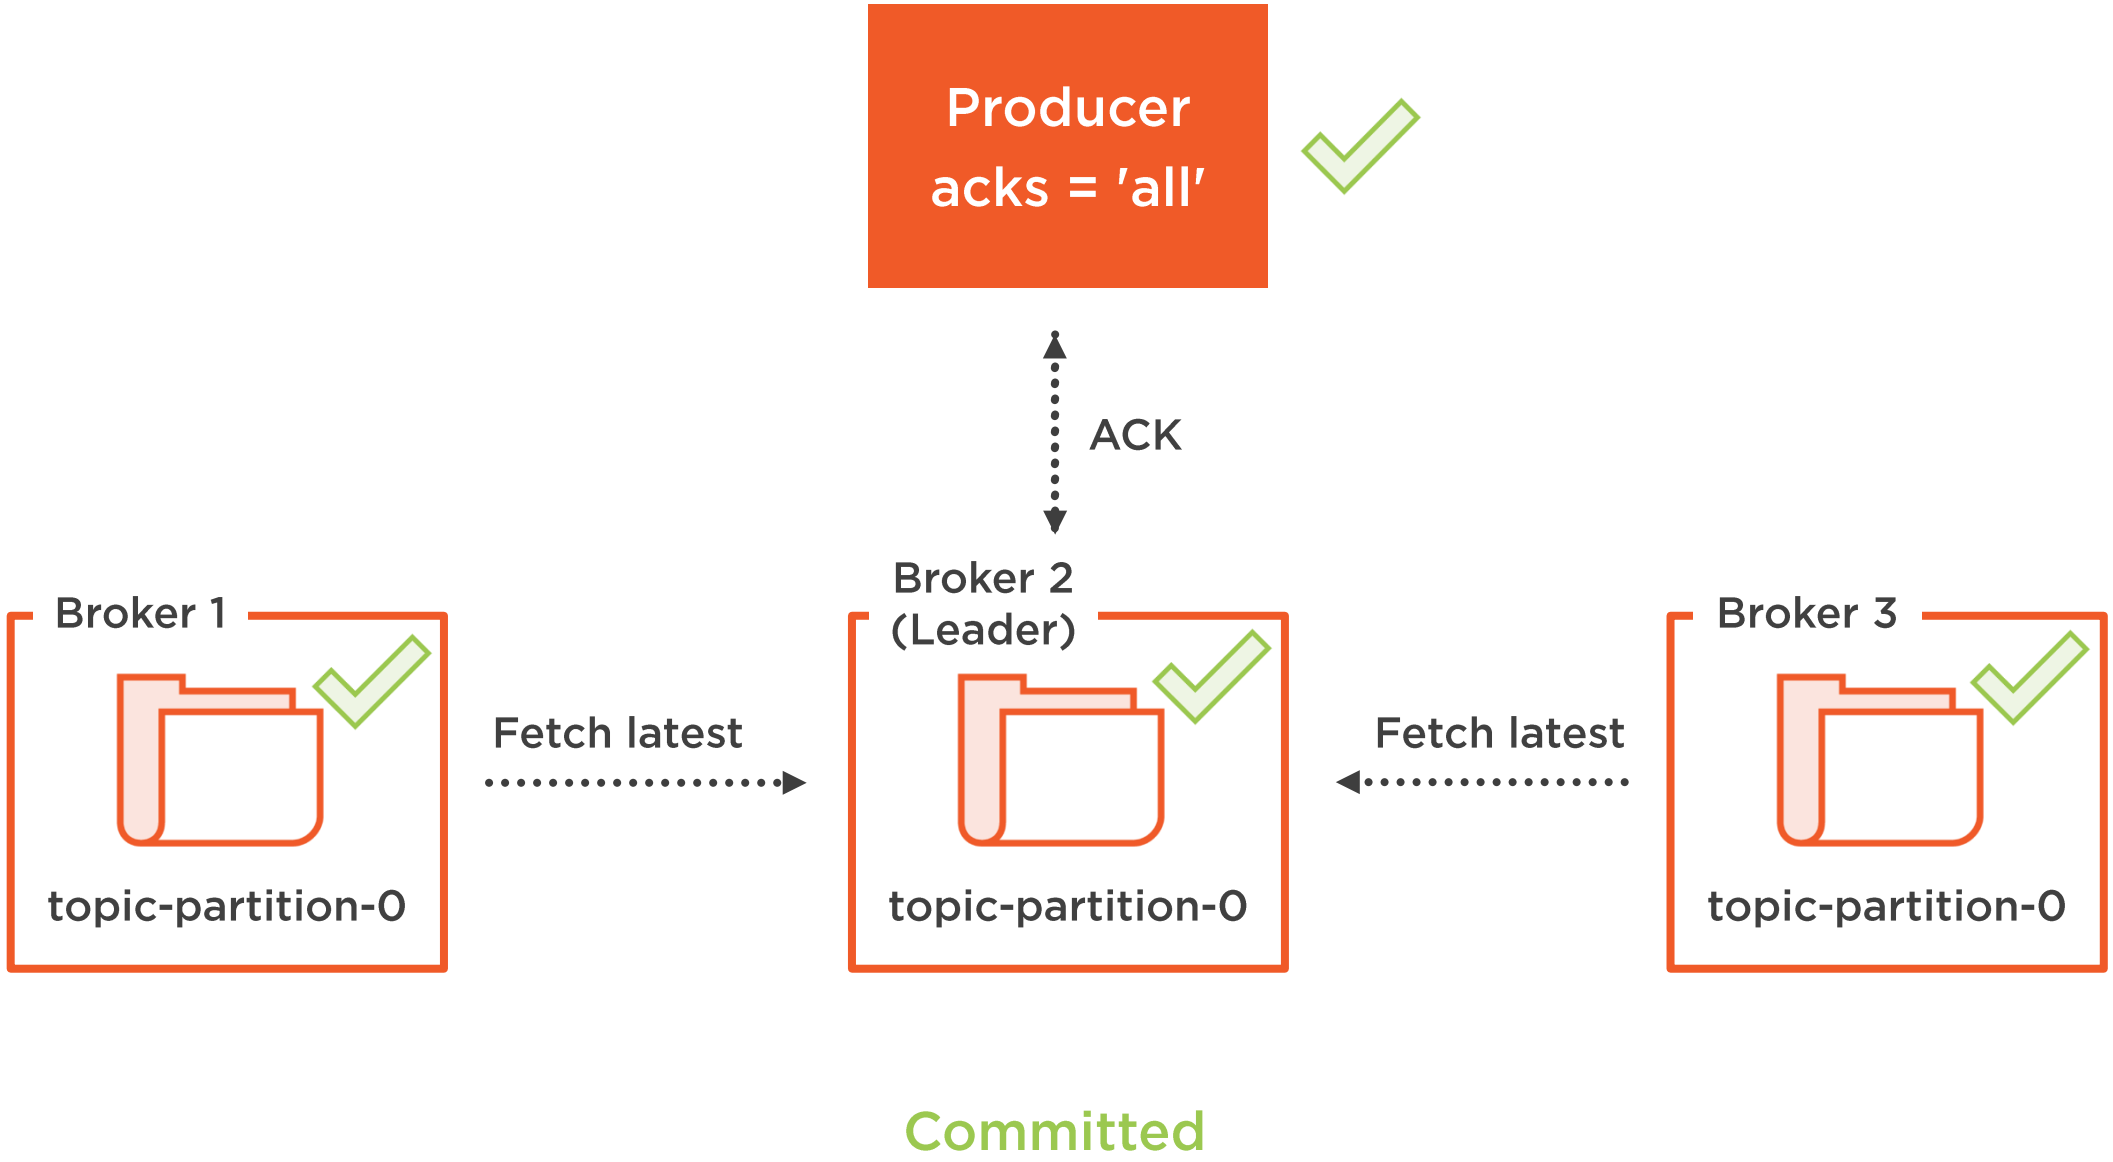
Producer
acks = 'all'
ACK
Broker 1
topic-partition-0
Fetch latest
Broker 2
(Leader)
topic-partition-0
Fetch latest
Broker 3
topic-partition-0
Committed

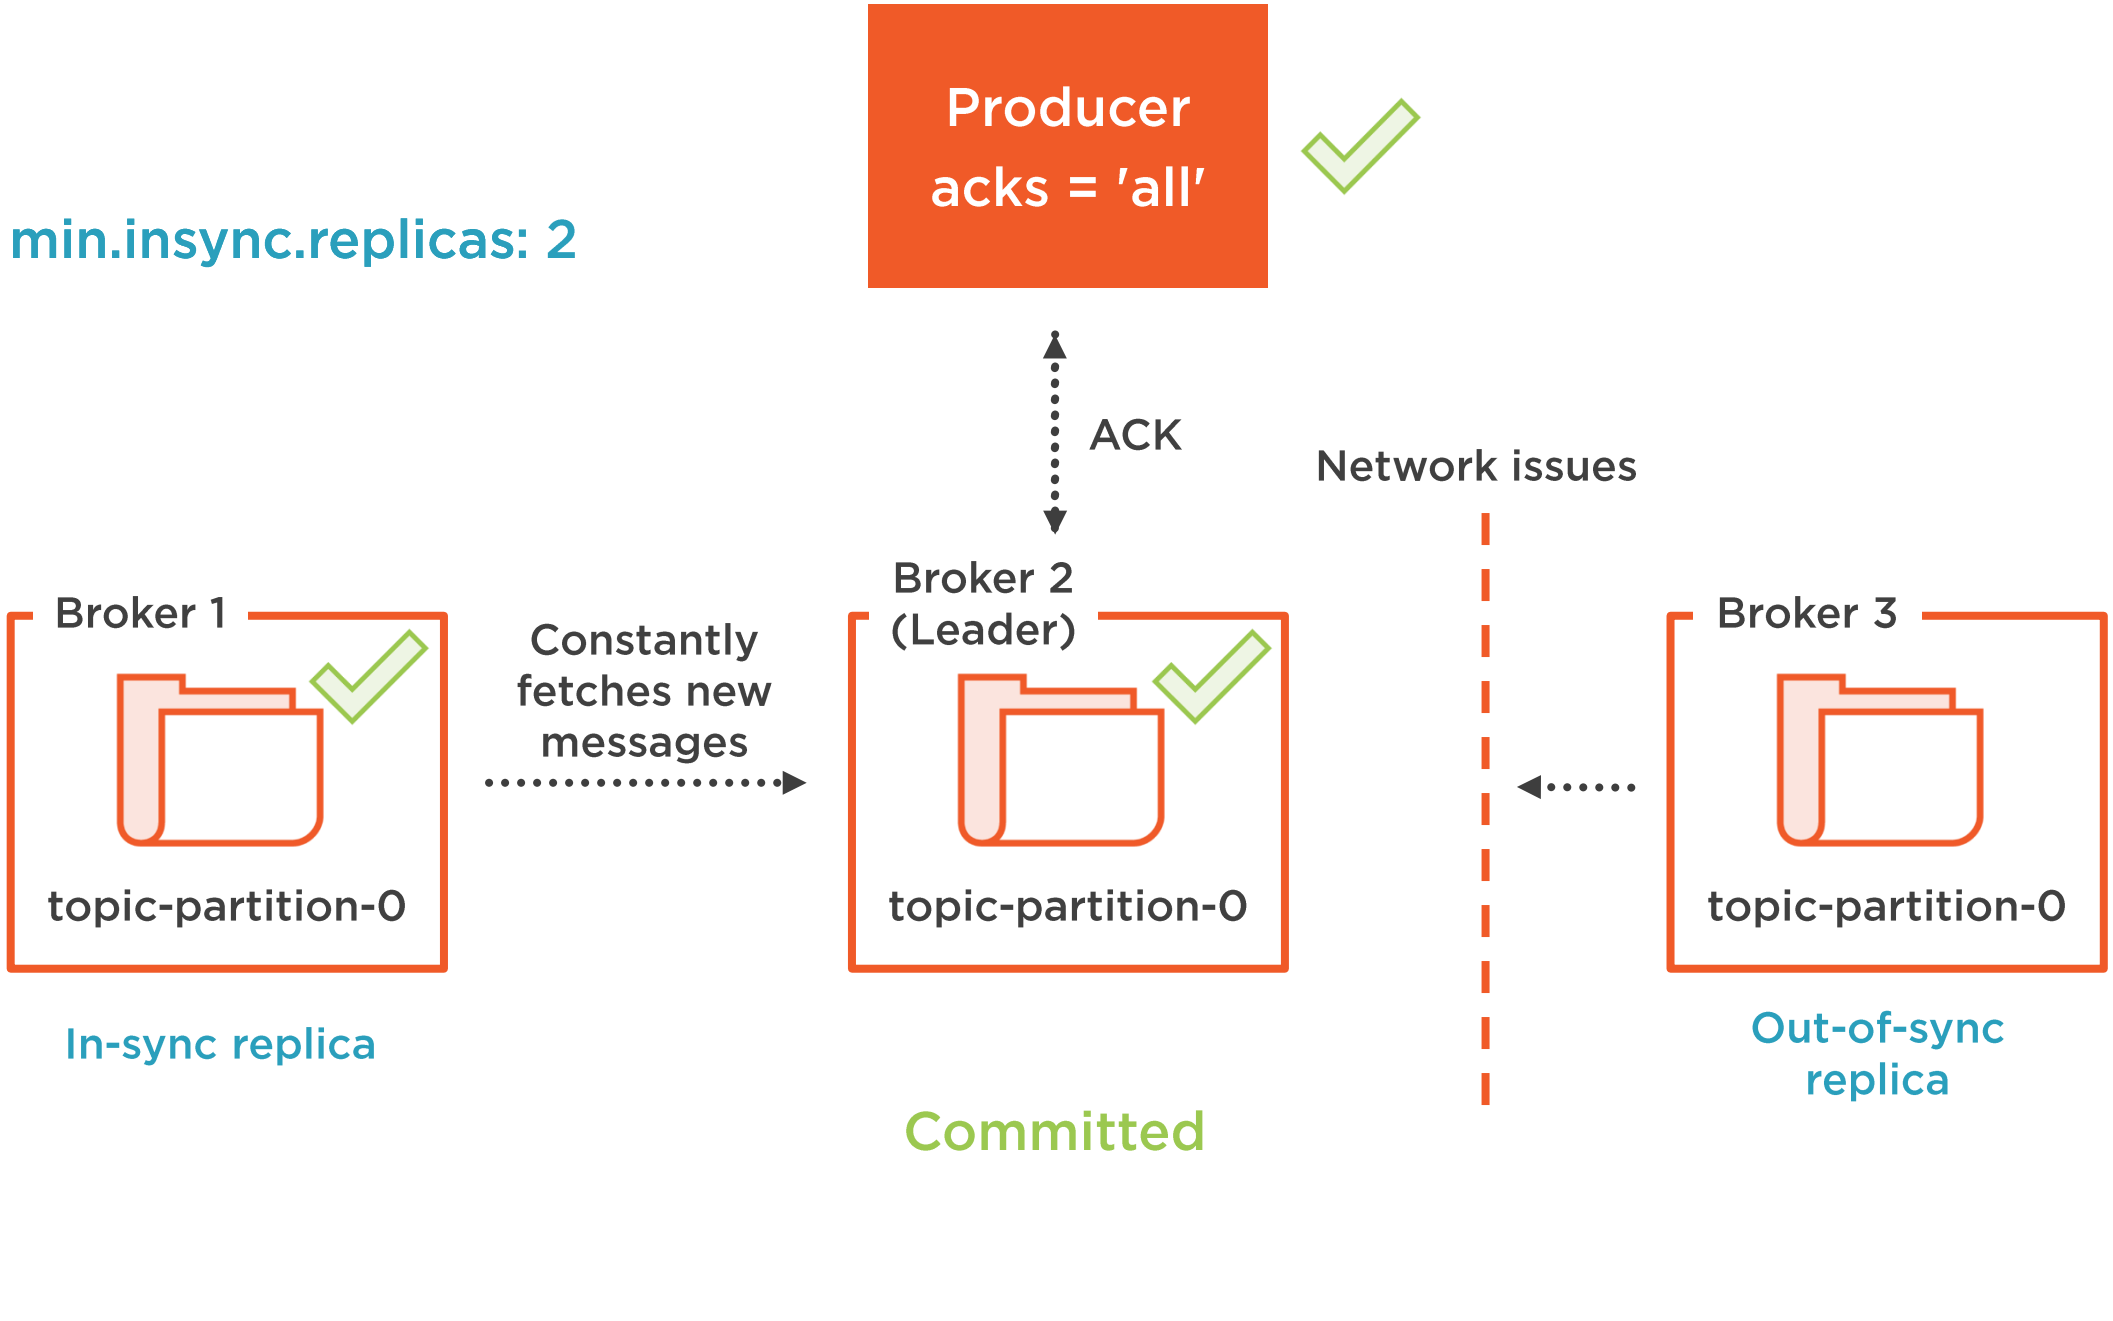
min.insync.replicas: 2
Producer
acks = 'all'
ACK
Network issues
Broker 1
Constantly
fetches new
messages
Broker 2
(Leader)
Broker 3
topic-partition-0
topic-partition-0
topic-partition-0
In-sync replica
Out-of-sync
replica
Committed

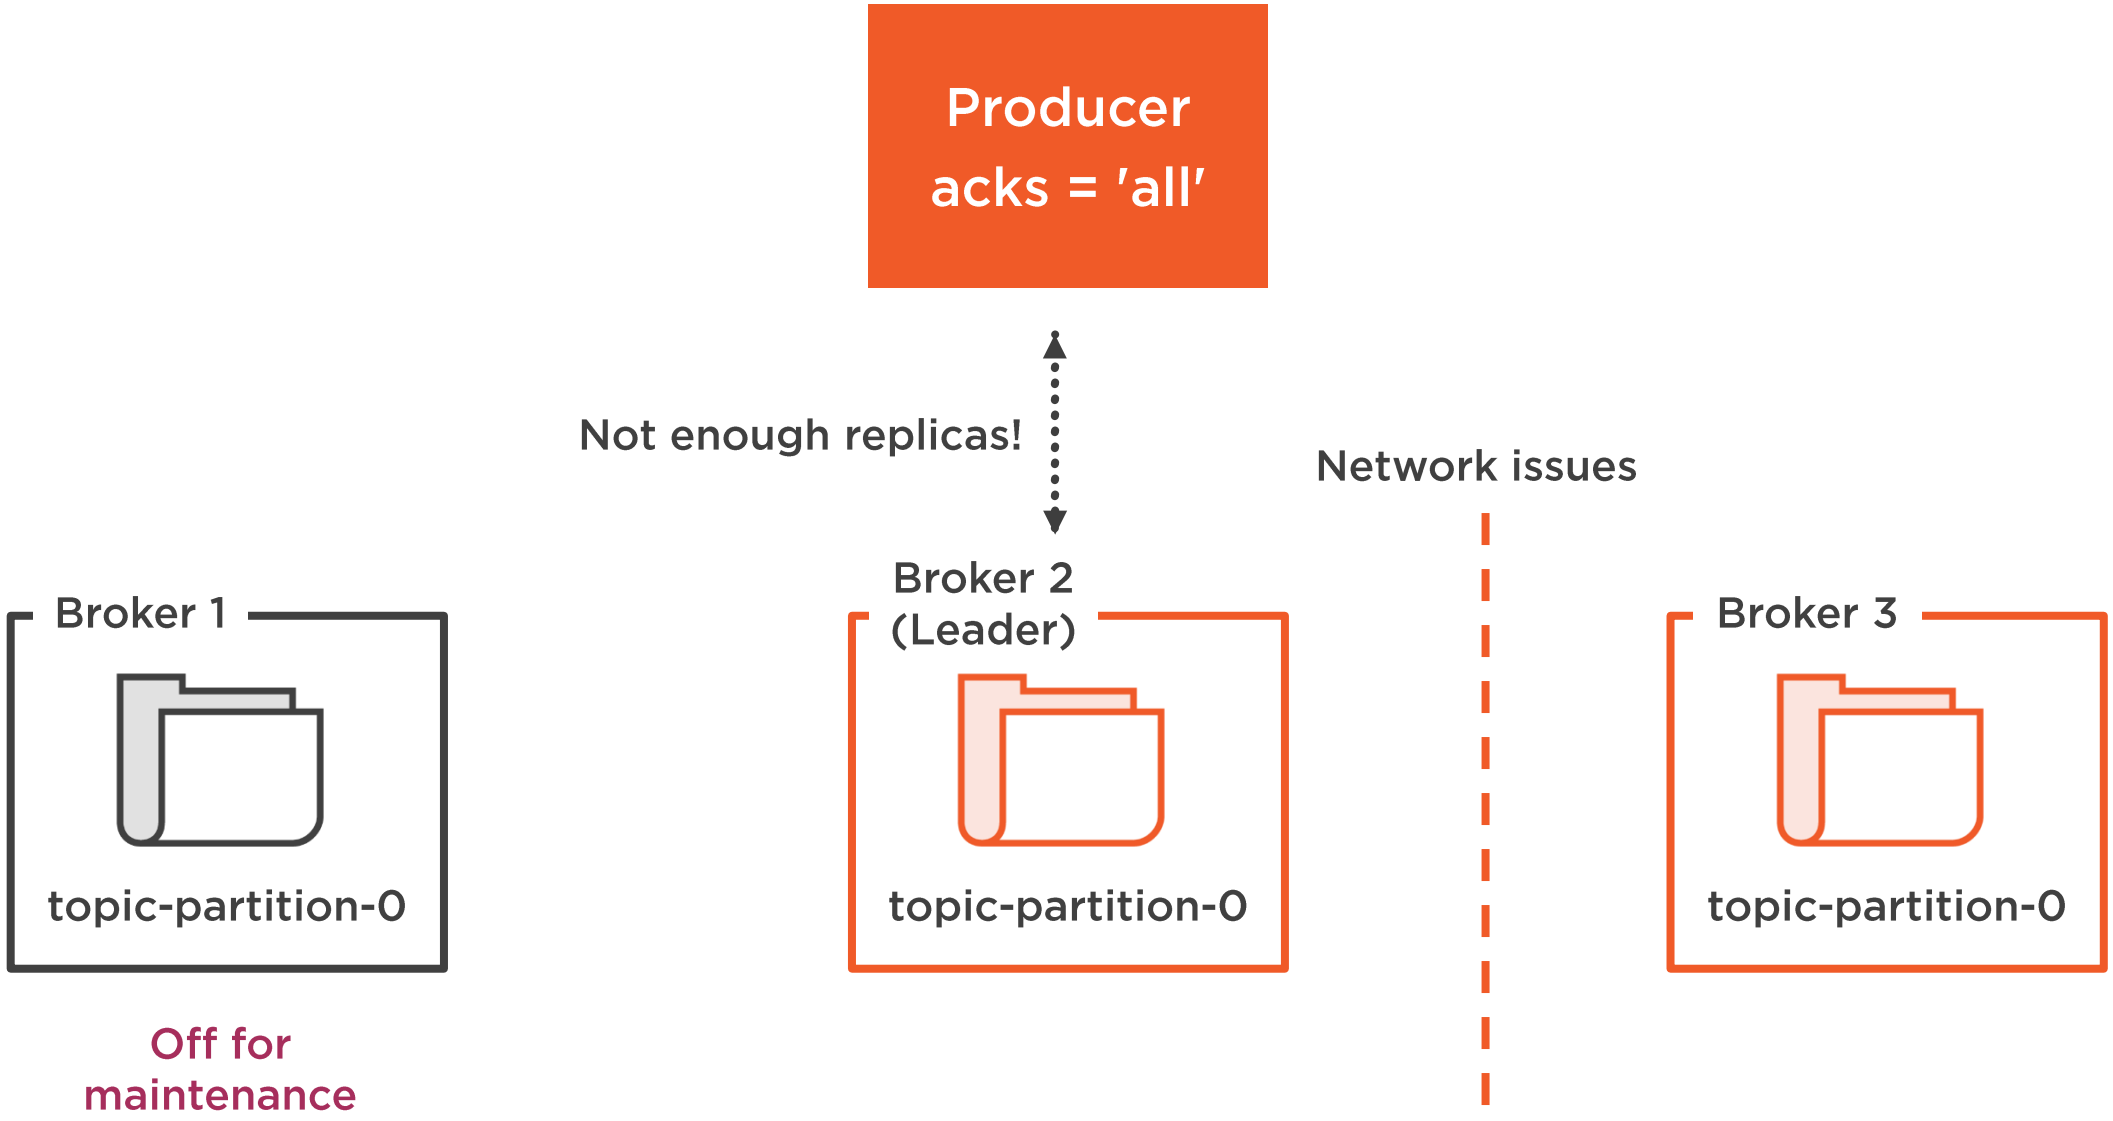
Producer
acks = 'all'
Not enough replicas!
Network issues
Broker 1
topic-partition-0
Off for
maintenance
Broker 2
(Leader)
topic-partition-0
Broker 3
topic-partition-0

Data Retention and Cluster Sizing

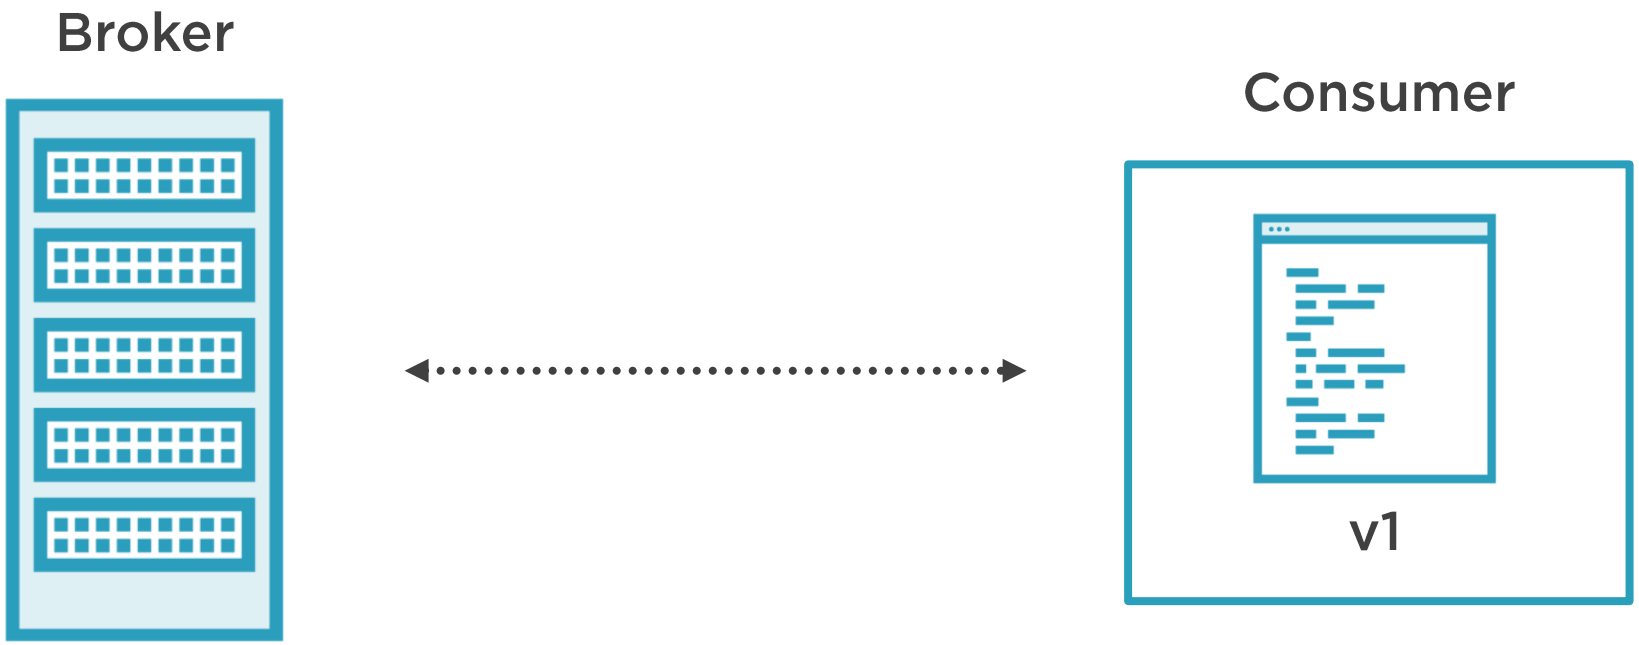
Broker
Consumer
v1

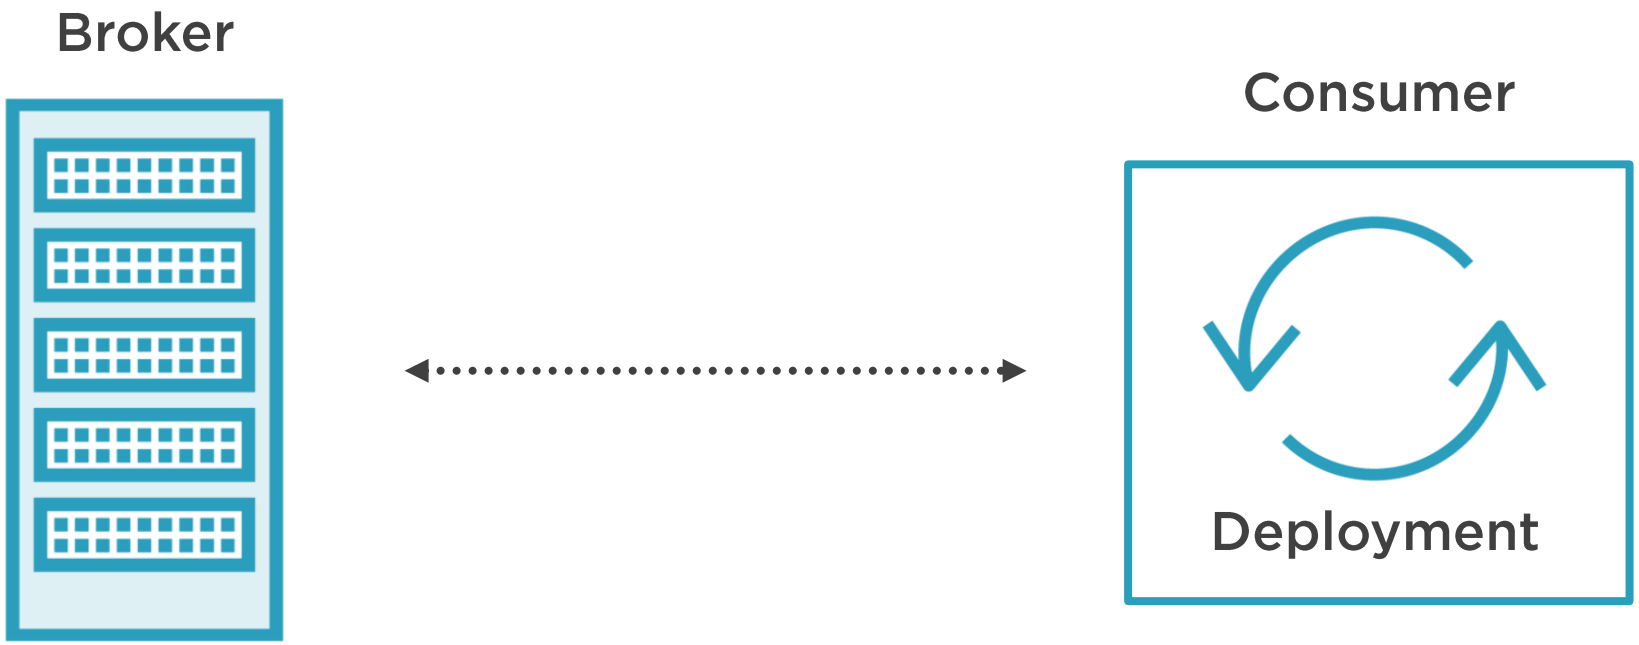
Broker
Consumer
Deployment

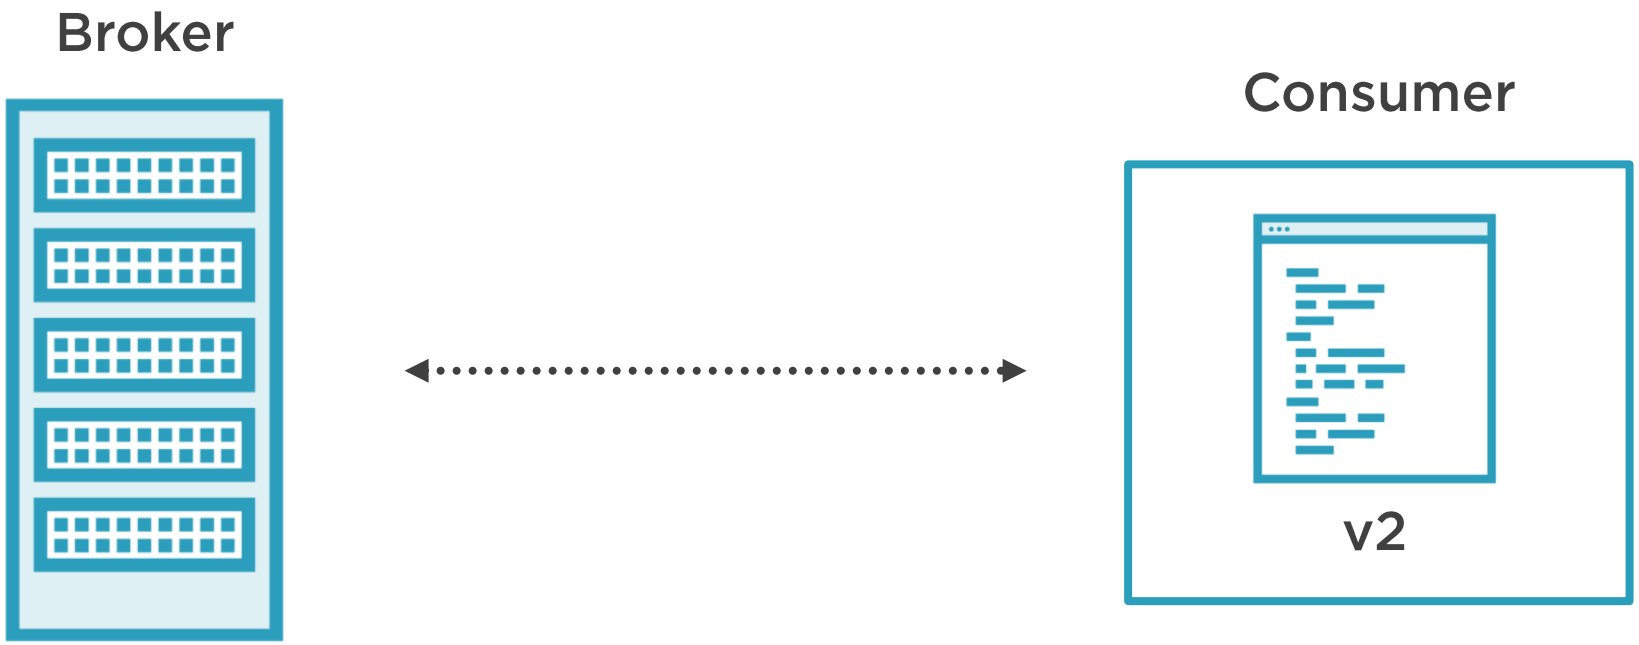
Broker
Consumer
v2

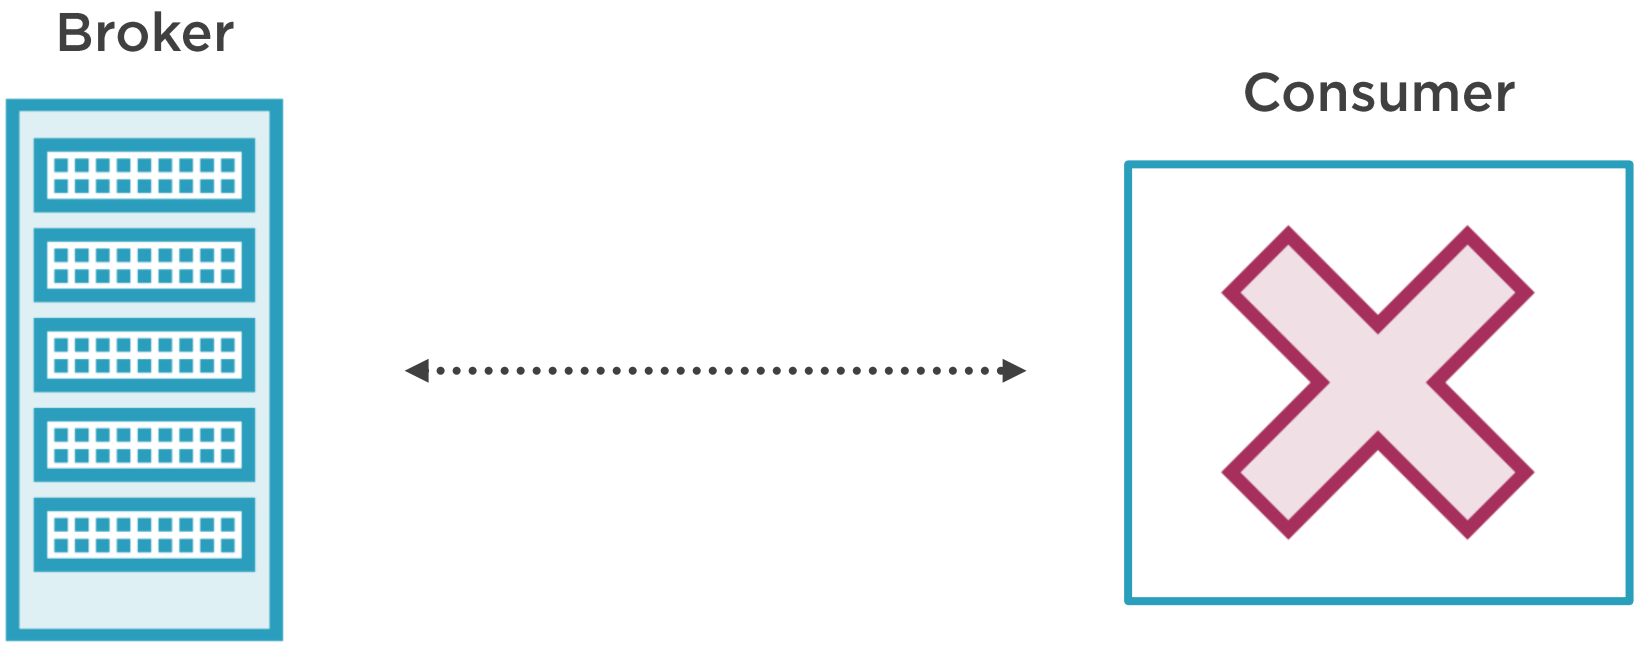
Broker
Consumer

Data Retention

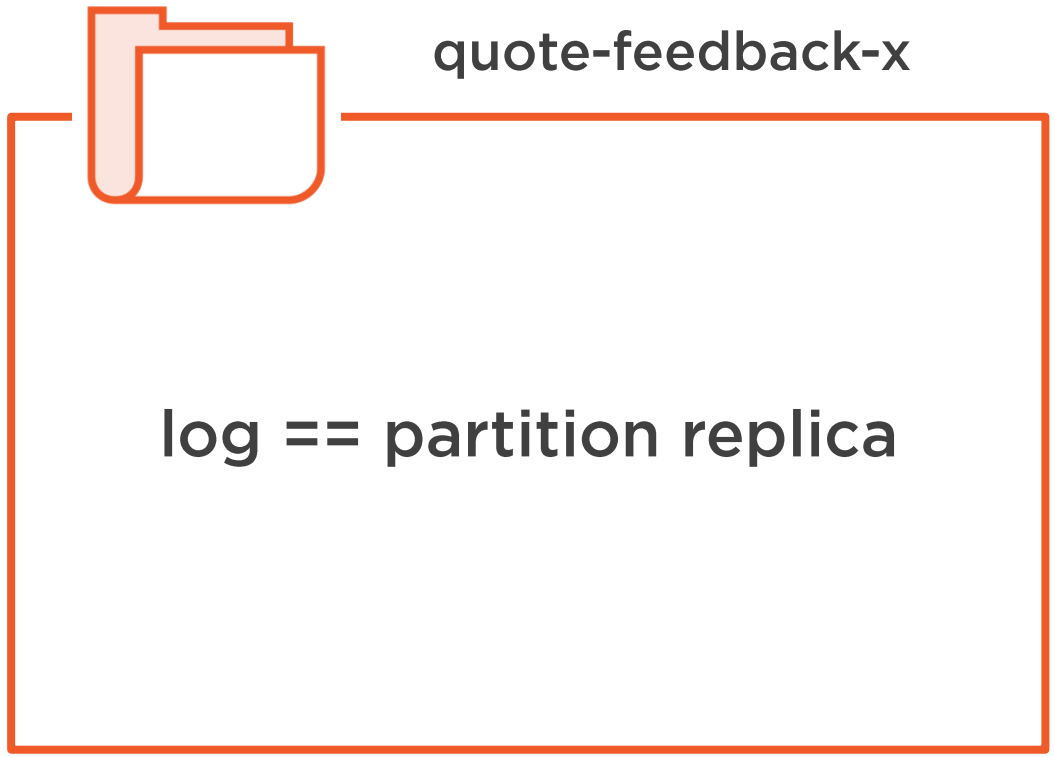

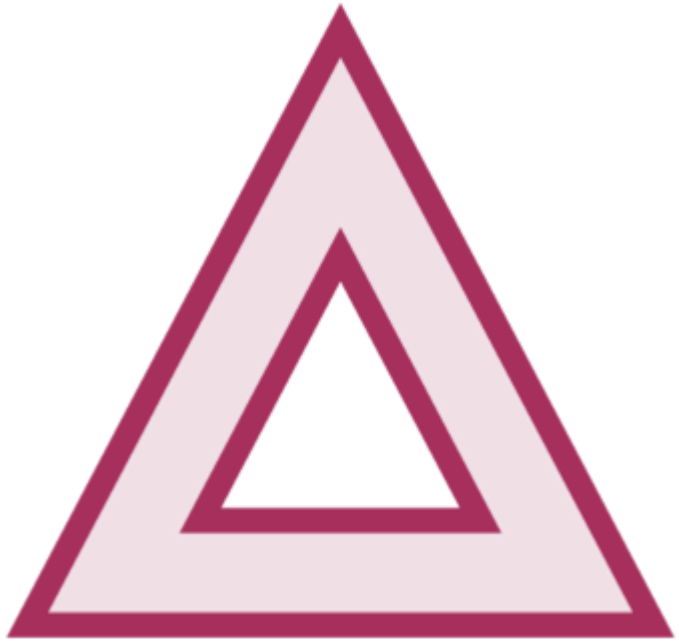

Think about possible log size when choosing the number of partitions for a topic

Data Retention

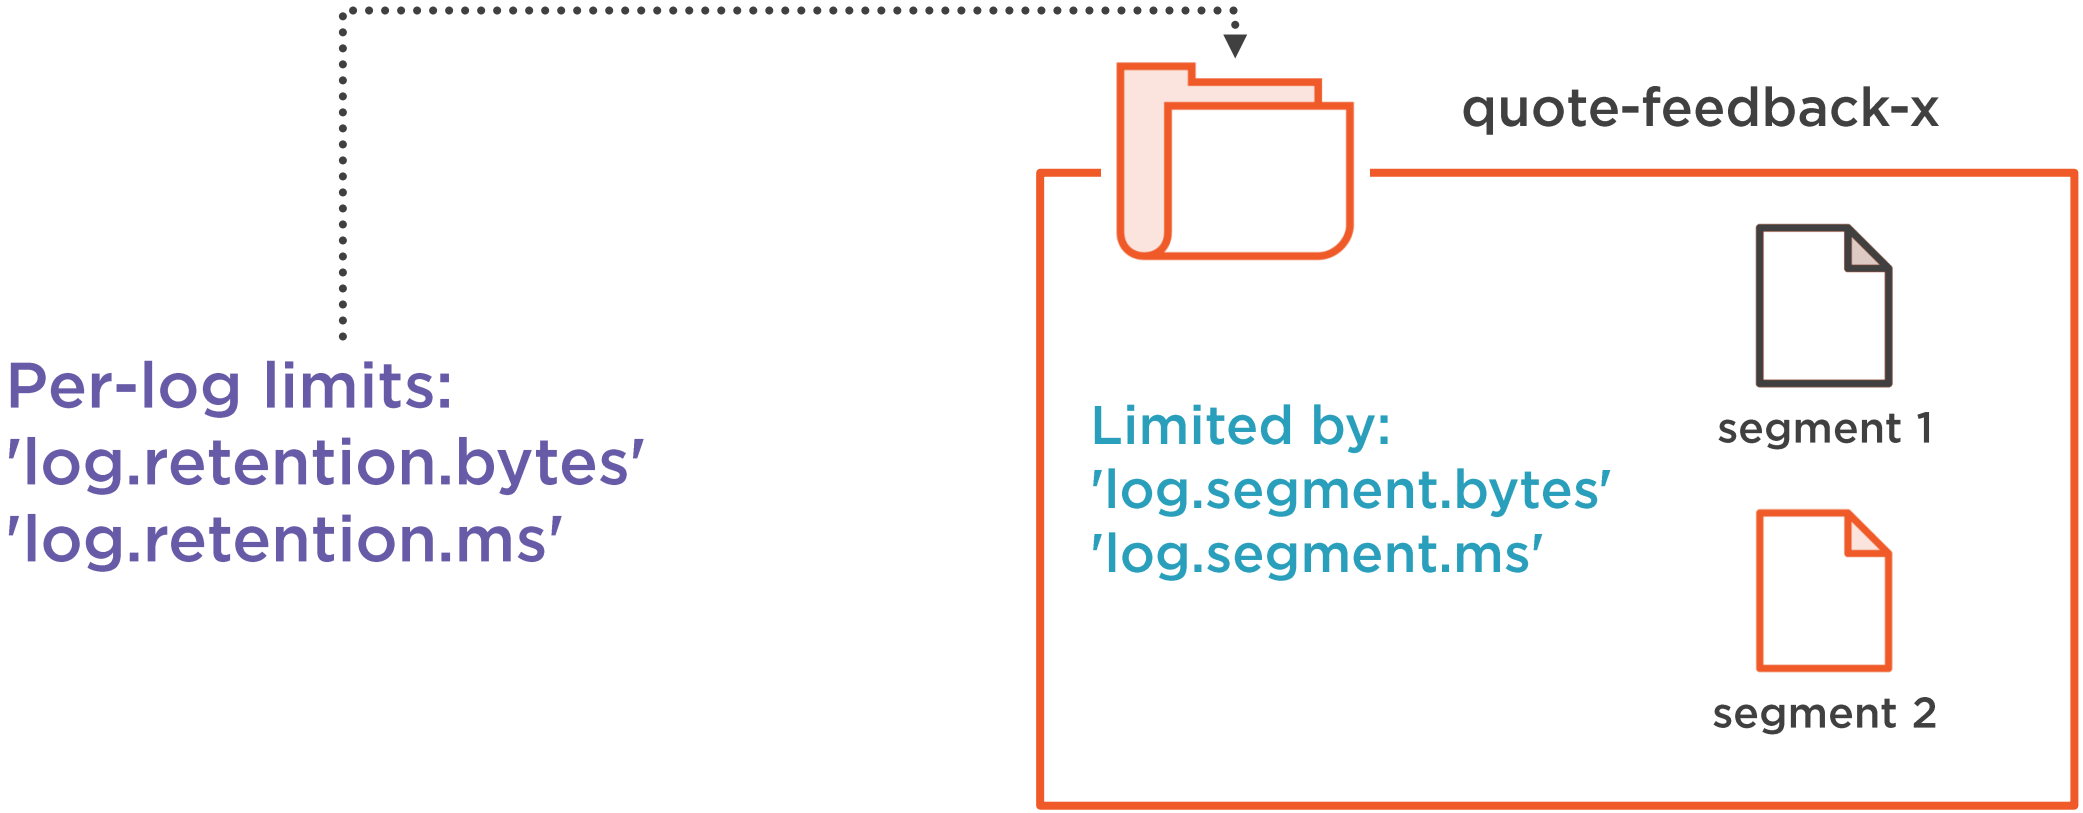

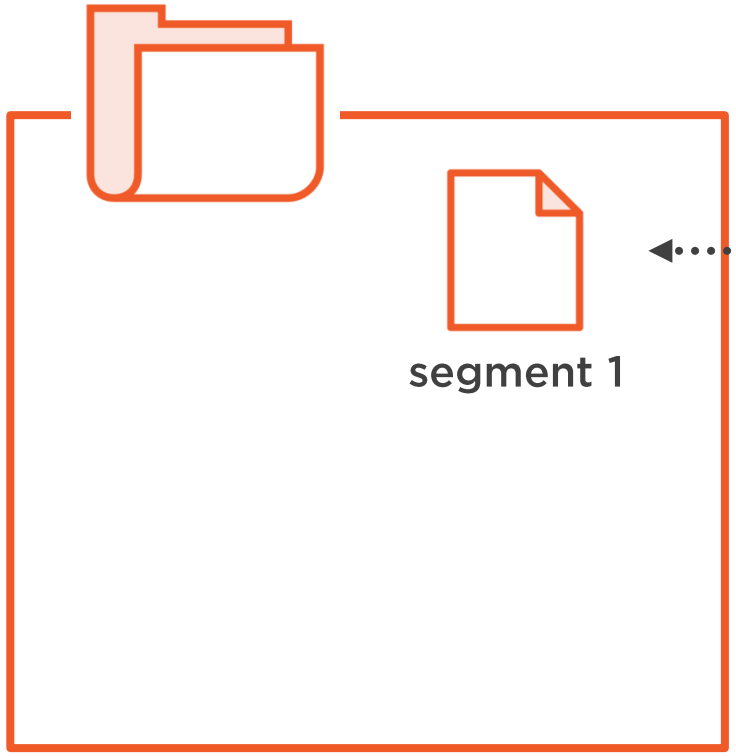

Day 1: ~500MB – cannot close

log.segment.bytes: 1GB
log.segment.ms: 604800000 (1 week)
log.retention.bytes: 2GB
log.retention.ms: 86400000 (1 day)

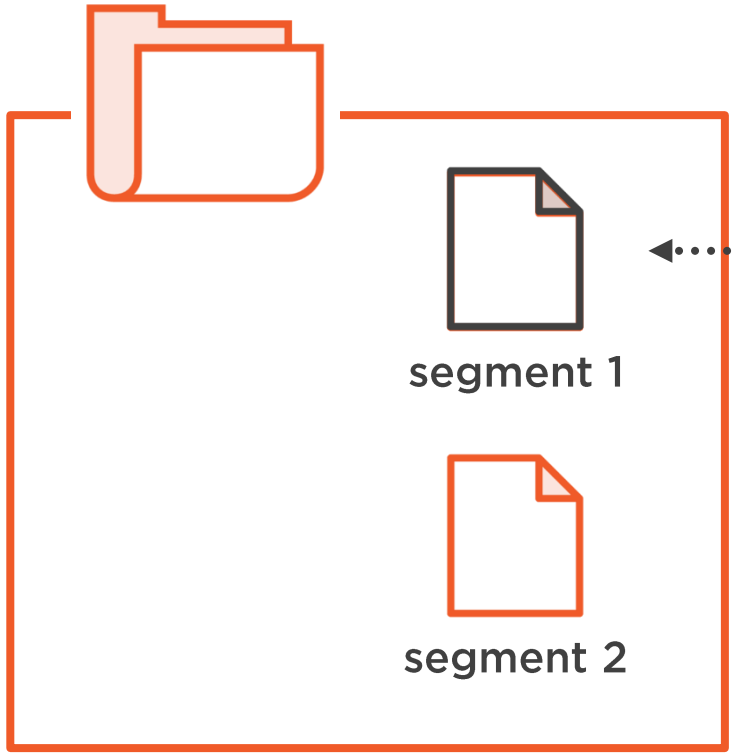

Day 2: > 1 GB – OK to close

log.segment.bytes: 1GB
log.segment.ms: 604800000 (1 week)
log.retention.bytes: 2GB
log.retention.ms: 86400000 (1 day)

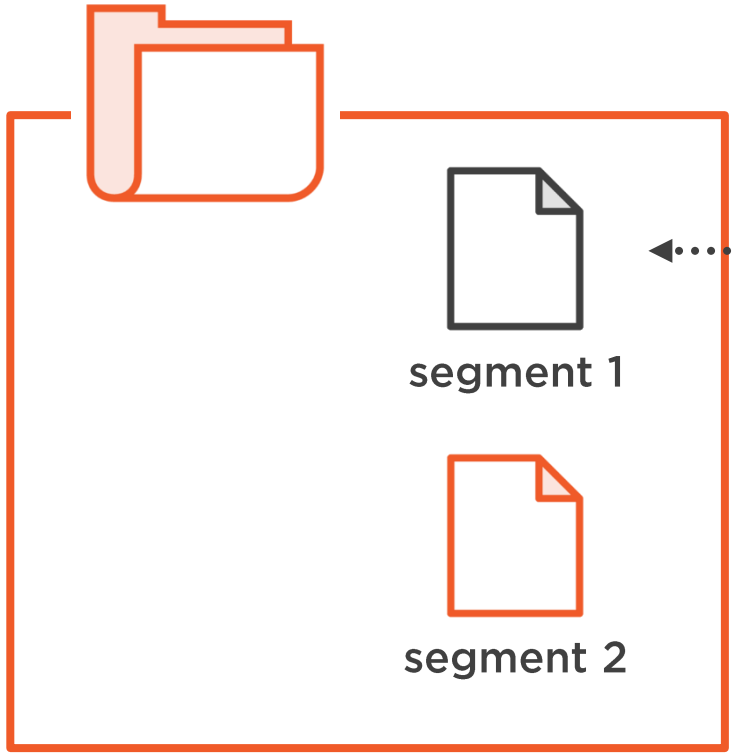

Day 2: Remove based on log.retention.ms

log.segment.bytes: 1GB
log.segment.ms: 604800000 (1 week)
log.retention.bytes: 2GB
log.retention.ms: 86400000 (1 day)

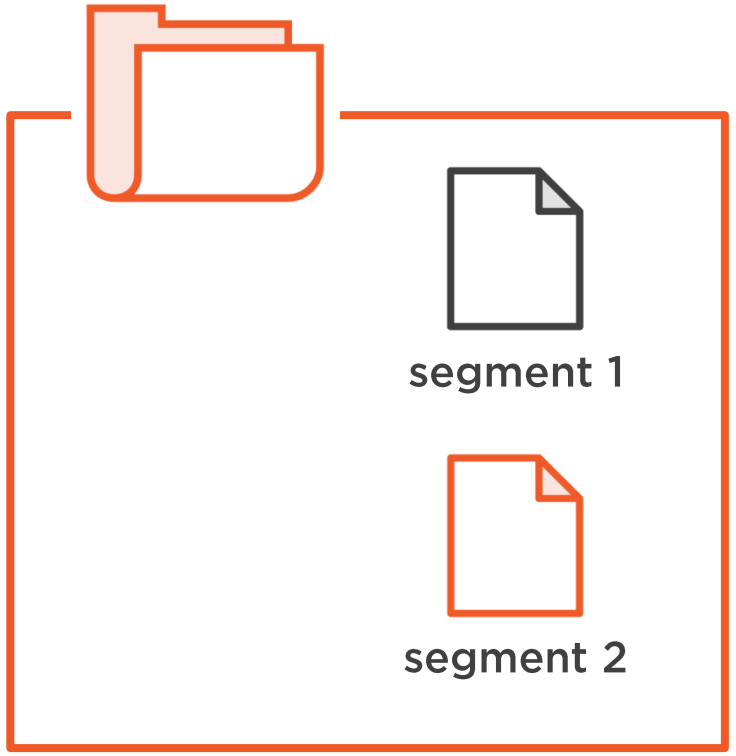

Defaults

log.segment.bytes: 1GB
log.segment.ms: 604800000 (1 week)
log.retention.bytes: 2GB
log.retention.ms: 86400000 (1 day)

Cluster Sizing

Usual bottlenecks:

- Disk size/throughput
- Network throughput

Always estimate the load
before deploying a cluster

Summary

- Why multiple brokers are needed
- Reliability guarantees
- Practical example
- Basics of data retention and cluster sizing